

CSPB 2270 Exam 1 Study Guide

BINARY SEARCH TREES

Overview

Binary Search Trees (BSTs) are a type of binary tree data structure that maintain a specific order of elements. In a BST, each node has a key or value, and the keys of nodes in the left subtree are less than the key of the current node, while the keys of nodes in the right subtree are greater. This property enables efficient searching, insertion, and deletion operations. BSTs offer fast average case performance for these operations, with a time complexity of $\mathcal{O}(\log(n))$ in balanced trees. They are commonly used in scenarios that require fast searching or ordered traversal of elements, such as dictionary implementations, range queries, and efficient data organization.

A Binary Search Tree (BST) and a linked list differ in their structure and organization of data. A BST is a binary tree where nodes are ordered based on their keys, allowing for efficient search, insertion, and deletion operations. In contrast, a linked list is a linear structure where elements are connected via pointers, without a specific ordering. Traversing a linked list requires sequentially following the next pointers, while a BST enables faster search and manipulation of elements due to its ordered nature.

A balanced tree is a tree data structure in which the heights of the left and right subtrees of any node differ by at most one. It ensures that the tree is evenly balanced and maintains efficient operations. Balanced trees are desirable because they provide improved performance for various operations such as searching, inserting, and deleting elements. By keeping the tree balanced, the height remains relatively low, which reduces the time complexity of operations and ensures that the worst-case scenario is avoided. Common examples of balanced trees include AVL trees, red-black trees, and B-trees.

BST's have some common terminology that are useful to describe certain aspects of the structure. The terminology that we use to describe a BST can be seen below:

- **Children:** Children are defined as nodes that are connected to a parent node, either left or right, in the hierarchical structure of the tree
- **Complete:** A BST where if all levels, except the last level, contain all possible nodes and all nodes in the last level are as far left as possible
- **Depth:** Depth is defined as the number of edges on the path from the root to the node in question
- **Edge:** An edge is a link from a parent node to its child node
- **Full:** A BST where every node contains 0 or 2 children
- **Height:** The height of a tree is the largest depth of any node in the tree
- **Internal Node:** An internal node is a node with at least one child
- **Leaf:** A leaf is a node with no children
- **Level:** The level of a BST is defined as nodes that have the same depth
- **Parent:** A parent is a node where one or more child nodes are directly connected to it
- **Perfect:** A BST where if all internal nodes have 2 children and all leaf nodes are at the same level
- **Root:** The root of a BST is referred to as the topmost node of the tree structure (The root node has no parent node)

The terms: **Full**, **Complete**, and **Perfect** define special types of BST's. The rest of the terms in the above list are components of the BST

BST's use a technique called recursion. Recursion is a programming technique where a function calls itself to solve a smaller instance of the same problem. It is useful in dealing with trees because trees are recursive data structures. Each node in a tree can be viewed as a smaller tree itself, consisting of its children and their descendants. Recursion allows us to traverse and manipulate tree structures by breaking down the problem into smaller sub problems that can be solved recursively. It simplifies the code by abstracting the complex tree operations into simpler recursive calls, leading to concise and elegant solutions. Additionally, recursion provides a natural way to handle tree-related algorithms, such as tree traversal, searching, insertion, deletion, and various other tree-based computations.

Common BST Operations

BST's require a number of operations for the functionality of it to be optimal. A lot of these operations can also be seen in linked lists. The operations that are used in BST's are the following:

- **Contains:** The contains function in a Binary Search Tree (BST) is used to check if a given value is present in the tree.
- **Get Node:** The getnode function in a Binary Search Tree (BST) is used to search for a node with a specified value in a tree.
- **Height:** The height operation in a BST recursively calculates the maximum number of edges from the root to any leaf node, providing a measure of the tree's vertical depth.
- **Initialize:** Initializing a node in a Binary Search Tree (BST) means creating a new node with a given value and setting its left and right child pointers to null or nullptr, indicating an empty subtree.
- **Insert:** The insert function in a Binary Search Tree (BST) is responsible for adding a new node to the tree while maintaining the binary search property, where values smaller than the current node are placed in the left subtree, and values larger than the current node are placed in the right subtree.
- **In Order:** An in-order traversal function in a Binary Search Tree (BST) visits all the nodes in the tree in ascending order of their values.
- **Post Order:** A post-order traversal function in a Binary Search Tree (BST) visits all the nodes in the tree in a specific order.
- **Pre Order:** A pre-order traversal function in a Binary Search Tree (BST) visits all the nodes in the tree in a specific order.
- **Remove:** The remove function in a Binary Search Tree (BST) is used to delete a node from the tree while preserving the binary search property. It involves finding the node to be removed, handling different cases based on the number of children the node has, and reorganizing the tree as necessary to maintain its structure and properties.
- **Size:** The size function in a Binary Search Tree (BST) is used to determine the number of nodes in a tree.
- **To Vector:** The tovector function in a Binary Search Tree (BST) is used to add values in a tree to a vector.

We can now look at each of these operations individually. The first operation that we will look at is the contains operation.

Contains In BST

Contains

Below is an example of the implementation of the contains operation in a BST in the context of C++:

```
bool BST::Contains(shared_ptr<bst_node> subtr, int data){
    if (subtr == nullptr) {
        return false;
    }
    else if (subtr->data == data) {
        return true;
    }
    else {
        return Contains(subtr->left, data) || Contains(subtr->right, data);
    }
}
```

The above function takes in a node that is present in a BST and searches for a specific value in the list. It utilizes recursion to search both the left and right subtrees of the current node and returns a boolean value determining if the value is present in the tree. The time complexity of this algorithm is the following:

Best Case Time Complexity	Worst Case Time Complexity
$\mathcal{O}(\log(n))$	$\mathcal{O}(n)$

The best case time complexity occurs when the tree is balanced. The worst case time complexity occurs when the node doesn't exist or if the node is at the leaf level of the tree.

The next operation that we will look at is the GetNode operation.

Get Node In BST

Get Node

Below is an example of the implementation of the Get Node operation in a BST in the context of C++:

```
shared_ptr<bst_node> BST::GetNode(shared_ptr<bst_node> sub, int data){
    if (sub == nullptr) {
        return shared_ptr<bst_node>(nullptr);
    }
    else if (sub->data == data) {
        return sub;
    }
    else {
        if (GetNode(sub->left, data)) {
            return GetNode(sub->left, data);
        }
        else {
            return GetNode(sub->right, data);
        }
    }
}
```

This operation searches a BST for a specific value and returns the node if it exists. Similar to the contains operation the GetNode operation utilizes recursion to search the left and right subtrees of the node that is fed into the function. The time complexity of this algorithm is the following:

Best Case Time Complexity	Worst Case Time Complexity
$\mathcal{O}(\log(n))$	$\mathcal{O}(n)$

The time complexity of this operation is the same as that of the operation that just searches for a value in a BST. When searching for the maximum or minimum value in a tree, the time complexity is the same for searching for a value in the tree. These time complexities are contingent upon the balancing of the tree.

The next operation that we will look at is the Height operation.

Height In BST

Height

Below is an example of the implementation of calculating the height of a BST in the context of C++:

```
int height(Node* root) {
    if (root == nullptr) {
        return 0; // An empty tree has height 0
    }
    else {
        // Recursively calculate the height of the left and right subtrees
        int leftHeight = height(root->left);
        int rightHeight = height(root->right);

        // Return the maximum height between the left and right subtrees,
        // plus 1 for the current node
        return std::max(leftHeight, rightHeight) + 1;
    }
}
```

The above function calculates the max number of edges to the leaf level in a BST. This function utilizes the concept of recursion to find which subtree has the most edges in it. This in turn is the height of the BST.

The next operation that we will look at is the Initialize operation.

Initializing In BST

Initializing

Below is an example of the implementation of initializing a node in a BST in the context of C++:

```
shared_ptr<bst_node> BST::InitNode(int data){
    shared_ptr<bst_node> ret(new bst_node());
    ret->data = data;
    ret->left = nullptr;
    ret->right = nullptr;
    return ret;
}
```

The above function initializes a node for a BST. It sets the value of the 'data' member to the input parameter that is passed in the function. This operation has a constant time complexity of $\mathcal{O}(1)$.

The next operation that we will look at is the insert operation.

Inserting In BST

Inserting

Below is an example of the implementation of inserting into a BST in the context of C++:

```
void BST::Insert(shared_ptr<bst_node> new_node){
    if (GetRoot() == nullptr) {
        SetRoot(new_node);
    }
    else {
        shared_ptr<bst_node> current_node = GetRoot();
        while (current_node != nullptr) {
            if (new_node->data < current_node->data) {
                if (current_node->left == nullptr) {
                    current_node->left = new_node;
                    current_node = nullptr;
                }
                else {
                    current_node = current_node->left;
                }
            }
            else {
                if (current_node->right == nullptr) {
                    current_node->right = new_node;
                    current_node = nullptr;
                }
                else {
                    current_node = current_node->right;
                }
            }
        }
    }
}
```

The insert operation inserts a node into a BST in the correct spot in the tree. This function operates by searching through a BST for where the node should be inserted. Similar to that of the search operation, the insert operation has the following time complexities:

Best Case Time Complexity	Worst Case Time Complexity
$\mathcal{O}(\log(n))$	$\mathcal{O}(n)$

The time complexities of this operation are again contingent upon the balancing of the tree. This operation effectively searches through the BST to try and find the appropriate spot for where it should be inserted. The best case scenario for inserting into a BST occurs when it is balanced and the worst case happens when the insertion happens at the leaf level.

The next operation that we will look at is the in order traversal of a BST.

In Order Traversal In BST

In Order

Below is an example of the implementation of the in-order traversal of a BST in the context of C++:

```
shared_ptr<bst_node> BST::InOrder(shared_ptr<bst_node> sub) {  
    if (sub == nullptr) {  
        return sub;  
    }  
    else {  
        if (InOrder(sub->left)) {  
            return InOrder(sub->left);  
        }  
        else {  
            return InOrder(sub->right);  
        }  
    }  
}
```

This function traverses a BST in order, meaning in ascending order. This function utilizes recursion to traverse the tree. It traverses in ascending order by traversing through all of the left subtrees and then the right subtrees. The time complexity of the in order traversal is $\mathcal{O}(n)$ regardless of the balancing of the tree.

The next operation that we will look at is the post order traversal of a BST.

Post Order Traversal In BST

Post Order

Below is an example of the implementation of post order traversal in a BST in the context of C++:

```
void postOrderTraversal(shared_ptr<bst_node> root) {  
    if (root == nullptr) {  
        return;  
    }  
  
    // Visit the left subtree  
    postOrderTraversal(root->left);  
  
    // Visit the right subtree  
    postOrderTraversal(root->right);  
  
    // Process the root node  
    std::cout << root->data << " ";  
}
```

The above function traverses a BST in post order fashion. In post-order traversal, the nodes of the BST are visited in the following order: left subtree, right subtree, and then the root. The recursive function is called first for the left subtree, then for the right subtree, and finally the root node is processed. Similar to the in order traversal, the time complexity of this operation is $\mathcal{O}(n)$ regardless of the balancing of the tree.

The next operation that we will look at is the pre order traversal of a BST.

Pre Order Traversal In BST

Pre Order

Below is an example of the implementation of pre order traversal in a BST in the context of C++:

```
void PreOrderTraversal(std::shared_ptr<bst_node> root) {  
    if (root == nullptr)  
        return;
```

```

// Process current node (root)
std::cout << root->data << " ";

// Recursively traverse left subtree
PreOrderTraversal(root->left);

// Recursively traverse right subtree
PreOrderTraversal(root->right);
}

```

The above function traverse a BST in a pre order fashion. In pre-order traversal, we visit the current node (root) first, then recursively traverse the left subtree, and finally recursively traverse the right subtree. Similar to that of the past two traversals of the tree, the time complexity of this operation is $\mathcal{O}(n)$ regardless of the balancing of the tree.

The next operation that we will look at is the remove operation of a BST.

Remove In BST

Remove

Below is an example of the implementation of the remove operation in a BST in the context of C++:

```

void BST::Remove(int data) {
    shared_ptr<bst_node> current_node = GetRoot();
    shared_ptr<bst_node> parent_node = nullptr;
    shared_ptr<bst_node> successor_node = nullptr;
    while (current_node != nullptr) {
        // Node found
        if (current_node->data == data) {
            // Remove a leaf node
            if (current_node->left == nullptr && current_node->right == nullptr) {
                if (parent_node == nullptr) {
                    SetRoot(nullptr);
                }
                else if (parent_node->left == current_node) {
                    parent_node->left = nullptr;
                }
                else {
                    parent_node->right = nullptr;
                }
            }
            // Remove left child only
            else if (current_node->right == nullptr) {
                if (parent_node == nullptr) {
                    SetRoot(current_node->left);
                }
                else if (parent_node->left == current_node) {
                    parent_node->left = current_node->left;
                }
                else {
                    parent_node->right = current_node->left;
                }
            }
            // Remove right child only
            else if (current_node->left == nullptr) {
                if (parent_node == nullptr) {
                    SetRoot(current_node->right);
                }
                else if (parent_node->left == current_node) {
                    parent_node->left = current_node->right;
                }
                else {
                    parent_node->right = current_node->right;
                }
            }
        }
    }
}

```

```

    }
}
// Remove node with two children
else {
    successor_node = current_node->right;
    while (successor_node->left != nullptr) {
        successor_node = successor_node->left;
    }
    int SuccessorData = successor_node->data;
    Remove(successor_node->data);
    current_node->data = SuccessorData;
}
break;
}
// Search to the right
else if (current_node->data < data) {
    parent_node = current_node;
    current_node = current_node->right;
}
// Search to the left
else {
    parent_node = current_node;
    current_node = current_node->left;
}
}
}
}

```

The above function removes a node from a BST and reorders the tree so that it still follows the same rules of a BST. The time complexity of this operation is the following:

Best Case Time Complexity	Worst Case Time Complexity
$\mathcal{O}(\log(n))$	$\mathcal{O}(n)$

Similar to other operations, the remove operation's time complexity is contingent upon the balancing of the tree.

The next operation that we will look at is the size operation of a BST.

Size Of BST

Size Of

Below is an example of the implementation of the size operation in a BST in the context of C++:

```

int BST::Size(shared_ptr<bst_node> subt){
    if (subt == nullptr) {
        return 0;
    }
    else {
        return 1 + Size(subt->left) + Size(subt->right);
    }
}

```

The above function implements recursion to count the number of nodes in a BST. The time complexity of this operation is $\mathcal{O}(n)$.

The last operation that we will look at is the to vector operation in a BST.

To Vector In BST

To Vector

Below is an example of the implementation of the to vector operation in a BST in the context of C++:

```

void BST::ToVector(shared_ptr<bst_node> subt, vector<int>& vec){

```

```
SetRoot(subt);
if (GetRoot() == nullptr) {
    return;
}
else {
    ToVector(subt->left, vec);
    vec.push_back(subt->data);
    ToVector(subt->right, vec);
}
}
```

The above function appends the data member of the nodes in a BST to a vector with the use of recursion. The time complexity of this operation is $\mathcal{O}(n)$ regardless of the balancing of the tree.

Binary Search Trees (BSTs) are a type of binary tree data structure in which each node has a key/value pair, and the keys in the left subtree are smaller than the key in the node, while the keys in the right subtree are larger. BSTs offer efficient search, insert, and delete operations with an average time complexity of $\mathcal{O}(\log(n))$, making them optimal for storing and retrieving data in sorted order. They provide an ordered and hierarchical structure that allows for fast searching and efficient traversal. Additionally, BSTs can be used to implement other data structures like sets, maps, and priority queues.



LINKED LISTS

Overview

In C++, a linked list is a data structure consisting of a sequence of nodes, where each node contains a value and a pointer/reference to the next node. Unlike arrays and vectors, linked lists do not store elements in contiguous memory locations. Instead, each node in a linked list is dynamically allocated and connected through pointers/references. This allows for efficient insertion and removal of elements at any position in the list, but it requires traversal from the beginning to access specific elements. Linked lists have a flexible size and can expand or shrink as needed by allocating or deallocating nodes, unlike arrays whose size is fixed. However, accessing elements in a linked list requires iterating through the nodes, while arrays and vectors provide direct access using indices.

Common Linked List Operations

Linked lists can be written a number of ways. In the context of this class the linked list was written as a class where the nodes of the linked list was written as a struct. Nodes in linked lists are analogous to elements in an array or vector, the main difference is that nodes contain data (An integer value in the context of this class) and a pointer that points to the next node in the linked list. In doubly linked lists the nodes will contain a pointer that points to the next and the previous nodes in the linked list.

Once a linked list has been created there are a number of operations that can be done on them. Here is a quick summary of some of the linked list operations:

- **Accessing & Modifying:** Accessing or modifying elements in a linked list involves traversing the list by following the node connections until reaching the desired position, and then performing the required operation on the accessed node, such as retrieving or updating its data value.
- **Adding:** Adding an element to a linked list involves creating a new node with the desired data, adjusting the node connections to include the new node in the appropriate position, and updating the necessary pointers to maintain the integrity of the list.
- **Checking For Empty:** Determining if a linked list is empty means checking whether the head node of the list is 'NULL', indicating that there are no elements in the list.
- **Checking Size:** Finding the size of a linked list involves traversing through the list and counting the number of nodes or elements present in the list.
- **Initializing:** Initializing a linked list involves creating the first node, setting its data value, and establishing the necessary connections to indicate the beginning of the list.
- **Removing:** Removing a node from a linked list involves adjusting the node connections to bypass the node to be removed, updating the necessary pointers, and deallocating the memory occupied by the removed node to maintain the integrity of the list.
- **Traversing:** Traversing a linked list means sequentially visiting each node in the list, starting from the head node, and accessing or performing operations on each node until the end of the list is reached.

We will now examine some of these operations in more detail. The first operation that we will look at is accessing and modifying the data in a linked list.

Accessing & Modifying Elements In Linked List

Accessing & Modifying

Below is an example of accessing and modifying elements in a linked list in the context of C++:

```
bool LinkedList::Contains(int data){
    shared_ptr<node> currPtr = top_ptr_;
    bool ret = false;
    while (currPtr != nullptr) {
        if (currPtr->data == data) {
            ret = true;
        }
        else {}
        currPtr = currPtr->next;
    }
    return ret;
}
```

}

The above is an example of how we access elements in a linked list. In summary, we start traversing the linked list from the head of the list and check if the current node matches what we are looking for. If it does match, then we modify a boolean value indicating that this element exists in the list. We could simply modify the value of linked list by mutating the data of the node if it passes the if statement. We could return the index of this element with slight modification that keeps track of the position of the current node that we are examining. Conversely, we could find the middle element by counting the total size of the list, and then adding some logic that would allow us to determine what the middle index of the linked list would be.

The next operation that we will examine is adding a node to a linked list.

Adding Elements In A Linked List

Adding Elements

Below is an example of adding a node to the linked list in the context of C++. Particularly, this function adds an element to the end of the linked list (The tail).

```
void LinkedList::Append(shared_ptr<node> new_node){
    new_node->data;
    new_node->next = nullptr;
    if (top_ptr_ == nullptr) {
        SetTop(new_node);
    }
    else {
        shared_ptr<node> currPtr = top_ptr_;
        while (currPtr->next != nullptr) {
            currPtr = currPtr->next;
        }
        currPtr->next = new_node;
    }
}
```

The above function will append a node to the end of the linked list. This function operates on the assumption that the node has already been initialized (Data member has been set). This function can be modified to handle the case where the value that is to be appended is an input parameter of this function. On the contrary, we can add an element at a certain index in the linked list.

Inserting Elements

Below is an example of adding a node to the linked list at a specified index in the context of C++:

```
void LinkedList::Insert(int offset, shared_ptr<node> new_node){
    shared_ptr<node> currPtr = top_ptr_;
    vector<shared_ptr<node>> ptrs;
    while (currPtr != nullptr) {
        ptrs.push_back(currPtr);
        currPtr = currPtr->next;
    }
    new_node->data;
    new_node->next = nullptr;
    if (offset == 0) {
        new_node->next = top_ptr_;
        SetTop(new_node);
    }
    else {
        currPtr = top_ptr_;
        for (int i = 0; i < ptrs.size(); i++) {
            if (i == offset - 1) {
                new_node->next = currPtr->next;
                currPtr->next = new_node;
            }
            else {}
            currPtr = currPtr->next;
        }
    }
}
```

```
    }  
    }  
}
```

The above function is adding a node to a linked list at a specific index of the linked list. Similar to the previous function that appends a node to the linked list, we can modify this function to append a value at a certain index by having an input parameter that pertains to the desired value to be inserted into the linked list.

The next operation that we will examine is checking if a linked list is empty.

Checking Empty Linked List

Checking Empty

Below is an example of checking if the linked list is empty in the context of C++:

```
bool LinkedList::IsEmpty() {  
    return (top_ptr_ == nullptr);  
}
```

The above function checks to see if the linked list is empty mainly by seeing if the head of the linked list is empty or not. This check was not necessarily a tenant of the assignment and that is why there are no comments explaining how the function operates.

The next operation that we will examine is checking for the size of a linked list.

Checking Size Of Linked List

Checking Size

Below is an example of checking for the size of a linked list in the context of C++:

```
int LinkedList::Size(){  
    shared_ptr<node> currPtr = top_ptr_;  
    int ret = 0;  
    while (currPtr != nullptr) {  
        ret++;  
        currPtr = currPtr->next;  
    }  
    return ret;  
}
```

The above function returns the size of a linked list. This function operates by traversing through the linked list and incrementing an integer value that indicates the size of the linked list.

The next operation that we will examine is initializing a node in a linked list.

Initializing A Node In A Linked List

Initializing

Below is an example of initializing a node in a linked list in the context of C++:

```
shared_ptr<node> LinkedList::InitNode(int data){  
    shared_ptr<node> ret(new node);  
    ret->data = data;  
    return ret;  
}
```

The above function initializes a node in a linked list to a specific value. This function does not necessarily update the pointer to the next node in the linked list but that is done in the other functions that have been previously defined.

The next operation that we will examine is removing a node from a linked list.

Removing A Node From A Linked List

Removing

The below function is an example of removing a node from a linked list in the context of C++:

```
void LinkedList::Remove(int offset){
    shared_ptr<node> currPtr = top_ptr_;
    vector<shared_ptr<node>> ptrs;
    while (currPtr != nullptr) {
        ptrs.push_back(currPtr);
        currPtr = currPtr->next;
    }
    if (offset == 0) {
        SetTop(top_ptr_->next);
    }
    else {
        currPtr = top_ptr_;
        for (int i = 0; i < ptrs.size(); i++) {
            if (i == offset - 1) {
                currPtr->next = currPtr->next->next;
                break;
            }
            else {}
            currPtr = currPtr->next;
        }
    }
}
```

The above function removes a node from a linked list at a specified offset of the linked list. We can remove a specific node from a linked list by searching for if the node exists in the linked list and then deleting the node from the linked list. This implementation requires a separate function that will do the removing of the node from the linked list.

The final common operation of the linked list is traversing the linked list. This operation is required for a lot of functions that were previously defined in the linked list. This is done by using a while loop that runs until the current pointer in the linked list is null or some other condition that should be met. We now will look at some extra operations that can be utilized in a linked list.

Extra Linked List Operations

There are some extra operations that can be done on a linked list. These involve operations from reversing a linked list to splitting a linked list. Here is a comprehensive list of extra operations in linked lists:

- **Concatenating:** Concatenating a linked list involves combining two or more linked lists into a single linked list, where the last node of the first list is connected to the first node of the second list, and so on, creating a continuous sequence of nodes.
- **Reversing:** Reversing a linked list involves changing the direction of the next pointers in each node, effectively flipping the order of the elements.
- **Sorting:** Sorting a linked list involves rearranging the nodes in a specific order, such as ascending or descending, based on the values contained in each node.
- **Splitting:** Splitting a linked list involves dividing a single linked list into two separate linked lists, typically based on a given condition or position.

We now can examine these operations in more detail. The first extra operation that we are looking at is the concatenating of a linked list.

Concatenating A Linked List

Concatenating

Below is an example of concatenating a linked list in the context of C++:

```

std::shared_ptr<node> concatenateLists(std::shared_ptr<node> head1,
                                       std::shared_ptr<node> head2) {
    if (head1 == nullptr)
        return head2;
    if (head2 == nullptr)
        return head1;
    std::shared_ptr<node> current = head1;
    while (current->next != nullptr) {
        current = current->next;
    }
    current->next = head2;
    return head1;
}

```

The above function takes in two heads of linked lists. It then traverses the linked list looking for the tail of the first list. It then updates the tail of the first linked list and has it point to the head of the second linked list.

The next extra operation that we are going to look at is reversing a linked list.

Reversing A Linked List

Reversing

Below is an example of reversing a linked list in the context of C++:

```

std::shared_ptr<node> reverseList(std::shared_ptr<node> head) {
    std::shared_ptr<node> current = head;
    std::shared_ptr<node> previous = nullptr;
    std::shared_ptr<node> next = nullptr;
    while (current != nullptr) {
        next = current->next;
        current->next = previous;
        previous = current;
        current = next;
    }
    return previous;
}

```

The above function reverses the nodes in a linked list. This is done by keeping track of the current, next, and previous nodes in the linked list. This entire operation could be done a lot easier if the linked list in question was a doubly linked list. But this function works on a singly linked list.

The next extra operation that we are going to look at is sorting a linked list. There are a multitude of ways that this can be accomplished, by choosing different sorting algorithms. The easiest way this is done is with the use of a bubble sort.

Sorting A Linked List

Sorting

Below is an example of sorting a linked list with the use of a bubble sort in the context of C++:

```

void bubbleSort(std::shared_ptr<node> head) {
    if (head == nullptr || head->next == nullptr) {
        return;
    }
    bool swapped;
    std::shared_ptr<node> current;
    std::shared_ptr<node> last = nullptr;
    do {
        swapped = false;
        current = head;
        while (current->next != last) {
            if (current->data > current->next->data) {
                // Swap the nodes
            }
        }
    }
}

```

```

        std::swap(current->data, current->next->data);
        swapped = true;
    }
    current = current->next;
}
last = current;
} while (swapped);
}

```

The above function sorts the data in a linked list with the use of a bubble sort algorithm. This algorithm does not require one to update the pointers of the nodes as it is simply mutating the data members of the nodes. One can sort the elements of the linked list in descending order by just changing the if statement that is present in the while loop.

The last extra operation that we are going to look at is splitting a linked list at a specific position.

Splitting A Linked List

Splitting

Below is an example of splitting a linked list in the context of C++:

```

std::pair<std::shared_ptr<node>, std::shared_ptr<node>>
    splitList(std::shared_ptr<node> head, int splitPosition) {
    std::shared_ptr<node> splitPtr = nullptr;
    std::shared_ptr<node> secondHead = nullptr;
    std::shared_ptr<node> current = head;
    int position = 0;
    // Traverse until the split position
    while (current != nullptr && position < splitPosition - 1) {
        current = current->next;
        position++;
    }
    if (current != nullptr) {
        splitPtr = current->next;
        current->next = nullptr; // Terminate the first list
        secondHead = splitPtr;
    }
    return {head, secondHead};
}

```

The above function splits a linked list at a specific index of the original linked list. This is done by returning a pair of head nodes after the function is done executing. In the function we iterate until we reach the specified index or if the index happens to be out of bounds. We then assign the head of the second linked list by grabbing the next node after the index position.

Other Implementations Of Linked Lists

There are many other implementations for linked lists. Some implementations of linked lists are the following:

- **Circular Linked List:** - Implementing a linked list as a circular linked list involves connecting the last node of the list to the first node, creating a circular structure.
- **Doubly Linked List:** - Implementing a linked list as a doubly linked list involves adding an additional pointer in each node that points to the previous node.
- **Stack Or Queue:** - Implementing a linked list as a stack or queue involves utilizing the respective operations of push/pop for a stack or enqueue/dequeue for a queue to insert or remove elements at one end of the linked list.

We now will look at these implementations one by one. The first implementation that we will look at is circular linked lists.

Circular Linked List

Circular Linked Lists

Below is an example of how to implement a circular linked list in the context of C++:

```
struct Node {
    int data;
    std::shared_ptr<Node> next;
    Node(int val) : data(val), next(nullptr) {}
};

void traverseCircularLinkedList(const std::shared_ptr<Node>& head) {
    if (head == nullptr) {
        return;
    }

    std::shared_ptr<Node> current = head;

    do {
        std::cout << current->data << " ";
        current = current->next;
    } while (current != head);

    std::cout << std::endl;
}

int main() {
    // Create a circular linked list with three nodes: 1 -> 2 -> 3 -> (back to 1)
    std::shared_ptr<Node> head = std::make_shared<Node>(1);
    std::shared_ptr<Node> second = std::make_shared<Node>(2);
    std::shared_ptr<Node> third = std::make_shared<Node>(3);

    head->next = second;
    second->next = third;
    third->next = head; // Make it circular

    // Traverse the circular linked list
    traverseCircularLinkedList(head); // Output: 1 2 3

    return 0;
}
```

The above is an example of how to make a linked list circular. The crux of making a linked list circular is to have the last node of the list point to the head of the list. This is done inside the main function above.

Next, we look at how to implement a doubly linked list.

Doubly Linked List

Doubly Linked Lists

Below is an example of doubly linked lists in the context of C++:

```
#include <iostream>
#include <memory>

struct Node {
    int data;
    std::shared_ptr<Node> prev;
    std::shared_ptr<Node> next;
    Node(int val) : data(val), prev(nullptr), next(nullptr) {}
};

void traverseDoublyLinkedListForward(const std::shared_ptr<Node>& head) {
    std::shared_ptr<Node> current = head;

    while (current != nullptr) {
```

```

        std::cout << current->data << " ";
        current = current->next;
    }

    std::cout << std::endl;
}

void traverseDoublyLinkedListBackward(const std::shared_ptr<Node>& tail) {
    std::shared_ptr<Node> current = tail;

    while (current != nullptr) {
        std::cout << current->data << " ";
        current = current->prev;
    }

    std::cout << std::endl;
}

int main() {
    // Create a doubly linked list with three nodes: 1 <-> 2 <-> 3
    std::shared_ptr<Node> head = std::make_shared<Node>(1);
    std::shared_ptr<Node> second = std::make_shared<Node>(2);
    std::shared_ptr<Node> third = std::make_shared<Node>(3);

    head->next = second;
    second->prev = head;
    second->next = third;
    third->prev = second;

    // Traverse the doubly linked list forward and backward
    traverseDoublyLinkedListForward(head); // Output: 1 2 3
    traverseDoublyLinkedListBackward(third); // Output: 3 2 1

    return 0;
}

```

The above is an example of how to create a doubly linked list. In this example, the main difference between our linked list class that we created previously is that the node structure has a previous pointer as well as a next pointer. This in essence is how a doubly linked list is created.

Lastly, we look at how to implement a stack and queue with the use of a linked list.

Stacks & Queues With Linked Lists

Stacks & Queues

Below is an example of how to implement a linked list with the stack and queue data structure. This is first done by looking at the structure of the nodes in the linked list in the context of C++:

```

struct Node {
    int data;
    std::shared_ptr<Node> next;
    Node(int val) : data(val), next(nullptr) {}
};

```

We then move on to looking how we implement the stack data structure with this node structure:

```

class Stack {
private:
    std::shared_ptr<Node> top;

public:
    Stack() : top(nullptr) {}
}

```



```
void push(int val) {
    std::shared_ptr<Node> newNode = std::make_shared<Node>(val);
    newNode->next = top;
    top = newNode;
}

void pop() {
    if (top) {
        std::shared_ptr<Node> temp = top;
        top = top->next;
        temp.reset();
    }
}

int peek() {
    if (top) {
        return top->data;
    }
    throw std::runtime_error("Stack is empty.");
}

bool isEmpty() {
    return top == nullptr;
}
};
```

Next, we look at the queue data structure to see how it would be implemented:

```
class Queue {
private:
    std::shared_ptr<Node> front;
    std::shared_ptr<Node> rear;

public:
    Queue() : front(nullptr), rear(nullptr) {}

    void enqueue(int val) {
        std::shared_ptr<Node> newNode = std::make_shared<Node>(val);
        if (!front) {
            front = newNode;
            rear = newNode;
        } else {
            rear->next = newNode;
            rear = newNode;
        }
    }

    void dequeue() {
        if (front) {
            std::shared_ptr<Node> temp = front;
            front = front->next;
            temp.reset();
        }
    }

    int peek() {
        if (front) {
            return front->data;
        }
        throw std::runtime_error("Queue is empty.");
    }
};
```

```
    bool isEmpty() {  
        return front == nullptr;  
    }  
};
```

And lastly, we look at how these would be implemented:

```
int main() {  
    Stack stack;  
    stack.push(10);  
    stack.push(20);  
    stack.push(30);  
  
    while (!stack.isEmpty()) {  
        std::cout << stack.peek() << " ";  
        stack.pop();  
    }  
    // Output: 30 20 10  
  
    Queue queue;  
    queue.enqueue(5);  
    queue.enqueue(15);  
    queue.enqueue(25);  
  
    while (!queue.isEmpty()) {  
        std::cout << queue.peek() << " ";  
        queue.dequeue();  
    }  
    // Output: 5 15 25  
  
    return 0;  
}
```

The above is only one way this can be achieved. There are numerous different ways this outcome can be produced.

Overall, linked lists are a very powerful data structure in object oriented programming. They allow for easy traversal, insertion, and deletion of elements. The biggest downside to linked lists is that someone cannot easily access elements inside the linked list without traversing the list itself. Doubly linked lists allow for easy traversal in each direction but require additional logic to update parameters that are found in the structure of the linked list.

MEMORY ALLOCATION

Overview

Memory allocation in C++ refers to the process of assigning and managing memory for variables, objects, and data structures during program execution. C++ provides two main ways of allocating memory: stack allocation and heap allocation. Stack allocation, performed automatically by the compiler, is used for local variables and function calls, and the memory is automatically reclaimed when variables go out of scope. Heap allocation, on the other hand, allows dynamic memory allocation using operators like 'new' and 'delete', providing flexibility in managing memory for objects that have a longer lifespan or unknown size at compile time. Proper memory allocation and deallocation are important for avoiding memory leaks, optimizing resource usage, and ensuring program stability and performance.

Stack Allocation

Stack allocation pertains to the memory allocation of a C++ program at runtime. It is used for local variables, as well as function calls. When these types of variables or calls go out of scope (scope refers to the region of a program where a particular variable or identifier is visible and can be accessed) then the memory is deallocated and reclaimed. The stack for a given program has a limited size and if one exceeds its capacity, then stack overflow can occur (stack overflow occurs when the call stack, which is a limited region of memory used for function calls and local variables, exceeds its allocated size). Some common examples of stack allocation can be seen below:

- **Array Variables:** When one defines an array with a fixed size, the memory for that array is allocated to the stack. This can apply to other data structures such as lists as well.
- **Function Calls:** When a function is called in a program, the function parameters and local variables are allocated to the stack. When this function is finished executing, the memory that was allocated is reclaimed.
- **Local Variables:** When a variable is defined inside a function or a block, the memory for that variable is allocated on the stack. Similarly to these other examples, when the local variable goes out of scope it is reclaimed.
- **Recursive Functions:** Since recursive functions are essentially calls to functions (the same function inside a function) the memory for these calls are allocated on the stack. The same is true for when the call terminates or goes out of scope in the context of deallocating the memory.

To demonstrate this further we will look at example of this type of memory allocation.

Stack Allocation Example

Below we have a simple example of stack allocation in C++:

```
#include <iostream>

int main() {
    int x = 5;    // Stack-allocated variable

    // Print the value of x
    std::cout << "The value of x: " << x << std::endl;

    {
        int y = 10;    // Stack-allocated variable inside a nested scope

        // Print the value of y
        std::cout << "The value of y: " << y << std::endl;
    }

    // y is out of scope here and no longer accessible

    return 0;
}
```

The above example demonstrates defining local variables to the stack. The variable 'y' in this case no longer becomes accessible when it leaves the scope of its definition. The same can be said for the variable 'x' when the program finishes execution, it is no longer available and thus the memory that it consumed is reclaimed.

Heap Allocation

Heap allocation pertains to the allocation of memory that is ‘dynamic’. Dynamic in this context simply means that the memory that is consumed by the variable type or data structure is allocated at runtime (runtime refers to when a program is executing not when it is compiled) and must be deallocated. Heap allocation is usually defined with keywords like ‘new’ and the deallocation is usually referred to with ‘delete’. When discussing heap allocation, one usually discusses the concept of pointers. Traditional pointers require the manual deallocation of memory when it is no longer needed. Smart pointers on the other hand operate similarly to that of stack allocation in that the memory that is consumed by a variable or data structure is reclaimed when that variable or data structure goes out of scope. Some examples of the types of heap allocation can be seen below:

Heap Allocation Example

Below is an example of heap allocation in C++:

```
#include <iostream>

int main() {
    // Heap allocation
    int* ptr = new int[5]; // Allocate an array of 5 integers on the heap

    // Assign values to the array elements
    for (int i = 0; i < 5; i++) {
        ptr[i] = i + 1;
    }

    // Print the array elements
    for (int i = 0; i < 5; i++) {
        std::cout << ptr[i] << " ";
    }
    std::cout << std::endl;

    // Heap reallocation
    int* newPtr = new int[10]; // Allocate a new array of 10 integers on the heap

    // Copy elements from the original array to the new array
    for (int i = 0; i < 5; i++) {
        newPtr[i] = ptr[i];
    }

    // Release the memory occupied by the original array
    delete[] ptr;

    // Print the elements of the new array
    for (int i = 0; i < 10; i++) {
        std::cout << newPtr[i] << " ";
    }
    std::cout << std::endl;

    // Release the memory occupied by the new array
    delete[] newPtr;

    return 0;
}
```

The above example demonstrates the allocation and deallocation of memory with the use of heap allocation. This memory is only being allocated at runtime instead of at compile. Failure to properly release memory from the stack can cause what is referred to as a memory leak. Memory leaks are when heap allocation is not properly freed and it can cause problems with the execution of program up to the point where a program can actually run out of memory. This can simply be avoided with the use of smart pointers.

Unlike stack allocation, the types of variables that can be allocated on the heap are limitless. The main issue that arises when dealing with heap allocation is when dealing with reference and pointer parameters. The differences between the two parameters can be seen below:

- **Pointer Parameter:** A pointer parameter is declared using the '*' symbol and requires the use of pointer syntax when passing arguments.
- **Reference Parameter:** A reference parameter is declared using the '&' symbol and allows the function to directly access and modify the value of the variable passed as an argument.

This can best be understood with an example, as seen below.

Pointer & Reference Parameter Example

Below is an example of working with pointer and reference parameters in C++:

```
#include <iostream>

// Reference parameter example
void modifyByReference(int& value) {
    value += 10;
}

// Pointer parameter example
void modifyByPointer(int* value) {
    (*value) += 10;
}

int main() {
    int num = 5;

    // Pass num as a reference parameter
    modifyByReference(num);
    std::cout << "After modifyByReference: " << num << std::endl;    // Output: 15

    // Pass the address of num as a pointer parameter
    modifyByPointer(&num);
    std::cout << "After modifyByPointer: " << num << std::endl;      // Output: 25

    return 0;
}
```

The above example works with both pointer and reference parameters. The syntax of each is slightly different but both work on the concept of accessing the place in memory for where the variable is defined in memory and modifying its value without having to create a copy of the variable to do so.

Although the above parameters achieve the same result, they have some distinct differences. These differences can be summed up below:

- **Addressing:** Pointers store the memory address of an object, allowing for operations like pointer arithmetic and direct manipulation of memory. References, on the other hand, are simply aliases for existing objects and do not have an address of their own.
- **Initialization:** Pointers can be declared without initialization, allowing them to be assigned later. References must be initialized when declared and cannot be reassigned to refer to a different object.
- **Nullability:** Pointers can be assigned a 'nullptr' value, indicating that they don't point to any valid memory address. References, on the other hand, must always refer to a valid object and cannot be null.
- **Reassignment:** Pointers can be reassigned to point to different objects, while references cannot be reseated to refer to a different object after initialization.
- **Syntax:** Pointers are declared using the '*' symbol, while references are declared using the '&' symbol. When passing a pointer parameter, you need to use pointer dereferencing (*) to access the value. With a reference parameter, you can directly use the reference as if it were the original variable.

Overall, stack and heap allocation are two different memory allocation mechanisms in computer programs. Stack allocation is used for local variables and function call frames and is managed automatically by the compiler. It is fast and efficient, but limited in size and scope. Heap allocation, on the other hand, allows dynamic memory allocation at runtime, providing flexibility but also requiring manual memory management. Heap allocation is suitable for objects that need to persist

beyond the scope of a function or have a variable size. It requires explicit allocation and deallocation, making it more prone to memory leaks and fragmentation. Choosing between stack and heap allocation depends on factors such as variable lifetime, size, and desired control over memory management.



SORTING ALGORITHMS

Overview

Sorting algorithms are fundamental algorithms used to arrange elements in a specific order, typically in ascending or descending order. There are various sorting algorithms available, each with its own characteristics and performance trade-offs. Some popular sorting algorithms include bubble sort, insertion sort, selection sort, merge sort, quicksort, and heapsort. These algorithms employ different techniques, such as comparison-based sorting, divide-and-conquer, or the use of auxiliary data structures, to efficiently sort elements. The choice of sorting algorithm depends on factors such as the size of the data, the desired time complexity, and the nature of the data being sorted. Efficient sorting algorithms have a significant impact on the performance of applications that require ordered data.

Types of Sorting Algorithms

There are a multitude of different types of sorting algorithms. Below is a table that lists the sorting algorithms with a brief description, the best, average, and worst case time complexity as well as the worst case data sequence of each algorithm:

Algorithm	Description	$\Omega(n)$	$\Theta(n)$	$\mathcal{O}(n)$	Worst Data Sequence
Bubble Sort	Bubble sort is a simple comparison-based sorting algorithm that repeatedly swaps adjacent elements if they are in the wrong order.	$\Omega(n^2)$	$\Theta(n^2)$	$\mathcal{O}(n^2)$	Reverse Order
Insertion Sort	Insertion sort is an efficient comparison-based sorting algorithm that builds the final sorted array one element at a time.	$\Omega(n)$	$\Theta(n^2)$	$\mathcal{O}(n^2)$	Reverse Order
Merge Sort	Merge sort is a divide-and-conquer algorithm that recursively divides the input array into smaller subarrays, sorts them, and then merges them to produce a sorted array.	$\Omega(n \log(n))$	$\Theta(n \log(n))$	$\mathcal{O}(n \log(n))$	Any
Quick Sort	Quick sort is a divide-and-conquer algorithm that selects a pivot element and partitions the array into two subarrays, one with elements less than the pivot and the other with elements greater than the pivot.	$\Omega(n)$	$\Theta(n^2)$	$\mathcal{O}(n^2)$	Reverse Order
Radix Sort	Radix sort is a non-comparative sorting algorithm that sorts elements based on their digits or bits.	$\Omega(n)$	$\Theta(d(n+b))$	$\mathcal{O}(d^2(n+b))$	Different Digits
Selection Sort	Selection sort is a simple comparison-based sorting algorithm that works by dividing the input list into two parts: a sorted sublist and an unsorted sublist.	$\Omega(n)$	$\Theta(n^2)$	$\mathcal{O}(n^2)$	Reverse Order
Shell Sort	Shell sort is an in-place comparison-based sorting algorithm that is an extension of the insertion sort algorithm.	$\Omega(n \log(n))$	$\Theta(n^{3/2})$	$\mathcal{O}(n^2)$	Reverse Order

Bubble Sort Algorithm

Overview

The bubble sort is a simple sorting algorithm that works by repeatedly comparing adjacent elements in an data set and swapping them if they are in the wrong order. The algorithm starts at the beginning of the data set and compares the first two elements. If the first element is greater than the second element, they are swapped. The algorithm then moves on to the next two elements and repeats the process. This continues until the end of the data set is reached. At this point, the largest element in the data set will be in the last position. The algorithm then starts again at the beginning of the data set and repeats the process. This time, the second largest element will be moved to the second-to-last position, and so on. The algorithm continues until all of the elements in the data set are sorted in ascending order.

Bubble Sort Algorithm

Below is the bubble sort algorithm in the context of C++:

```
void Sorting::bubblesort(vector<int>& data){
// Iterate through the vector "data"
for (int i = 0; i < data.size() - 1; i++) {
    // Compare adjacent elements in the vector
```

```

for (int j = 0; j < data.size() - i - 1; j++) {
    // Check to see if next element in vector is greater than that of the
    // current element
    if (data.at(j) > data.at(j+1)) {
        // Create a temporary integer value for that of the current element
        int temp_val = data.at(j);
        // Assign the current element with that of the next element
        data.at(j) = data.at(j+1);
        // Assign the next element with that of the current element
        data.at(j+1) = temp_val;
    }
    else {}
}
}
}

```

The bubble sort algorithm is a relatively inefficient sorting algorithm when discussing runtime complexity. The runtime complexity of the bubble sort algorithm is as follows:

Best Case $\Omega(n)$	Average Case $\Theta(n)$	Worst Case $\mathcal{O}(n)$
$\Omega(n)$	$\Theta(n^2)$	$\mathcal{O}(n^2)$

At its best, the bubble sort algorithm will run at a linear runtime complexity. In the average and worst case, the bubble sort algorithm will run at a quadratic runtime complexity. These runtime complexities are not considered great, but the given the simplicity of the algorithm it is rather efficient.

Merge Sort Algorithm

Overview

The merge sort algorithm is a divide-and-conquer algorithm for sorting a data set. This algorithm sorts a data set by recursively breaking it down into smaller and smaller sub-data set until they are sorted. Once the sub-data sets have been sorted, they are sorted back together into one final sub-data set. The merge sort algorithm utilizes an auxiliary data-set (A temporary data-set) to store the sorted elements of the original data set. This auxiliary data-set is then copied back to the original data set. Similar to the quick sort algorithm, the merge sort algorithm requires the use of a separate algorithm in its implementation.

The algorithm that is used in the merge sort algorithm is the merge algorithm. This algorithm is responsible for taking two data sets and combining them into one data set. The merge algorithm can be seen below:

Merge Algorithm

Below is an example of the merge algorithm in the context of C++:

```

void Sorting::merge(vector<int>& left, vector<int>& right, vector<int>& result) {
    // Calculate the size of the needed vector for the merged vector
    int mergedSize = left.size() + right.size();
    // Set the positions of each vector
    int mergePos = 0;
    int leftPos = 0;
    int rightPos = 0;
    // Create a temporary vector that will hold the merged values
    vector<int> mergedNumbers(mergedSize);
    // Continue while loop as long as there are elements left in each vector to
    // be merged
    while (leftPos < left.size() && rightPos < right.size()) {
        // Insert the element in the left vector into the merged vector if it is less
        // than that of the right vector
        if (left.at(leftPos) <= right.at(rightPos)) {
            mergedNumbers.at(mergePos) = left.at(leftPos);
            ++leftPos;
        }
        // Insert the element in the right vector into the merged vector if it is
        // greater than that of the left vector
    }
}

```



```

        else {
            mergedNumbers.at(mergePos) = right.at(rightPos);
            ++rightPos;
        }
        // Increment the merge position
        ++mergePos;
    }
    // Merge the left vector into the mergedNumbers vector
    while (leftPos < left.size()) {
        mergedNumbers.at(mergePos) = left.at(leftPos);
        ++leftPos;
        ++mergePos;
    }
    // Merge the right vector into the mergedNumbers vector
    while (rightPos < right.size()) {
        mergedNumbers.at(mergePos) = right.at(rightPos);
        ++rightPos;
        ++mergePos;
    }
    // Copy the elements from the mergedNumbers vector into the result vector
    for (mergePos = 0; mergePos < mergedSize; ++mergePos) {
        result.push_back(mergedNumbers.at(mergePos));
    }
}

```

The main responsibility of the above algorithm is to sort the elements of the two data sets, insert them into the auxiliary data set, and then copy that auxiliary data set back to the original data set. The general procedure of this algorithm can be explained below:

- First calculate the needed size of the auxiliary data set and initialize the indexes of the left, right, and auxiliary data set
- The sub-data sets are then traversed through where the elements of the left and right sub-data sets are compared, after the comparison, the appropriate element is inserted into the auxiliary data set
- The above procedure is repeated until the index of either the left or right sub-data set reaches the back of the respective data set
- The remaining elements of the left and right sub-data sets are then added to auxiliary data set
- The auxiliary data set is then copied back to the resultant data set
- This process will repeat until the left and right sub-data sets are sorted

The merge algorithm is where the sorting of the merge sort algorithm occurs as well as where the merging of two sub-data sets occurs. The recursive nature of the merge algorithm will sort each halves of the original data set.

The next algorithm is the implementation of the merge sort algorithm. The merge algorithm is responsible for creating recursive calls to the input data set so that it can be sorted in a timely manner. This algorithm can be seen below:

Merge Sort Algorithm

Below is an example of the merge sort algorithm in the context of C++:

```

void Sorting::mergesort(vector<int>& data){
    // The vector is empty or only has one element
    if (data.size() <= 1) {
        return;
    }
    // The vector is greater than 1
    else {
        // Calculate the midpoint of the vector
        int midPoint = data.size() / 2;
        // Create a left half of the vector

```

```

        vector<int> leftHalf(data.begin(), data.begin() + midPoint);
        // Create a right half of the vector
        vector<int> rightHalf(data.begin() + midPoint, data.end());
        // Merge the left and right half of the vector, recursively
        mergesort(leftHalf);
        mergesort(rightHalf);
        // Create an empty vector
        vector<int> result;
        // Call the merge method and merge the two halves together
        merge(leftHalf, rightHalf, result);
        // Copy the results to data
        data = result;
    }
}

```

This algorithm takes an input data set and recursively splits the input data set until it reaches a size of one. Once the sub-data set reaches a size of one, it then gets passed in as an input parameter into the merge algorithm so that it can eventually be copied back to the original data set that is passed into the algorithm. The general procedure for how this data set works is the following:

- The original data set that is passed into the algorithm is first checked if the data set is of size one, if it is, then the data set is returned, otherwise it continues on to the else block statement
- The midpoint of the data set is then calculated such that it can be split into two halves
- A left and right data set is then created to represent each half of the split data data set
- These sub-data sets are recursively fed into the function until the returned data set is of size one
- Once the returned sub-data set is of size one, it is fed into the merge algorithm so that the two sub-data sets can be combined into one data set
- The above process is repeated until the algorithm no longer merges two halves together, indicating that the data set has been sorted

This algorithm effectively splits an input data set until the returned data set is of size one. Once this happens, the merge algorithm then effectively merges two separate sub-data sets in a manner that they will be sorted. This process repeats over and over until all elements from the original data set have been sorted.

Similar to the quick sort algorithm, the merge sort algorithm is a highly efficient algorithm. Its recursive nature allows for less complex runtimes which makes it an optimal choice for sorting elements. The runtime complexities of this algorithm can be seen below:

Best Case $\Omega(n)$	Average Case $\Theta(n)$	Worst Case $\mathcal{O}(n)$
$\Omega(n \log(n))$	$\Theta(n \log(n))$	$\mathcal{O}(n \log(n))$

Unlike the quick sort algorithm, the best case runtime complexity for the merge sort algorithm occurs when the list is already sorted. Because of nature of the merge sort algorithm, if the list is already sorted then it typically will only have to make n comparisons when sorting the list. On the contrary, the worst case time complexity for the merge sort algorithm occurs when the data set is sorted in reverse order. This is due to the splitting of the data sets and the recursive nature of the algorithm.

On the contrary, the space complexity of the merge sort algorithm tends to be of the order n , $\mathcal{O}(n)$. This is because each time the recursive function is called it has to occupy memory. If the merge sort algorithm uses the 'divide and conquer in place' strategy, where the sorted halves of the data set are stored in the original data set, the space complexity will decrease.

Quick Sort Algorithm

Overview

The quick sort algorithm is a divide-and-conquer algorithm for sorting a data set. This algorithm works by repeatedly partitioning the data set around a pivot element, and then recursively sorting the two sub-data sets created by the partition. The pivot element that is chosen for this algorithm is typically the middle element of the data set, but it is not necessarily required for the algorithm to work. When the partitioning of elements occurs, elements that are smaller than the pivot are moved to the left of the pivot. On the contrary, the elements that are larger than the pivot move to the right of the pivot. Once this partition is completed, the two sub-data sets are recursively sorted. The quick sort algorithm is often regarded as

a very efficient sorting algorithm and works well on large data sets.

The quick sort algorithm requires a separate algorithm than the main algorithm for it to operate correctly. The secondary algorithm that is required is the 'quicksort_partition' algorithm. The primary role of this algorithm is to iterate through the vector and partition it so that elements can be sorted in a manner that is appropriate. This algorithm can be seen below.

Quick Sort Partition Algorithm

Below is the quick sort partition algorithm in the context of C++:

```
int Sorting::quicksort_partition(vector<int>& data, int low_idx, int high_idx){
    // Define variables to track elements in the "data" vector
    int mid_idx = low_idx + (high_idx - low_idx) / 2;
    int pivot = data.at(mid_idx);
    bool finished = false;
    // Begin traversing through the vector
    while (!finished) {
        // Increment lower index until value greater than or equal to that of the
        // pivot is found
        while (data.at(low_idx) < pivot) {
            low_idx += 1;
        }
        // Decrement higher index until value less than or equal to that of the
        // pivot is found
        while (pivot < data.at(high_idx)) {
            high_idx -= 1;
        }
        // If no elements remain, then exit the loop
        if (low_idx >= high_idx) {
            finished = true;
        }
        // Begin the process of swapping values
        else {
            // Keep track of the current value at the lower index
            int temp_val = data.at(low_idx);
            // Swap the values of lower and higher indexes
            data.at(low_idx) = data.at(high_idx);
            data.at(high_idx) = temp_val;
            // Increment and decrement index values
            low_idx += 1;
            high_idx -= 1;
        }
    }
    return high_idx;
}
```

The general method for how this algorithm works is the following:

- The function calculates the middle index of the 'data' vector
- This function then sets the pivot element at the element of the middle index
- We then iterate through the the data set until we find a value on the left side of the pivot that is either greater than or equal to that of the pivot value
- After this, we iterate through the data set until we find a value on the right side of the pivot that is less than or equal to that of the pivot value
- We then check to see if the two sub-vectors (The values to left and right of the pivot) are sorted with the 'if (low_idx ≥ high_idx)' statement
 - If this statement evaluates to true, we exit the while loop
 - If this statement is false, we swap the values of the lower index with that of the value at the higher index, and proceed to continue partitioning the data set
- The above process will repeat until the the previous if statement evaluates to true

- ‘high_idx’ is returned from the partition algorithm to represent the index where elements to the left of it are less than or equal to it as well as the elements to the right of it are greater than or equal to it

The above algorithm correctly places the elements that are less than that of the pivot to the left of it and the values that are greater than or equal to that of the pivot to the right of it. This process repeats itself until the condition previously mentioned is achieved. Once this condition is met, then we recursively sort the elements on the left and right of the pivot.

The next algorithm, which is the implementation of the quick sort algorithm, recursively sorts the sub-data sets that are created with the partitioning algorithm. This is done by making recursive calls to itself while utilizing the ‘quicksort_partition’ algorithm to partition the left and right sub data sets of the data set. This algorithm can be seen below.

Quick Sort Algorithm

Below is the quick sort algorithm in the context of C++:

```
void Sorting::quicksort(vector<int>& data, int low_idx, int high_idx){
    // Check to see if the lower index is greater than or equal to that of higher index
    if (low_idx >= high_idx) {
        return;
    }
    else {
        // Determine the pivot point of the current vector
        int low_end_idx = quicksort_partition(data, low_idx, high_idx);
        // Recursively sort the lower half of the vector
        quicksort(data, low_idx, low_end_idx);
        // Recursively sort the higher half of the vector
        quicksort(data, low_end_idx + 1, high_idx);
    }
}
```

The general method for how this algorithm works is the following:

- First check to see if there is only one element in the data set that is being sorted, otherwise continue to the next logic block
- If the data set is greater in size than one element, then we partition the data set such that it will have a lower half and upper half that can be sorted after
- The function then proceeds to finding the partition point (The point where elements to the left of it are less than it and the elements to the right are greater than or equal to it) with the ‘quicksort_partition’ algorithm
- After the partition point has been found, the algorithm proceeds to recursively sort the left and right halves of the data set that are set with the help of the partition point until the partition reaches size one

The goal of the quick sort algorithm is to partition a data set to the point such that the elements to the left of the pivot are less than or equal to the pivot and the elements to the right of the pivot point are greater than or equal to that of the pivot. This process will repeat until the data set that is being partitioned is of size one. After the data set has been partitioned, it is sorted recursively with a call to ‘quicksort’.

The quick sort algorithm is a very efficient algorithm for data sets of all sizes. Its recursive nature allows for different sizes of data sets to be sorted efficiently in a timely matter implicating that it is an optimal choice for a sorting algorithm. The runtime complexity of the quick sort algorithm is as follows:

Best Case $\Omega(n)$	Average Case $\Theta(n)$	Worst Case $\mathcal{O}(n)$
$\Omega(n \log(n))$	$\Theta(n \log(n))$	$\mathcal{O}(n^2)$

The quick sort algorithm typically will have a runtime complexity of $\Theta(n \log(n))$ with the worst runtime complexity of $\mathcal{O}(n^2)$. The worst case runtime complexity occurs when the data set that is being fed into the algorithm is sorted in reverse order.

VECTORS

C++ Vector Overview

In C++, a vector is a dynamic array-like container provided by the Standard Template Library (STL). It is a sequence container that can dynamically resize itself to accommodate elements. Unlike an array, which has a fixed size determined at compile time, a vector can grow or shrink in size as elements are added or removed. Vectors offer various advantages over arrays, such as automatic memory management, the ability to easily change their size, and built-in functions for common operations like inserting, erasing, and accessing elements. Vectors provide similar functionality to arrays but with added flexibility and convenience.

How Do You Declare And Initialize A Vector With Values?

Initializing A Vector

The way we initialize a vector in C++ with values is like the following:

```
#include <vector>

int main() {
    // Declare and initialize a vector of integers
    std::vector<int> numbers = {1, 2, 3, 4, 5};

    // Declare and initialize a vector of strings
    std::vector<std::string> names = {"Alice", "Bob", "Charlie"};

    // Declare an empty vector
    std::vector<double> values;

    return 0;
}
```

Vector Operations In C++

Here are some common vector operations in C++. First, we look at how we access elements in a vector.

Accessing Vector Elements

Accessing Elements

Here is how someone accesses elements in a vector:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Accessing elements using the indexing operator []
int firstElement = numbers[0]; // Access the first element (1)
int thirdElement = numbers[2]; // Access the third element (3)

// Accessing elements using the at() member function
int secondElement = numbers.at(1); // Access the second element (2)
int fourthElement = numbers.at(3); // Access the fourth element (4)
```

Here is how someone adds elements to a vectors.

Adding Vector Elements

Adding To End

Adding an element at the end of a vector with **push.back()**:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};
```

```
// Add the number 6 to the end of the vector
numbers.push_back(6);

// Print the modified vector
for (int i = 0; i < numbers.size(); i++) {
    std::cout << numbers[i] << std::endl;
}
```

Here is how vectors can be cleared.

Clearing Vectors

Clearing All Elements

Clearing all the elements from a vector with **clear()**:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Remove all elements from the vector
numbers.clear();

// Print the modified vector
for (int i = 0; i < numbers.size(); i++) {
    std::cout << numbers[i] << std::endl;
}
```

If you want to find the capacity of a vector, this is how its done.

Finding Capacity Of Vector

Finding Capacity

Finding the capacity of the vector with **capacity()**:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Find the capacity of the vector
int capacity = numbers.capacity();

// Print the capacity of the vector
std::cout << "The capacity of the vector is: " << capacity << std::endl;
```

To find the size of vector, this is how its done.

Finding Size Of Vector

Finding Size

Finding the size of the vector with **size()**:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Find the size of the vector
int size = numbers.size();

// Print the size of the vector
std::cout << "The size of the vector is: " << size << std::endl;
```

Next, we look at how we modify elements in a vector.

Modifying Vector Elements

Modifying Elements

Here is how someone modifies elements in a vector:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Modifying elements using the indexing operator []
numbers[0] = 10; // Modify the value of the first element to 10
numbers[2] = 30; // Modify the value of the third element to 30

// Modifying elements using the at() member function
numbers.at(1) = 20; // Modify the value of the second element to 20
numbers.at(3) = 40; // Modify the value of the fourth element to 40
```

There are multiple ways to remove an element at an index of a vector, here is one.

Removing Element Of Vector

Removing Element At Index

Removing an element at an index of a vector with `erase()`:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Remove the element at index 2 from the vector
int element_to_remove = numbers[2];
numbers.erase(numbers.begin() + 2);

// Print the modified vector
for (int i = 0; i < numbers.size(); i++) {
    std::cout << numbers[i] << std::endl;
}
```

Another way to remove an element is to remove the last element of a vector.

Removing Last Element Of Vector

Removing Last Element

Removing the last element from a vector with `pop_back()`:

```
std::vector<int> numbers = {1, 2, 3, 4, 5};

// Remove the last element from the vector
int last_element = numbers.back();
numbers.pop_back();

// Print the modified vector
for (int i = 0; i < numbers.size(); i++) {
    std::cout << numbers[i] << std::endl;
}
```

Vectors can be sorted as well, here is how this is done.

Sorting Vectors

Sorting

Vectors can be sorted with sorting algorithms or the built in STL function `sort()`:

```
std::vector<int> numbers = {5, 2, 8, 1, 9};
```



```
// Sort the vector in ascending order
std::sort(numbers.begin(), numbers.end());

// Print the sorted vector
for (const auto& num : numbers) {
    std::cout << num << " ";
}
std::cout << std::endl;
```

Here is a summary of some common vector operations.

- **Accessing Elements:**
 - Using the indexing operator `[]` to access individual elements by their index.
 - Using the `at()` member function to access elements with bounds checking.
 - Using iterators to traverse the vector and access elements.
- **Iterating & Traversing:**
 - Using range-based **for** loops or iterators to iterate over the elements of the vector.
 - Using algorithms like `for_each()`, `transform()`, or `accumulate()` from the **algorithm** library to perform operations on each element.
- **Modifying Elements:**
 - Using the assignment operator `=` to modify individual elements directly.
 - Using the `push_back()` member function to add elements at the end of the vector.
 - Using the `pop_back()` member function to remove the last element from the vector.
 - Using the `insert()` member function to insert elements at a specified position.
 - Using the `erase()` member function to remove elements at a specified position or range.
- **Size & Capacity:**
 - Using the `size()` member function to get the current number of elements in the vector.
 - Using the `empty()` member function to check if the vector is empty.
 - Using the `resize()` member function to change the size of the vector.
 - Using the `reserve()` member function to allocate memory for a specified number of elements.
- **Sorting & Manipulating Order:**
 - Using the `sort()` function from the **algorithm** library to sort the elements in the vector.
 - Using the `reverse()` function from the **algorithm** library to reverse the order of elements.
 - Using algorithms like `find()`, `count()`, or `binary_search()` from the **algorithm** library to search for elements.

Benefits Of Vectors Over Arrays

There are many benefits to using a vector over an array. Some examples of this are dynamic size, automatic memory management, convenient functions, range checking, and copy and assignment. These are not all of the advantages, rather just some common examples. Diving into these examples a little bit more:

- **Automatic Memory Management:** Vectors handle memory allocation and deallocation automatically. When elements are added or removed, the vector manages the underlying memory for you. Arrays, on the other hand, require manual memory management.
- **Convenient Functions:** Vectors provide a range of built-in functions that make working with elements more convenient. For example, vectors have functions like `push_back()` to add elements at the end, `pop_back()` to remove elements from the end, and `insert()` to insert elements at specific positions. Arrays lack such built-in functions, requiring manual manipulation of elements.
- **Copy and Assignment:** Vectors can be easily copied or assigned to another vector using the assignment operator (`=`) or the copy constructor. This simplifies the process of creating copies or working with multiple vectors. Arrays do not have built-in copy or assignment mechanisms, requiring manual copying of each element.

- **Dynamic Size:** Vectors have a flexible and resizable size, allowing elements to be added or removed easily. In contrast, arrays have a fixed size determined at compile-time, making it difficult to change their size dynamically.
- **Range Checking:** Vectors perform bounds checking to ensure that accessing elements stays within the valid range. This helps prevent accessing elements outside the vector's size, which can lead to runtime errors. Arrays do not perform such range checking, allowing access to elements beyond the declared size, which can lead to undefined behavior.

Overall, vectors provide more flexibility, convenience, and safety compared to arrays in terms of size management, memory handling, and operations on elements.

C++ Vector Functions Recap

Here is a quick summary of vector functions in C++:

- **back():** The 'back()' method in C++ returns a reference to the last element in a vector.
- **capacity():** The 'capacity()' method in C++ returns the current storage capacity of a vector, which represents the maximum number of elements that the vector can hold without requiring reallocation of memory.
- **clear():** The 'clear()' method in C++ removes all elements from a vector, effectively emptying it.
- **erase():** The 'erase()' method in C++ removes one or a range of elements from a vector, based on the specified position or iterator.
- **find():** The 'find()' function in C++ searches for a specified element in a vector and returns an iterator pointing to the first occurrence of the element.
- **front():** The 'front()' method in C++ returns a reference to the first element in a vector.
- **pop_back():** The 'pop_back()' method in C++ removes the last element from a vector.
- **push_back():** The 'push_back()' method in C++ adds an element to the end of a vector.
- **resize():** The 'resize()' method in C++ changes the size of the vector, either increasing or decreasing it.
- **size():** The 'size()' method in C++ returns the number of elements in a vector, representing the current size of the vector.