

Exceptional Control Flow

- Up to now: two mechanisms for changing control flow:
 - Jumps and branches (e.g. if/else, switch, for, while)
 - Call and return (e.g. function call, return)Both react to changes in *program state*
- Insufficient for a useful system:
 - Difficult to react to changes in *system state*
 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - User hits Ctrl-C at the keyboard
 - System timer expires
- System needs mechanisms for “exceptional control flow”

Exceptional Control Flow

1. Exceptions :

- Low level, implemented with hardware and OS
 - a. **interrupts**, e.g. I/O interrupts (caused by disk data arrival, network data arrival, ..)
 - b. **traps**, intentional, e.g. system-call, break-points
 - c. **faults**, unintentional but possibly recoverable, e.g. page fault
 - d. **aborts**, unintentional and unrecoverable

2. Process context switch

Shell Lab

- Implemented by OS software and hardware timer
- control flow from one process to another, scheduled by OS kernel

3. Signals

- Implemented by OS software
- a small message sent to a process from kernel

4. Nonlocal jumps: `setjmp()` and `longjmp()`

- Implemented by C runtime library

Process

- What is a *process*?
 - A process is **an instance of a running program**
 - 2 key abstractions:
 - each process seems to have **exclusive use of the CPU**
 - each process seems to have **exclusive use of main memory**
 - The reality is:
 - Process executions interleaved, (alternatively executed)
 - Address spaces managed by virtual memory system
 - difference with a **program**:
 - program can be an executable (e.g. binary file), saved in disk
 - process is an instance of a program running in memory

Process basic states

From a simple perspective, we can think of a process as being in one of three states

- Running (Runnable)
 - Process is either executing, or waiting to be executed and will eventually be *scheduled* (i.e., chosen to execute) by the kernel
- Stopped
 - Process execution is *suspended* and will not be scheduled until further notice (i.e. SIGCONT signal)
- Terminated
 - Process is stopped permanently

Four basic process control functions

- `fork()`
- `exec()` (and other variants such as `execve()`)
- `exit()`
- `wait()` (and variants like `waitpid()`)

Standard on all UNIX-based systems

Running Processes

- **int exec()**

- to load and run **another** program in the **current** process
 - program is changed but process ID is not
- **Never returns on successful execution**
- More useful variant is **int execve()**

```
int execve(char *filename, char *argv[], char *envp[])
```

- **filename**: executable file to run
- **argv[]**: argument list, `argv[0] == filename`
- **envp[]**: environment variable list
 - in format of "name=value" strings (e.g., `USER=droh`)

The global variable, "environ", which saves all environment variables.

Returning from Processes

- **`void exit(int status)`**

- Terminates the current process with a return status (error code)
- Normal return with status 0 (other numbers indicate an error)
- OS frees resources such as heap memory and open file descriptors and so on...
- Reduce to a **zombie** state
 - process is terminated but there's still its entry in the process table (managed by OS)
 - Must wait to be **reaped** by the parent process (or the init process if the parent died)
 - To reap a child, process, use `wait()` or `waitpid()`
 - `exit` is called **once** but **never** returns!

Waiting for Processes

- **pid_t wait(int *child_status)**

- suspends current process until **one of its children** terminates
- return value is the pid of the child process that terminated
 - When wait returns a pid > 0, child process has been **reaped**
- update child_status, can be used to check the exit status and reason:
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED` ...
- More useful variant is **waitpid()**
 - `pid_t waitpid(pid_t pid, int *status, int options)`
 - can specify a child by **pid**, can use more **options**
- search online for more

Process Example (1)

```
pid_t child_pid = fork();

if (child_pid == 0){
    /* only child comes here */

    printf("Child!\n");

    exit(0);
}
else{

    printf("Parent!\n");
}
```

- What are the possible output (assuming fork succeeds) ?
 - Child!
Parent!
 - Parent!
Child!
- How to get the child to always print first?

How to guarantee the order of execution?

```
int status;
pid_t child_pid = fork();

if (child_pid == 0){
    /* only child comes here */

    printf("Child!\n");

    exit(0);
}
else {
    waitpid(child_pid, &status, 0);
    printf("Parent!\n");
}
```

- Waits until the child has terminated.
- Parent can inspect exit status of child using 'status'
 - WEXITSTATUS(status)
- Output always:
Child!
Parent!

Process Example (2)

```
int status;
pid_t child_pid = fork();
char* argv[] = {"/bin/ls", "-l", NULL};
char* env[] = {..., NULL};

if (child_pid == 0){
    /* only child comes here */

    execve("/bin/ls", argv, env);
    printf("Child");
    /* will child reach here? */
}
else{
    waitpid(child_pid, &status, 0);

    ... parent continue execution...
}
```

- Will child reach here?




Process Graphs

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

fork.c



events in child/parent **are partially ordered!**
arrows determine relative orders,
if no arrow, order is unknown

Signals

- A *signal* is a small message that notifies a process that an event
 - sent from the kernel (sometimes at the request of another process)
 - sent to a process
 - signal type is identified by small integer ID's (1-30)
 - only information in a signal is its ID and the fact that it arrived

<i>ID</i>	<i>Name</i>	<i>Default Action</i>	<i>Corresponding Event</i>
2	SIGINT	Terminate	Interrupt (e.g., ctrl-c from keyboard)
9	SIGKILL	Terminate	Kill program (cannot override or ignore)
11	SIGSEGV	Terminate & Dump	Segmentation violation
14	SIGALRM	Terminate	Timer signal
17	SIGCHLD	Ignore	Child stopped or terminated
	SIGTSTP	default action: Stop	Ctrl-Z from keyboard
	SIGCONT	default action: continue run for a stopped process	

Sending signals

- Kernel *sends* (delivers) a signal to a *destination process* by updating some state in the context of the destination process
- Kernel sends a signal for one of the following reasons:
 - Kernel has detected a system event such as Ctrl-C (SIGINT), divide-by-zero (SIGFPE), or the termination of a child process (SIGCHLD)
 - The user used a `kill` command (in terminal)
 - Another program called the `kill()` function

```
int kill(pid_t pid, int sig);    //send signal "sig" to process pid
                                // if send signal to a process group, use -pid, pid means the PGID
```

Receiving signals

- A destination process *receives* a signal when it is forced by the kernel to react in some way to the delivery of the signal
- Three possible ways to react:
 - *Ignore* the signal (do nothing)
 - *Terminate* the process (with optional core dump)
 - *Catch* the signal by executing a user-level function called *signal handler*

Receiving signals

- Signal Handler

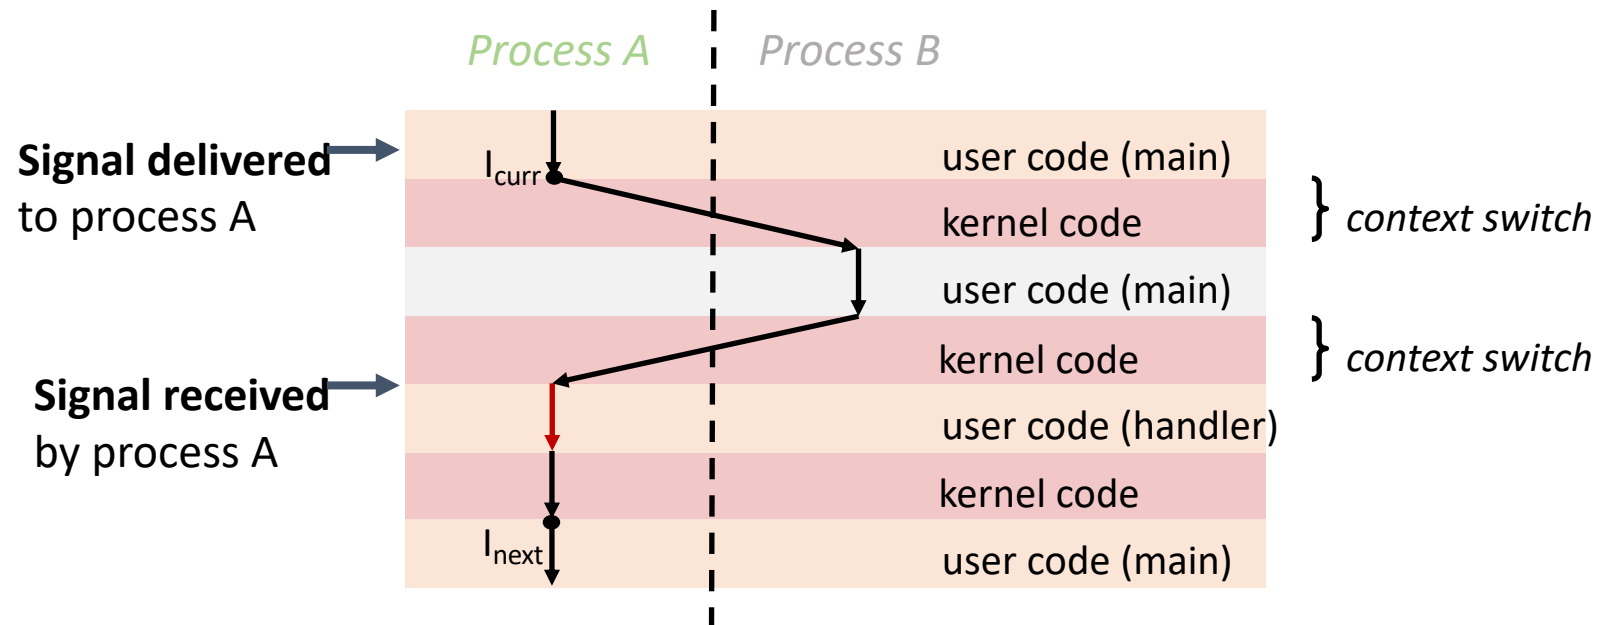
- a user-level function to catch a specific signal
- **Separate** flow of control in the same process
- Resumes normal flow of control upon returning
- Can be called **anytime** when the appropriate signal is fired
- implement in the form: `void handler(int signum){ ... }`
- install the handler by:

```
handler_t *signal(int signum, handler_t *handler)
```

```
//signal function will modify the default action associated with signum
```


Receiving signals

- Receiving a signal is **non-queuing**
 - There is only one bit in the context per signal
 - Receiving 1 or 300 SIGINTs looks the same to the process
- Signals are received at a **context switch**



Receiving signals

- Blocking/Unblocking signals
 - Sometimes code needs to run through a section that can't be interrupted
 - use `sigprocmask()` to block/unblock/set blocking masks
 - support functions:
 - `sigemptyset()`, `sigaddset()`, `sigfillset()`, `sigdelset()`

```
sigset_t mask, prev_mask;

Sigemptyset(&mask);
Sigaddset(&mask, SIGINT);

/* Block SIGINT and save previous blocked set */
Sigprocmask(SIG_BLOCK, &mask, &prev_mask);

    /* Code region that will not be interrupted by SIGINT */

/* Restore previous blocked set, unblocking SIGINT */
Sigprocmask(SIG_SETMASK, &prev_mask, NULL);
```

Tsh functionalities and specifications (1/7)

- **command line**, consists of a “name” and 0-or-more “arguments”

e.g. `“./myspin 10”` .name: `./myspin`, 1 argument: `“10”`

if name is a **built-in command**:

tsh executes the command immediately in “tsh” process

a built-in command corresponds to a function in `tsh.c` -- you need to implement these functions

if name is a pathname of an **executable**:

tsh forks a child process, and load/run the executable in the child

Hints:

1. This part is left for you to implement.
2. **parseline()** is already implemented, to parse the cmdline into `char** argv`
3. use **fork()** to fork a child process, and use **execve()** to load/run an executable

Tsh functionalities and specifications (2/7)

- **background job & foreground job**

If command line ends with “&”,

it is a **background** job, i.e. shell process won't wait the child process

otherwise, it's a **foreground** job,

i.e. shell process will wait the child process to terminate/stopped

Hints:

1. **parseline()** will return you a “bg/fg” boolean indicator.

2. implement your waiting function in **waitfg()**

A Recommendation (provided in the writeup) :

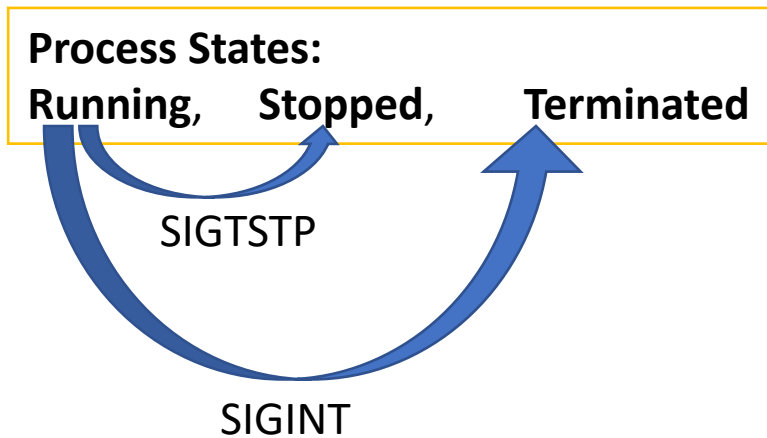
- In `waitfg`, use a busy loop around the `sleep` function.

loop iterates **until this process is no longer the foreground process**, e.g terminated/stopped/move to background

Tsh functionalities and specifications (3/7)

- **Catch Signals**

- Typing **CTRL+C** should cause a **SIGINT** signal to be sent to the current **foreground job** (i.e. the foreground process and its descendent processes), which is to **terminate** these processes.
- Typing **CTRL+Z** should cause a **SIGTSTP** signal to be sent to the current **foreground job**, which is to **stop** these processes.



Tsh functionalities and specifications (4/7)

- **Catch Signals**

Hints:

1. you need implement handlers for SIGINT, SIGTSTP **to catch the SIGINT/SIGTSTP** signals
 - implement: `sigint_handler, sigtstp_handler`.
 - default action for SIGINT/SIGTSTP is to terminate/stop its own process (tsh), but here you need to terminate/stop the foreground processes.
2. to send signals to some process, use **kill()**
`int kill(pid_t pid, int sig);` //pass **-pid** to send sig to the **process group** whose GID==pid
3. create a separate process group for the child process, otherwise all the children processes will be in the same group with “tsh”.
call “setpgid(0,0)” between fork() and execve() in the child process.
`setpgid(pid_t pid, pid_t pgid)` // sets the PGID of the process specified by pid to pgid

Tsh functionalities and specifications (5/7)

- Built-in Commands

- **quit**: terminates “tsh”
- **jobs**: list the existing (running/stopped) jobs
- **bg** <job>: restart <job> by sending it SIGCONT signal, then run in background
- **fg** <job>: restart <job> by sending it SIGCONT signal, then run in foreground
 - <job> can be process id, or job id (with a “J” prefix, e.g. bg J1)

Tsh functionalities and specifications (6/7)

- Tsh should reap all of its zombie children.
 - A child process terminated but not reaped becomes a zombie process
 - So tsh should keep records of all its existing jobs and **reap** them when they terminates
 - Besides, if any job terminates/stops because of signals (e.g. SIGINT/SIGTSTP), the **tsh** should recognize the event and **print a message** with JobID, PID and signal number which terminates/stops the job.

Tsh functionalities and specifications (7/7)

- Tsh should reap all of its zombie children.

Hints:

1. the global var “jobs” is used to keep the list of jobs, you can use **addjob()**, **deletejob()** to update this “jobs”.
2. **SIGCHLD** is a signal which is sent to a process when any of its children terminated/stopped
You can catch SIGCHLD by implement: `sigchld_handler()`
3. You can put a **waitpid()** in `sigchld_handler()` to reap terminated processes.
`pid_t waitpid(pid_t pid, int* status, int options)`
if **pid==-1**, means wait for any child
options:
WUNTRACED: returns also when the process is **stopped**. (Otherwise, only return for terminated)
WNOHANG: returns 0 **immediately** is no stopped/terminated child process. (Otherwise, will wait)
4. the **status** output from **waitpid()** can be used to test the exit status/signals
e.g. **WIFSTOPPED(status)**, **WIFEXITED(status)**, **WIFSIGNALED(status)**, ...

Other hints

- In `eval()`, the parent must **block SIGCHLD signals** before it forks the child, and then **unblock these signals**, after it calls `addjob` (to add job in the “jobs”).

Also remember to unblock SIGCHLD signals for the child before it calls `execve()`.

in Parent:

```
//block SIGCHLD
```

```
fork()
```

```
...
```

```
addjob()
```

```
//unblock SIGCHLD
```

in Child:

```
//block SIGCHLD
```

```
fork()
```

```
...
```

```
//unblock SIGCHLD
```

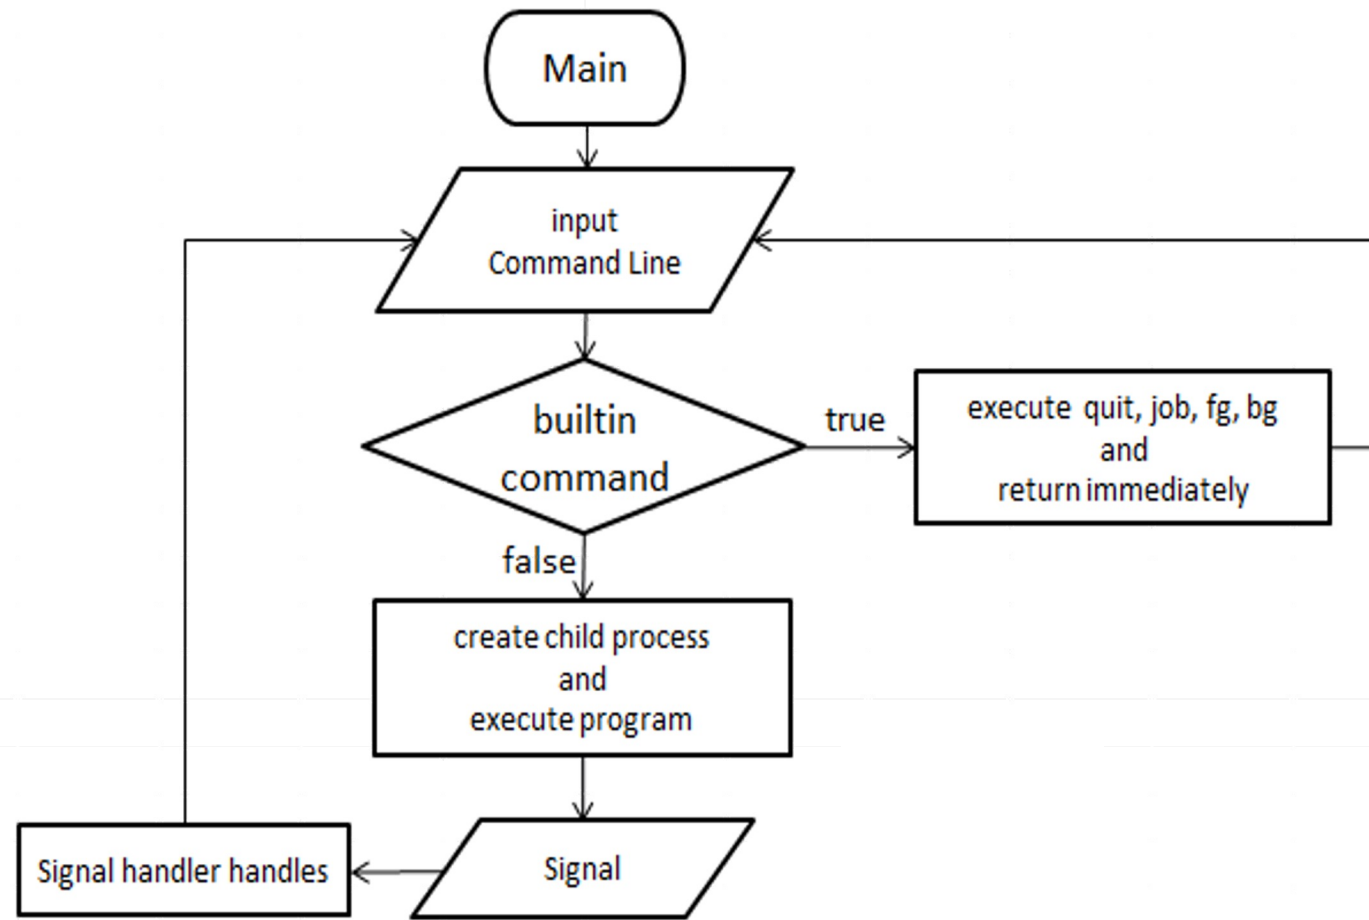
```
execve()
```

- This is to avoid the race condition where the child is reaped by `sigchld_handler` (and thus removed from the job list) *before* the parent calls `addjob()`
- related function: **`sigprocmask()`**, **`sigemptyset()`**, **`sigaddset()`**..

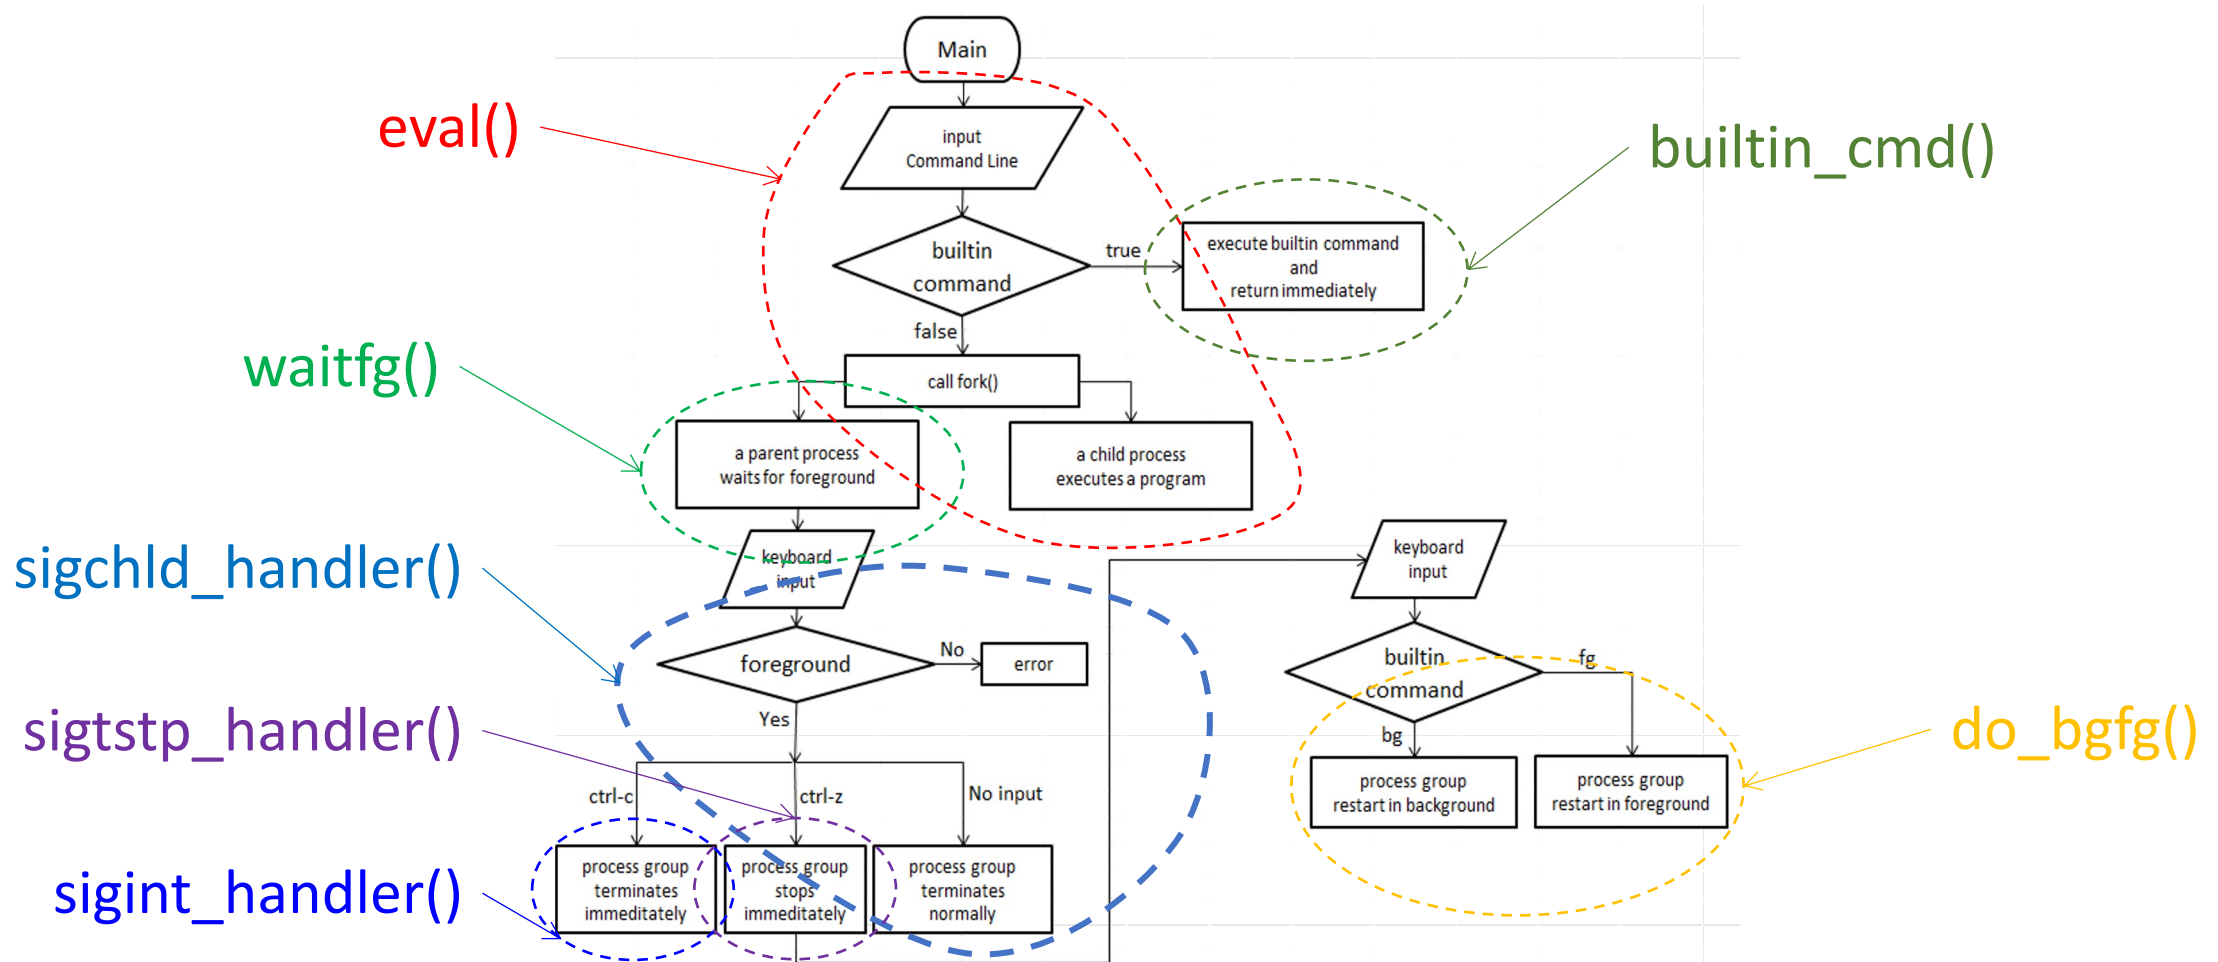
Introduction to Shell Lab

- Implementation of a simple Unix shell program (mainly job control): [tsh.cc](#)
 - eval: Main routine that parses and interprets the command line. [70 lines]
 - builtin_cmd: Recognizes and interprets the built-in commands: quit, fg, bg, and jobs. [25 lines]
 - do_bgfg: Implements the bg and fg built-in commands. [50 lines]
 - waitfg: Waits for a foreground job to complete. [20 lines]
 - sigchld_handler: Catches SIGCHLD signals. 80 lines]
 - sigint_handler: Catches SIGINT (ctrl-c) signals. [15 lines]
 - sigtstp_handler: Catches SIGTSTP (ctrl-z) signals. [15 lines]

Logical Flowchart of TSH (1/2)



Logical Flowchart of TSH (2/2)



* Figure 8.24 in the textbook is a good reference to start the shell lab

Example Codes in the Textbook

- Figure 8.24

```
code/ecf/shellx.c
1  /* eval - Evaluate a command line */
2  void eval(char *cmdline)
3  {
4      char *argv[MAXARGS]; /* Argument list execve() */
5      char buf[MAXLINE];    /* Holds modified command line */
6      int bg;               /* Should the job run in bg or fg? */
7      pid_t pid;           /* Process id */
8
9      strcpy(buf, cmdline);
10     bg = parseline(buf, argv);
11     if (argv[0] == NULL)
12         return; /* Ignore empty lines */
13
14     if (!builtin_command(argv)) {
15         if ((pid = Fork()) == 0) { /* Child runs user job */
16             if (execve(argv[0], argv, environ) < 0) {
17                 printf("%s: Command not found.\n", argv[0]);
18                 exit(0);
19             }
20         }
21
22         /* Parent waits for foreground job to terminate */
23         if (!bg) {
24             int status;
25             if (waitpid(pid, &status, 0) < 0)
26                 unix_error("waitfg: waitpid error");
27         }
28         else
29             printf("%d %s", pid, cmdline);
30     }
31     return;
32 }
```

```
33
34 /* If first arg is a builtin command, run it and return true */
35 int builtin_command(char **argv)
36 {
37     if (!strcmp(argv[0], "quit")) /* quit command */
38         exit(0);
39     if (!strcmp(argv[0], "&")) /* Ignore singleton & */
40         return 1;
41     return 0; /* Not a builtin command */
42 }
```

code/ecf/shellx.c

Figure 8.24 eval evaluates the shell command line.