



Exceptional Control Flow: Signals and Nonlocal Jumps

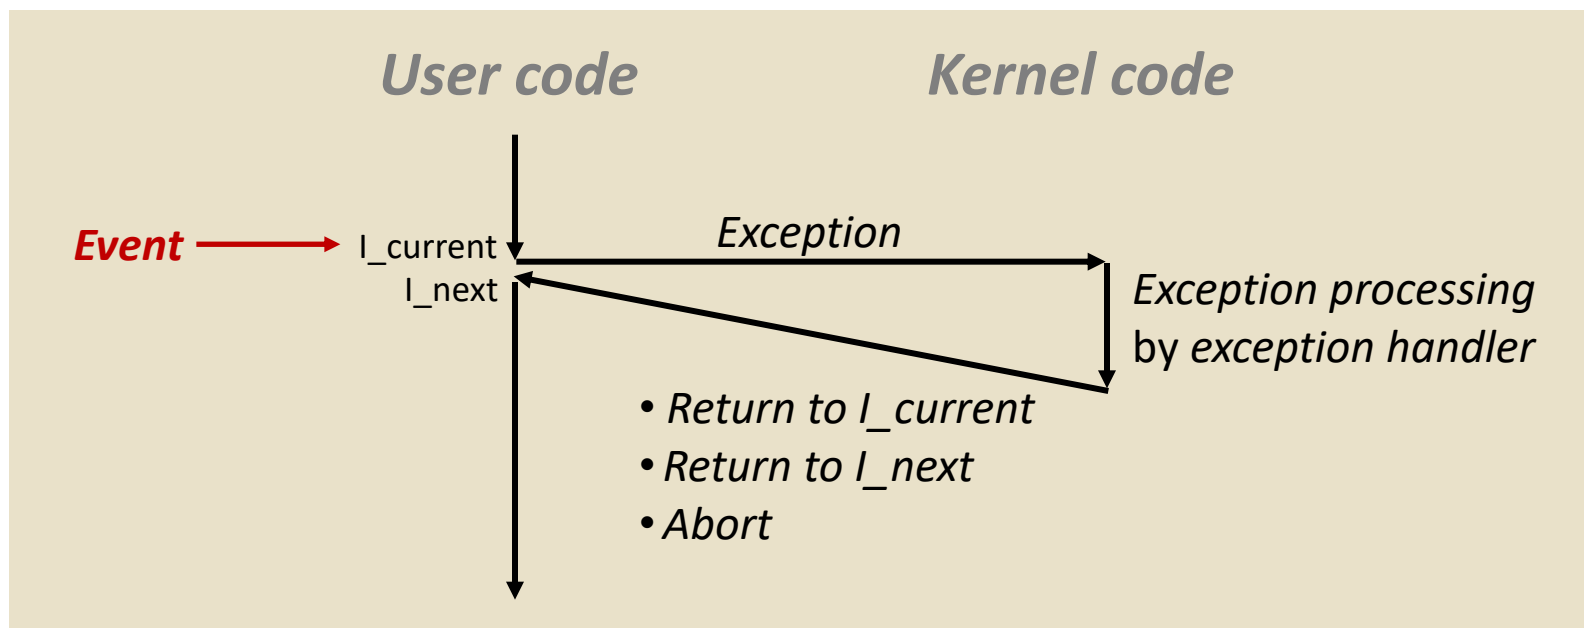
These slides adapted from materials provided by the textbook authors.

Signals and Nonlocal Jumps

- **Review**
- Shells
- Signals
- Nonlocal jumps

Exceptions

- An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)
 - Kernel is the memory-resident part of the OS
 - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



Exceptions & Interrupts

■ Interrupts

- Asynchronous
- Cause: Signal from I/O device
- Returns control to next instruction

■ Traps

- Synchronous
- Cause: Intentional exception (system calls)
- Returns control to next instruction

■ Faults

- Synchronous
- Cause: Potentially recoverable error
- Either re-executes faulting (current) instruction or aborts

■ Aborts

- Synchronous
- Cause: Unrecoverable error
- Never returns

Review: syscalls

■ **getpid, getppid**

- Returns pid of calling process (`getpid`), or its parent (`getppid`).

■ **fork**

- Copies current process state.
- “Call once, return twice”

■ **exit**

- Terminates calling process.
- Never returns.

■ **execve**

- Replaces calling process' state with code/state for new program.
- Only returns if there was an error.

■ **wait, waitpid**

- Doesn't return until a child process terminates.

■ **setjmp, longjmp**

ECF Exists at All Levels of a System

■ Exceptions

- Hardware and operating system kernel software

■ Process Context Switch

- Hardware timer and kernel software

■ Signals

- Kernel software and application software

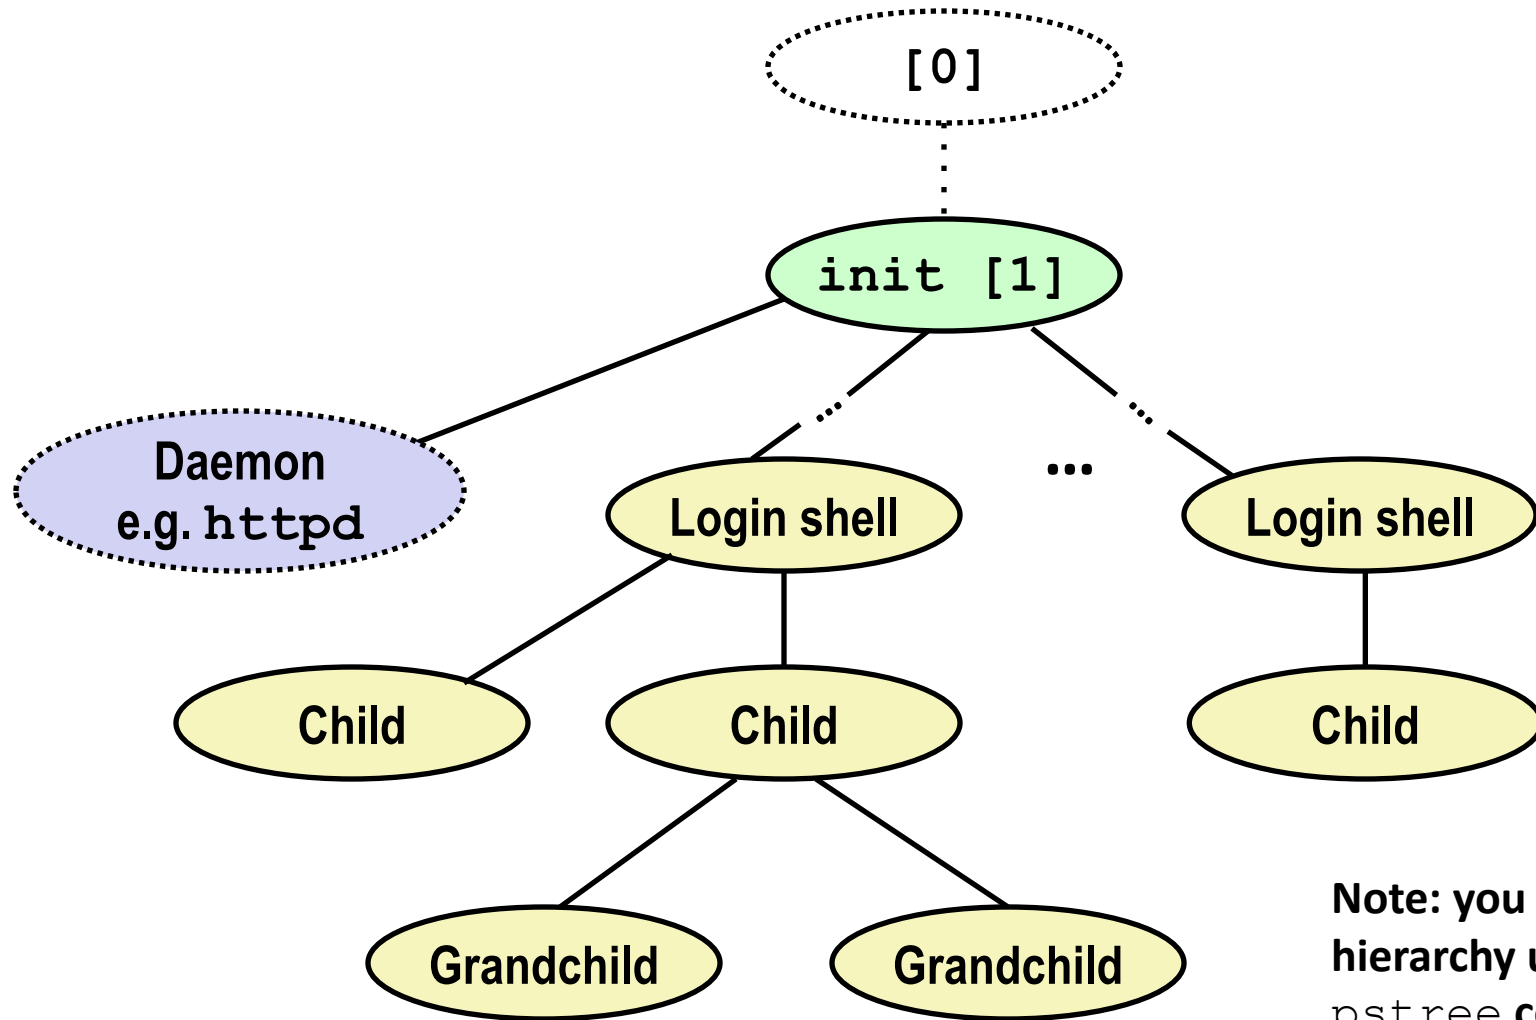
} **Previous**

} **New**

Signals and Nonlocal Jumps

- Review
- **Shells**
- Signals
- Nonlocal jumps

Linux Process Hierarchy



Note: you can view the hierarchy using the Linux `ps tree` command

Shell Programs

- A *shell* is an application program that runs programs on behalf of the user.

- `sh` Original Unix shell (Stephen Bourne, AT&T Bell Labs, 1977)
- `csch/tcsch` BSD Unix C shell
- `bash` “Bourne-Again” Shell (default Linux shell)

```
int main()
{
    char cmdline[MAXLINE]; /* command line */

    while (1) {
        /* read */
        printf("> ");
        fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
}
```

shellex.c

*Execution is a
sequence of
read/evaluate
steps*

Simple Shell eval Function

```
void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = Fork()) == 0) { /* Child runs user job */
            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
}
```

Problem with Simple Shell Example

- **Our example shell correctly waits for and reaps foreground jobs**
- **But what about background jobs?**
 - Will become zombies when they terminate
 - Will never be reaped because shell (typically) will not terminate
 - Will create a memory leak that could run the kernel out of memory

ECF to the Rescue!

■ Solution: Exceptional control flow

- The kernel will interrupt regular processing to alert us when a background process completes
- In Unix, the alert mechanism is called a *signal*