



College of Engineering & Applied Sciences

CSPB 3202

Introduction To Artificial Intelligence

Assignment 4 - Games

UNIVERSITY OF COLORADO


2024

Introduction To Artificial Intelligence - Assignment 4 - Games

1 Assignment 4 - Games	3
Assignment 4 - Games	3
Problem 1.....	4
Problem 2.....	6
Problem 3.....	8
Problem 4.....	10
Problem 5 - Reflex Agent.....	12
Problem 6 - Minimax Agent	15
Problem 7 - Alpha-Beta Pruning	18
Problem 8 - Expectimax Agent	21
Problem 9 - Evaluation Function	23



Assignment 4 - Games



Assignment 4 - Games

I have neither given nor received unauthorized assistance.

Taylor Larrechea


The original assignment can be found [here](#) and [here](#).



Problem 1

Problem Statement


Consider the zero-sum game tree shown below. Triangles that point up, such as at the top node (root), represent choices for the maximizing player; triangles that point down represent choices for the minimizing player. Assuming both players act optimally, fill in the minimax value of each node.




Solution

When we say ‘minimizing player’, we are constituting that that given player must pick the minimum value amongst its children. So the triangles that point down must pick the minimum value from its children. Conversely, the ‘maximizing player must pick the maximum value from its children.

We first start by finding the minimum value of all the downward facing triangles, and then once this is completed, we pick the maximum value from the downward facing triangles for the root node (the upward facing triangle.)





Problem 2

Problem Statement

Which nodes can be pruned from the game tree above through alpha-beta pruning? If no nodes can be pruned, explain why not. Assume the search goes from left to right; when choosing which child to visit first, choose the left-most unvisited child.



Solution

In the leftmost subtree, we evaluate the values 10, 8, and 3. Alpha is updated to 10 and beta is updated to 3 after evaluating these nodes

In the middle subtree, since alpha was previously set to 10, when we evaluate 2, beta is updated to 2. As beta is now less than or equal to alpha, further nodes in this subtree (15 and 7) are pruned.

In the rightmost subtree, alpha remains 10 and beta remains 2. Evaluating the nodes with values 6, 5, and 4 does not result in any updates that would trigger further pruning. Therefore, no values are pruned from this subtree.

The nodes that are pruned are those in the middle subtree with values 15 and 7. This is because, with alpha set to 10, beta is continuously updated to smaller values, leading to the pruning of larger values as they are no longer relevant to the optimal decision-making process.


Nodes **15** and **7** are pruned.



Problem 3


Problem Statement


Again, consider the same zero-sum game tree, except that now, instead of a minimizing player, we have a chance node that will select one of the three values uniformly at random. Fill in the expectimax value of each node. The game tree is redrawn below for your convenience.



Solution

In the context of random chance nodes, we aren't essentially 'randomly' selecting a value from the possible children. Instead, we are taking the average of the children and assigning that as the value for the node. If we perform this choice for each subtree, and then take the maximum value from the children of the root node, we can determine the expectimax value of each node in this scenario. This scenario can be found below.





Problem 4

Problem Statement

Which nodes can be pruned from the game tree above through alpha-beta pruning? If no nodes can be pruned, explain why not.



Solution

Since we are using random chance nodes for this problem, we cannot necessarily prune any nodes from the game tree with alpha-beta pruning. This is because the random chance nodes do not pick values systematically like the minimizing player in the minimax algorithm. Therefore, we cannot prune any nodes from the game tree using alpha-beta pruning in this scenario.

No nodes can be pruned from the game tree.



Problem 5 - Reflex Agent

Problem Statement

Improve the ReflexAgent in `multiAgents.py` to play respectably. The provided reflex agent code provides some helpful examples of methods that query the GameState for information. A capable reflex agent will have to consider both food locations and ghost locations to perform well. Your agent should easily and reliably clear the testClassic layout:

```
1 $ python pacman.py -p ReflexAgent -l testClassic
2
```

Try out your reflex agent on the default `mediumClassic` layout with one ghost or two (and animation off to speed up the display):

```
1 $ python pacman.py --frameTime 0 -p ReflexAgent -k 1
2 $ python pacman.py --frameTime 0 -p ReflexAgent -k 2
3
```

How does your agent fare? It will likely often die with 2 ghosts on the default board, unless your evaluation function is quite good. Try the reciprocal of important values (such as distance to food) rather than just the values themselves. The evaluation function you're writing is evaluating state-action pairs; in later parts of the assignment, you'll be evaluating states.

Command line options that may be useful:

- Default ghosts are random; you can also play for fun with slightly smarter directional ghosts using `-g DirectionalGhost`.
- If the randomness is preventing you from telling whether your agent is improving, you can use `-f` to run with a fixed random seed (same random choices every game).
- You can also play multiple games in a row with `-n`.
- Turn off graphics with `-q` to run lots of games quickly.

Problem 1. Improve the Pacman ReflexAgent behavior as described above. We will run your agent on the `openClassic` layout 10 times. You will receive 0 points if your agent times out, or never wins. You will receive 1 point if your agent wins at least 5 times, or 2 points if your agent wins all 10 games. You will receive an addition 1 point if your agent's average score is greater than 500, or 2 points if it is greater than 1000.

You can try your agent out under these conditions with:

```
1 $ python autograder.py -q q1
2 $ python autograder.py -q q1 --no-graphics # disables graphics
3
```

Don't spend too much time on this question, though, as the meat of the assignment lies ahead!

Solution

```
1 class ReflexAgent(Agent):
2     """
3     A reflex agent chooses an action at each choice point by examining
4     its alternatives via a state evaluation function.
5
6     The code below is provided as a guide. You are welcome to change
7     it in any way you see fit, so long as you don't touch our method
8     headers.
9     """
10
11
12 def getAction(self, gameState):
13     """
14     You do not need to change this method, but you're welcome to.
15
```

```

16     getAction chooses among the best options according to the evaluation function.
17
18     Just like in the previous project, getAction takes a GameState and returns
19     some Directions.X for some X in the set {North, South, West, East, Stop}
20     """
21     # Collect legal moves and successor states
22     legalMoves = gameState.getLegalActions()
23
24     # Choose one of the best actions
25     scores = [self.evaluationFunction(gameState, action) for action in legalMoves]
26     bestScore = max(scores)
27     bestIndices = [index for index in range(len(scores)) if scores[index] == bestScore]
28     chosenIndex = random.choice(bestIndices) # Pick randomly among the best
29
30     "Add more of your code here if you want to"
31
32     return legalMoves[chosenIndex]
33
34 """ evaluationFunction is the function that evaluates the current state of the game and returns a
35 score
36 Input:
37     self - The object pointer
38     currentGameState - The current state of the game
39     action - The action to be taken
40 Algorithm:
41     * Get the current position of the pacman
42     * Get the food, ghosts and capsules from the current state
43     * Calculate the distance of the pacman from the food
44     * Iterate over the total number of ghosts
45     * If the ghost is scared, add the score
46     * If the ghost is not scared, subtract the score
47     * Calculate the distance of the pacman from the capsules
48     * If the capsule distances are not empty
49     * Get the minimum distance of the pacman from the capsules
50     * Calculate the score
51     * Get the remaining food
52     * Subtract the remaining food from the score
53     * Calculate the final score
54     * Add the score of the current state
55     * Add the food score
56     * Add the ghost score
57     * Add the capsule score
58     * Add the food count score
59     * Return the final score
60 Output:
61     finalScore - The score of the current state
62 """
63 def evaluationFunction(self, currentGameState, action):
64     """
65     Design a better evaluation function here.
66
67     The evaluation function takes in the current and proposed successor
68     GameStates (pacman.py) and returns a number, where higher numbers are better.
69
70     The code below extracts some useful information from the state, like the
71     remaining food (newFood) and Pacman position after moving (newPos).
72     newScaredTimes holds the number of moves that each ghost will remain
73     scared because of Pacman having eaten a power pellet.
74
75     Print out these variables to see what you're getting, then combine them
76     to create a masterful evaluation function.
77     """
78     "*** YOUR CODE HERE ***"
79     # Useful information you can extract from a GameState (pacman.py)
80     successorGameState = currentGameState.generatePacmanSuccessor(action)
81     pacmanPos = successorGameState.getPacmanPosition()
82     food = successorGameState.getFood()
83     ghosts = successorGameState.getGhostStates()
84     capsules = currentGameState.getCapsules()
85     # Initialize the variables
86     foodScore = 0
87     ghostScore = 0
88     capsuleScore = 0
89     foodDistances = [manhattanDistance(pacmanPos, foodPos) for foodPos in food.asList()]
90     # If the food distances are not empty
91     if foodDistances:
92         nearestFoodDist = min(foodDistances) # Get the minimum distance of the pacman from the food
93         foodScore = 1.0 / nearestFoodDist # Calculate the score
94     # Iterate over the total number of ghosts
95     for ghost in ghosts:
96         ghostDist = manhattanDistance(pacmanPos, ghost.getPosition()) # Calculate the distance of the
97         pacman from the ghosts
98         # If the ghost is scared, add the score
99         if ghost.scaredTimer > 0:
100             ghostScore += 2.0 / (ghostDist + 1)
101         # Otherwise, subtract the score
102         else:
103             ghostScore -= 1.0 / (ghostDist + 1)
104     # Calculate the distance of the pacman from the capsules
105     capsuleDistances = [manhattanDistance(pacmanPos, capsule) for capsule in capsules]
106     # If the capsule distances are not empty
107     if capsuleDistances:

```

```
106     nearestCapsuleDist = min(capsuleDistances) # Get the minimum distance of the pacman from the
capsules
107     capsuleScore = 1.0 / (nearestCapsuleDist + 1) # Calculate the score
108     # Get the remaining food
109     remainingFood = len(food.asList())
110     # Subtract the remaining food from the score
111     foodCountScore = -remainingFood
112     # Calculate the final score
113     finalScore = (
114         successorGameState.getScore() +
115         foodScore * 10 +
116         ghostScore * 20 +
117         capsuleScore * 10 +
118         foodCountScore * 5
119     )
120     return finalScore
121
```



Problem 6 - Minimax Agent

Problem Statement

Now you will write an adversarial search agent in the provided `MinimaxAgent` class stub in `multiAgents.py`. Your minimax agent should work with any number of ghosts, so you'll have to write an algorithm that is slightly more general than what you've previously seen in lecture. In particular, your minimax tree will have multiple min layers (one for each ghost) for every max layer.

Your code should also expand the game tree to an arbitrary depth. Score the leaves of your minimax tree with the supplied `self.evaluationFunction`, which defaults to `scoreEvaluationFunction`. `MinimaxAgent` extends `MultiAgentSearchAgent`, which gives access to `self.depth` and `self.evaluationFunction`. Make sure your minimax code makes reference to these two variables where appropriate as these variables are populated in response to command line options. It's worth noting that a single search ply is considered to be one Pacman move and all the ghosts' responses, so depth 2 search will involve Pacman and each ghost moving two times.

Problem 2. Using your adversarial search agent in the class above, we will be checking your code to determine whether it explores the correct number of game states. This is the only way reliable way to detect some very subtle bugs in implementations of minimax. As a result, the autograder will be very picky about how many times you call `GameState.generateSuccessor`. If you call it any more or less than necessary, the autograder will mark your code incorrect.

To test and debug your code, run `python autograder.py -q q2` (remember from 1.1 how to run it without graphics, too!). This will show what your algorithm does on a number of small trees, as well as a pacman game. The correct implementation of minimax will lead to Pacman losing the game in some tests. This is not a problem: as it is correct behaviour, it will pass the tests. The evaluation function for the Pacman test in this part is already written (`self.evaluationFunction`). You shouldn't change this function, but recognize that now we're evaluating states rather than actions, as we were for the reflex agent. Look-ahead agents evaluate future states whereas reflex agents evaluate actions from the current state.

- The minimax values of the initial state in the `minimaxClassic` layout are 9, 8, 7, -492 for depths 1, 2, 3 and 4 respectively. Note that your minimax agent will often win (665/1000 games for us) despite the dire prediction of depth 4 minimax (e.g. `python pacman.py -p MinimaxAgent -l minimaxClassic -a depth=4`).
- Pacman is always agent 0, and the agents move in order of increasing agent index
- All states in minimax should be `GameStates`, either passed in to `getAction` or generated via `GameState.generateSuccessor`. In this problem, you will not be abstracting to simplified states.
- On larger boards such as `openClassic` and `mediumClassic` (the default), you'll find Pacman to be good at not dying, but quite bad at winning. He'll often thrash around without making progress. He might even thrash around right next to a dot without eating it because he doesn't know where he'd go after eating that dot. Don't worry if you see this behavior, question 5 will clean up all of these issues.
- When Pacman believes that his death is unavoidable, he will try to end the game as soon as possible because of the constant penalty for living. Sometimes, this is the wrong thing to do with random ghosts, but minimax agents always assume the worst: `python pacman.py -p MinimaxAgent -l trappedClassic -a depth=3`. Make sure you understand why Pacman rushes the closest ghost in this case.

Solution

```
1 class MinimaxAgent(MultiAgentSearchAgent):
2     """
3     Your minimax agent (question 2)
4     """
```

```

5  def getAction(self, gameState):
6      """
7      Returns the minimax action from the current gameState using self.depth
8      and self.evaluationFunction.
9
10     Here are some method calls that might be useful when implementing minimax.
11
12     gameState.getLegalActions(agentIndex):
13         Returns a list of legal actions for an agent
14         agentIndex=0 means Pacman, ghosts are >= 1
15
16     gameState.generateSuccessor(agentIndex, action):
17         Returns the successor game state after an agent takes an action
18
19     gameState.getNumAgents():
20         Returns the total number of agents in the game
21     """
22     """ *** YOUR CODE HERE *** """
23     """ minimax - Calculates the minimax value of the current state
24         Input:
25         agentIndex - The index of the agent
26         depth - The depth of the search
27         gameState - The current state of the game
28         Algorithm:
29         * If the game is won or lost or the depth is reached
30             * Return the evaluation function of the current state
31         * If the agent index is 0
32             * Return the max value
33         * Otherwise
34             * Return the min value
35         Output:
36         The minimax value of the current state
37     """
38     def minimax(agentIndex, depth, gameState):
39         # If the game is won or lost or the depth is reached
40         if gameState.isWin() or gameState.isLose() or depth == self.depth:
41             return self.evaluationFunction(gameState), None
42         # If the agent index is 0
43         if agentIndex == 0:
44             return maxValue(agentIndex, depth, gameState)
45         # Otherwise
46         else:
47             return minValue(agentIndex, depth, gameState)
48     """ maxVal - Calculates the max value of the current state
49         Input:
50         agentIndex - The index of the agent
51         depth - The depth of the search
52         gameState - The current state of the game
53         Algorithm:
54         * Initialize the variables
55         * Iterate over the legal actions of the agent
56             * Get the successor of the current state
57             * Calculate the value of the successor
58             * If the value is greater than the max value
59             * Update the max value
60             * Update the best action
61         * If the depth is 0
62             * Return the max value and the best action
63         * Otherwise
64             * Return the max value
65         Output:
66         The max value of the current state
67     """
68     def maxValue(agentIndex, depth, gameState):
69         # Initialize the variables
70         v = float("-inf")
71         bestAction = None
72         # Iterate over the legal actions of the agent
73         for action in gameState.getLegalActions(agentIndex):
74             # Get the successor of the current state
75             successor = gameState.generateSuccessor(agentIndex, action)
76             # Calculate the value of the successor
77             value, _ = minimax((agentIndex + 1) % gameState.getNumAgents(), depth + ((agentIndex + 1)
78 // gameState.getNumAgents()), successor)
79             # If the value is greater than the max value
80             if value > v:
81                 v = value
82                 bestAction = action
83         # If the depth is 0
84         if depth == 0:
85             return v, bestAction
86         # Otherwise
87         else:
88             return v, None
89     """ minVal - Calculates the min value of the current state
90         Input:
91         agentIndex - The index of the agent
92         depth - The depth of the search
93         gameState - The current state of the game
94         Algorithm:
95         * Initialize the variables
96         * Iterate over the legal actions of the agent

```



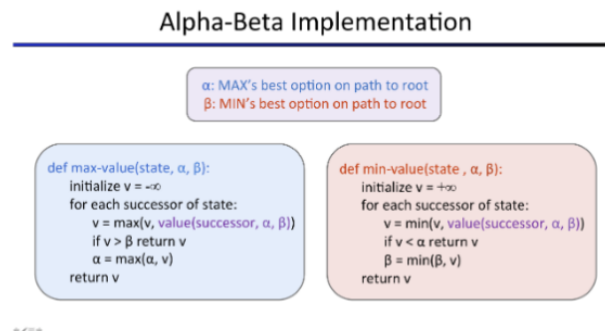
```
96         * Get the successor of the current state
97         * Calculate the value of the successor
98         * If the value is less than the min value
99         * Update the min value
100        * Return the min value
101        Output:
102        The min value of the current state
103    """
104    def minValue(agentIndex, depth, gameState):
105        # Initialize the variables
106        v = float("inf")
107        # Iterate over the legal actions of the agent
108        for action in gameState.getLegalActions(agentIndex):
109            # Get the successor of the current state
110            successor = gameState.generateSuccessor(agentIndex, action)
111            # Calculate the value of the successor
112            value, _ = minimax((agentIndex + 1) % gameState.getNumAgents(), depth + ((agentIndex + 1)
113            // gameState.getNumAgents()), successor)
114            # If the value is less than the min value
115            if value < v:
116                v = value
117        return v, None
118    # Initialize the variables
119    _, action = minimax(0, 0, gameState)
120    # Return the action
121    return action
```



Problem 7 - Alpha-Beta Pruning

Problem Statement

Make a new agent that uses alpha-beta pruning to more efficiently explore the minimax tree, in `AlphaBetaAgent`. Again, your algorithm will be slightly more general than the pseudocode from lecture, so part of the challenge is to extend the alpha-beta pruning logic appropriately to multiple minimizer agents. The pseudo-code below represents the algorithm you should implement for this question.



You should see a speed-up (perhaps depth 3 alpha-beta will run as fast as depth 2 minimax). Ideally, depth 3 on `smallClassic` should run in just a few seconds per move or faster: `python pacman.py -p AlphaBetaAgent -a depth=3 -l smallClassic`.

The `AlphaBetaAgent` minimax values should be identical to the `MinimaxAgent` minimax values, although the actions it selects can vary because of different tie-breaking behavior. Again, the minimax values of the initial state in the `minimaxClassic` layout are 9, 8, 7 and -492 for depths 1, 2, 3 and 4 respectively.

Problem 3. Play the game with your new agent in `AlphaBetaAgent`. We will test your code on a number of small trees, as well as a pacman game. Because we check your code to determine whether it explores the correct number of states, it is important that you perform alpha-beta pruning without reordering children. In other words, successor states should always be processed in the order returned by `GameState.getLegalActions`. Again, do not call `GameState.generateSuccessor` more than necessary.

You must not prune on equality in order to match the set of states explored by our autograder. Indeed, alternatively, but incompatible with our autograder, would be to also allow for pruning on equality and invoke alpha-beta once on each child of the root node, but this will not match the autograder. The correct implementation of alpha-beta pruning will lead to Pacman losing some of the tests. This is not a problem: as it is correct behavior, it will pass the tests

Solution

```

1  class AlphaBetaAgent(MultiAgentSearchAgent):
2      """
3          Your minimax agent with alpha-beta pruning (question 3)
4      """
5      def getAction(self, gameState):
6          """
7          Returns the minimax action using self.depth and self.evaluationFunction
8          """
9          """ *** YOUR CODE HERE *** """
10         """ alphaBeta - Calculates the alpha-beta value of the current state
11             Input:
12                 agentIndex - The index of the agent
13                 depth - The depth of the search
14                 gameState - The current state of the game
15                 alpha - The alpha value
16                 beta - The beta value
17             Algorithm:
18                 * If the game is won or lost or the depth is reached
19                 * Return the evaluation function of the current state
        
```

```

20         * If the agent index is 0
21         * Return the max value
22         * Otherwise
23         * Return the min value
24     Output:
25     The alpha-beta value of the current state
26 """
27 def alphaBeta(agentIndex, depth, gameState, alpha, beta):
28     # If the game is won or lost or the depth is reached
29     if gameState.isWin() or gameState.isLose() or depth == self.depth:
30         return self.evaluationFunction(gameState), None
31     # If the agent index is 0
32     if agentIndex == 0:
33         return maxValue(agentIndex, depth, gameState, alpha, beta)
34     # Otherwise
35     else:
36         return minValue(agentIndex, depth, gameState, alpha, beta)
37 """
38 """ maxValue - Calculates the max value of the current state
39 Input:
40     agentIndex - The index of the agent
41     depth - The depth of the search
42     gameState - The current state of the game
43     alpha - The alpha value
44     beta - The beta value
45 Algorithm:
46     * Initialize the variables
47     * Iterate over the legal actions of the agent
48         * Get the successor of the current state
49         * Calculate the value of the successor
50         * If the value is greater than the max value
51         * Update the max value
52         * Update the best action
53         * If the value is greater than beta
54         * Return the max value and the best action
55         * Update the alpha value
56     * If the depth is 0
57         * Return the max value and the best action
58     * Otherwise
59         * Return the max value
60 Output:
61     The max value of the current state
62 """
63 def maxValue(agentIndex, depth, gameState, alpha, beta):
64     # Initialize the variables
65     v = float("-inf")
66     bestAction = None
67     # Iterate over the legal actions of the agent
68     for action in gameState.getLegalActions(agentIndex):
69         # Get the successor of the current state
70         successor = gameState.generateSuccessor(agentIndex, action)
71         # Calculate the value of the successor
72         value, _ = alphaBeta((agentIndex + 1) % gameState.getNumAgents(), depth + ((agentIndex +
73 1) // gameState.getNumAgents()), successor, alpha, beta)
74         # If the value is greater than the max value
75         if value > v:
76             v = value
77             bestAction = action
78         # If the value is greater than beta
79         if v > beta:
80             return v, bestAction
81         # Update the alpha value
82         alpha = max(alpha, v)
83     # If the depth is 0
84     if depth == 0:
85         return v, bestAction
86     # Otherwise
87     else:
88         return v, None
89 """
90 """ minValue - Calculates the min value of the current state
91 Input:
92     agentIndex - The index of the agent
93     depth - The depth of the search
94     gameState - The current state of the game
95     alpha - The alpha value
96     beta - The beta value
97 Algorithm:
98     * Initialize the variables
99     * Iterate over the legal actions of the agent
100         * Get the successor of the current state
101         * Calculate the value of the successor
102         * If the value is less than the min value
103         * Update the min value
104         * If the value is less than alpha
105         * Return the min value
106         * Update the beta value
107     * Return the min value
108 Output:
109     The min value of the current state
110 """
111 def minValue(agentIndex, depth, gameState, alpha, beta):
112     # Initialize the variables
113     v = float("inf")

```

```
11     bestAction = None
12     # Iterate over the legal actions of the agent
13     for action in gameState.getLegalActions(agentIndex):
14         # Get the successor of the current state
15         successor = gameState.generateSuccessor(agentIndex, action)
16         # Calculate the value of the successor
17         value, _ = alphaBeta((agentIndex + 1) % gameState.getNumAgents(), depth + ((agentIndex +
1) // gameState.getNumAgents()), successor, alpha, beta)
18         # If the value is less than the min value
19         if value < v:
20             v = value
21             bestAction = action
22         # If the value is less than alpha
23         if v < alpha:
24             return v, bestAction
25         # Update the beta value
26         beta = min(beta, v)
27     return v, None
28 # Initialize the variables
29 _, action = alphaBeta(0, 0, gameState, float("-inf"), float("inf"))
30 # Return the action
31 return action
32
```



Problem 8 - Expectimax Agent

Problem Statement

Minimax and alpha-beta are great, but they both assume that you are playing against an adversary who makes optimal decisions. As anyone who has ever won tic-tac-toe can tell you, this is not always the case. In this question you will implement the **ExpectimaxAgent**, which is useful for modeling probabilistic behavior of agents who may make suboptimal choices.

As with the search and constraint satisfaction problems covered so far in this class, the beauty of these algorithms is their general applicability. To expedite your own development, we've supplied some test cases based on generic trees. You can debug your implementation on small the game trees using the command: `python autograder.py -q q4`.

Debugging on these small and manageable test cases is recommended and will help you to find bugs quickly. *Make sure when you compute your averages that you use floats. Integer division in Python 2 (if you're using python 2) truncates, so that $1/2 = 0$, unlike the case with floats where $1.0/2.0 = 0.5$.*

Once your algorithm is working on small trees, you can observe its success in Pacman. Random ghosts are of course not optimal minimax agents, and so modeling them with minimax search may not be appropriate. **ExpectimaxAgent**, however will no longer take the min over all ghost actions, but the expectation according to your agent's model of how the ghosts act. To simplify your code, assume you will only be running against an adversary which chooses amongst their `getLegalActions` uniformly at random. To see how the **ExpectimaxAgent** behaves in Pacman, run:

```
1 python pacman.py -p ExpectimaxAgent -l minimaxClassic -a depth=3
2
```

You should now observe a more cavalier approach in close quarters with ghosts. In particular, if Pacman perceives that he could be trapped but might escape to grab a few more pieces of food, he'll at least try.

Investigate the results of these two scenarios:

```
1 python pacman.py -p AlphaBetaAgent -l trappedClassic -a depth=3 -q -n 10
2 python pacman.py -p ExpectimaxAgent -l trappedClassic -a depth=3 -q -n 10
3
```

Problem 4. We will test your **ExpectimaxAgent** in the two scenarios listed above.

You should find that your **ExpectimaxAgent** wins about half the time, while your **AlphaBetaAgent** always loses. Make sure you understand why the behavior here differs from the minimax case.

Solution

```
1 class ExpectimaxAgent(MultiAgentSearchAgent):
2     """
3     Your expectimax agent (question 4)
4     """
5     def getAction(self, gameState):
6         """
7         Returns the expectimax action using self.depth and self.evaluationFunction
8
9         All ghosts should be modeled as choosing uniformly at random from their
10        legal moves.
11        """
12        """ *** YOUR CODE HERE *** """
13        """ expectimax - Calculates the expectimax value of the current state
14        Input:
15            agentIndex - The index of the agent
16            depth - The depth of the search
17            gameState - The current state of the game
18        Algorithm:
19            * If the game is won or lost or the depth is reached
20              * Return the evaluation function of the current state
21            * If the agent index is 0
22              * Return the max value
23            * Otherwise
24              * Return the exp value
25        Output:
26        The expectimax value of the current state
```

```

27 """
28 def expectimax(agentIndex, depth, gameState):
29     # If the game is won or lost or the depth is reached
30     if gameState.isWin() or gameState.isLose() or depth == self.depth:
31         return self.evaluationFunction(gameState), None
32     # If the agent index is 0
33     if agentIndex == 0:
34         return maxValue(agentIndex, depth, gameState)
35     # Otherwise
36     else:
37         return expValue(agentIndex, depth, gameState)
38 """
39 """ maxValue - Calculates the max value of the current state
40 Input:
41     agentIndex - The index of the agent
42     depth - The depth of the search
43     gameState - The current state of the game
44 Algorithm:
45     * Initialize the variables
46     * Iterate over the legal actions of the agent
47         * Get the successor of the current state
48         * Calculate the value of the successor
49         * If the value is greater than the max value
50         * Update the max value
51         * Update the best action
52     * If the depth is 0
53         * Return the max value and the best action
54     * Otherwise
55         * Return the max value
56 Output:
57     The max value of the current state
58 """
59 def maxValue(agentIndex, depth, gameState):
60     # Initialize the variables
61     v = float("-inf")
62     bestAction = None
63     # Iterate over the legal actions of the agent
64     for action in gameState.getLegalActions(agentIndex):
65         # Get the successor of the current state
66         successor = gameState.generateSuccessor(agentIndex, action)
67         # Calculate the value of the successor
68         value, _ = expectimax((agentIndex + 1) % gameState.getNumAgents(), depth + ((agentIndex +
69 1) // gameState.getNumAgents()), successor)
70         # If the value is greater than the max value
71         if value > v:
72             v = value
73             bestAction = action
74     # If the depth is 0
75     if depth == 0:
76         return v, bestAction
77     # Otherwise
78     else:
79         return v, None
80 """
81 """ expValue - Calculates the exp value of the current state
82 Input:
83     agentIndex - The index of the agent
84     depth - The depth of the search
85     gameState - The current state of the game
86 Algorithm:
87     * Initialize the variables
88     * Iterate over the legal actions of the agent
89         * Get the successor of the current state
90         * Calculate the value of the successor
91         * Add the value to the total value
92     * Return the total value
93 Output:
94     The exp value of the current state
95 """
96 def expValue(agentIndex, depth, gameState):
97     # Initialize the variables
98     v = 0
99     actions = gameState.getLegalActions(agentIndex)
100     p = 1.0 / len(actions)
101     # Iterate over the legal actions of the agent
102     for action in actions:
103         # Get the successor of the current state
104         successor = gameState.generateSuccessor(agentIndex, action)
105         # Calculate the value of the successor
106         value, _ = expectimax((agentIndex + 1) % gameState.getNumAgents(), depth + ((agentIndex +
107 1) // gameState.getNumAgents()), successor)
108         v += p * value
109     return v, None
110 """
111 """
112 # Initialize the variables
113 _, action = expectimax(0, 0, gameState)
114 # Return the action
115 return action

```

Problem 9 - Evaluation Function

Problem Statement

Write a better evaluation function for Pacman in the provided function `betterEvaluationFunction`. The evaluation function should evaluate states, rather than actions like your reflex agent evaluation function did. You may use any tools at your disposal for evaluation. With depth 2 search, your evaluation function should clear the `smallClassic` layout with one random ghost more than half the time and still run at a reasonable rate.

Problem 5. We will test your `ExpectimaxAgent` with `betterEvaluationFunction`. The grader will run your agent on the `smallClassic` layout 10 times. We will give points to your evaluation function in the following way: if you win at least once without timing out the autograder, you receive 1 points. Any agent not satisfying these criteria will receive 0 points. +1 for winning at least 5 times, +2 for winning all 10 times +1 for an average score of at least 500, +2 for an average score of at least 1000 (including scores on lost games) +1 if your games take on average less than 30 seconds on the autograder machine. The additional points for average score and computation time will only be awarded if you win at least 5 times.

As for your reflex agent evaluation function, you may want to use the reciprocal of important values (such as distance to food) rather than the values themselves. One way you might want to write your evaluation function is to use a linear combination of features. That is, compute values for features about the state that you think are important, and then combine those features by multiplying them by different values and adding the results together. You might decide what to multiply each feature by based on how important you think it is.

Solution

```

1  """ betterEvaluationFunction - Calculates the better evaluation function of the current state
2  Input:
3      currentGameState - The current state of the game
4  Algorithm:
5      * Useful information from the game state
6      * Base score is the current game state score
7      * Calculate the reciprocal of the distance to the nearest food
8      * Calculate the reciprocal of the distance to the nearest ghost
9      * Calculate the reciprocal of the distance to the nearest capsule
10     * Calculate remaining food score (negative because more food is worse)
11     * Combine all the scores with appropriate weights
12     * Return the final score
13  Output:
14     finalScore - The score of the current state
15  """
16  def betterEvaluationFunction(currentGameState):
17      """
18          Your extreme ghost-hunting, pellet-nabbing, food-gobbling, unstoppable
19          evaluation function (question 5).
20
21          DESCRIPTION: This evaluation function balances the immediate need to avoid
22          ghosts, collect food, and consume power pellets. It uses the distances to
23          the nearest food, ghost, and power pellet, and the number of remaining food
24          pellets to compute a score. The function prioritizes food collection while
25          avoiding ghosts, and it seeks power pellets to turn ghosts into vulnerable targets.
26      """
27      # Useful information from the game state
28      pacmanPos = currentGameState.getPacmanPosition()
29      food = currentGameState.getFood()
30      ghosts = currentGameState.getGhostStates()
31      capsules = currentGameState.getCapsules()
32      # Base score is the current game state score
33      score = currentGameState.getScore()
34      # Calculate the reciprocal of the distance to the nearest food
35      foodDistances = [manhattanDistance(pacmanPos, foodPos) for foodPos in food.asList()]
36      if foodDistances:
37          nearestFoodDist = min(foodDistances)
38          foodScore = 1.0 / nearestFoodDist
39      else:
40          foodScore = 0

```

```
41 # Calculate the reciprocal of the distance to the nearest ghost
42 ghostScore = 0
43 for ghost in ghosts:
44     ghostDist = manhattanDistance(pacmanPos, ghost.getPosition())
45     if ghost.scaredTimer > 0:
46         ghostScore += 2.0 / (ghostDist + 1) # Incentivize approaching scared ghosts
47     else:
48         ghostScore -= 1.0 / (ghostDist + 1) # Discourage approaching active ghosts
49 # Calculate the reciprocal of the distance to the nearest capsule
50 capsuleDistances = [manhattanDistance(pacmanPos, capsule) for capsule in capsules]
51 if capsuleDistances:
52     nearestCapsuleDist = min(capsuleDistances)
53     capsuleScore = 1.0 / (nearestCapsuleDist + 1)
54 else:
55     capsuleScore = 0
56 # Calculate remaining food score (negative because more food is worse)
57 remainingFood = len(food.asList())
58 foodCountScore = -remainingFood
59 # Combine all the scores with appropriate weights
60 finalScore = (
61     score +
62     foodScore * 10 +
63     ghostScore * 20 +
64     capsuleScore * 10 +
65     foodCountScore * 5
66 )
67 return finalScore
68
```

