

Exam 4 Notes

Memory Management Overview

Memory management is a crucial function of an operating system, enabling efficient allocation, deallocation, and retrieval of memory for processes. It ensures optimal utilization of available memory while minimizing fragmentation and improving access speed. Concepts such as paging, fragmentation, and the use of translation look-aside buffers (TLBs) play significant roles in modern memory systems.

By effectively managing memory, operating systems support multi-programming, reduce access delays, and ensure processes can execute without interference, even in systems with limited resources.

Key Concepts in Memory Management

- **Paging:** Divides memory into fixed-size blocks to eliminate external fragmentation.
- **Fragmentation Types:** Includes internal and external fragmentation.
- **TLB:** Reduces the number of direct memory references.

Fragmentation

Fragmentation occurs when memory is inefficiently utilized. Internal fragmentation happens when allocated memory exceeds process requirements, leaving unused space within a block. External fragmentation arises when free memory exists but is non-contiguous, preventing its use for large processes.

Internal vs. External Fragmentation

- **Internal Fragmentation:** Unused memory within an allocated block.
- **External Fragmentation:** Non-contiguous free blocks prevent allocation for large processes.

Page Faults and TLBs

Page faults occur when a requested memory page is not found in the physical memory and must be retrieved from disk. Translation look-aside buffers (TLBs) improve performance by caching frequently accessed page table entries, minimizing memory lookup times.

Page Faults and TLBs

- **Page Faults:** Occur when memory pages are absent in RAM.
- **TLBs:** Cache page table entries to reduce memory reference delays.

Storage Types and Volatility

Memory technologies vary in speed, volatility, and purpose. Registers and main memory are volatile storage types that lose data when powered off, whereas secondary storage retains data persistently.

Volatile vs. Non-Volatile Storage

- **Volatile Storage:** Includes registers and main memory; data is lost without power.
- **Non-Volatile Storage:** Includes secondary storage; retains data persistently.

Paging and Fragmentation Solutions

Paging solves external fragmentation by dividing memory into fixed-size blocks called pages. This method ensures that free memory blocks can be used efficiently, regardless of their physical location. However, internal fragmentation can still occur if a process does not fully utilize a page.

Paging and Fragmentation

- **Paging Solves External Fragmentation:** Allocates memory in fixed-size pages, making contiguous blocks unnecessary.

- **Internal Fragmentation:** Still possible when processes leave unused space within a page.

Shared Pages and Memory Optimization

Shared pages enable multiple processes to access the same memory regions, such as libraries, without duplication. These pages must be thread-safe and reentrant to prevent data corruption. While sharing pages optimizes memory usage, it does not directly improve memory access times.

Shared Pages

- **Shared Code:** Must be reentrant and thread-safe for safe use by multiple processes.
- **Library Sharing:** Allows memory efficiency by avoiding redundant code duplication.

Inverted Page Tables

Inverted page tables address the problem of sparse page tables by maintaining a single entry for each physical memory page. This reduces memory overhead for large virtual address spaces and ensures efficient memory mapping.

Inverted Page Tables

- **Sparse Page Table Solution:** Maps physical pages directly, reducing table size.
- **Efficiency:** Optimizes memory mapping for large address spaces.

Summary of Key Concepts

- **Fragmentation:** Internal fragmentation wastes memory within blocks; external fragmentation arises from non-contiguous free blocks.
- **Paging:** Resolves external fragmentation by using fixed-size memory blocks.
- **TLBs and Page Faults:** Improve memory performance by caching page table entries and handling absent pages.
- **Volatile vs. Non-Volatile Memory:** Differentiates storage based on data retention without power.
- **Inverted Page Tables:** Optimize memory for sparse page table scenarios.

Effective memory management ensures optimal system performance, minimizing fragmentation and improving resource utilization through paging, TLBs, and shared pages.

Virtual Memory Overview

Virtual memory is a memory management technique that enables systems to execute processes larger than the available physical memory. By dividing programs into pages and mapping them to frames in physical memory, virtual memory allows for efficient use of limited resources and provides the illusion of a large contiguous address space. Key concepts such as page faults, page replacement algorithms, and working sets are crucial to understanding how virtual memory operates.

Virtual memory reduces the need for complete program loading into physical memory, enabling multitasking and enhancing overall system performance.

Key Concepts in Virtual Memory

- **Page Fault:** Occurs when a requested page is not in physical memory.
- **Page Replacement Algorithms:** Determine which pages to evict when physical memory is full.
- **Working Set:** The set of pages a process accesses within a given timeframe.

Page Replacement Algorithms

Page replacement algorithms manage which pages are removed from physical memory when a page fault occurs. Common algorithms include FIFO (First-In, First-Out), LRU (Least Recently Used), and OPTIMAL. Each algorithm has unique trade-offs in terms of complexity and efficiency.

Comparison of Page Replacement Algorithms

- **FIFO**: Evicts the oldest page in memory; simple but prone to Belady's anomaly.
- **LRU**: Evicts the least recently accessed page; avoids Belady's anomaly but requires tracking usage history.
- **OPTIMAL**: Evicts the page that will not be used for the longest time in the future; ideal but impractical to implement.

Belady's Anomaly

Belady's anomaly describes a counterintuitive scenario where increasing the number of available page frames leads to more page faults. This anomaly occurs in algorithms like FIFO but not in stack-based algorithms like LRU.

Belady's Anomaly

- **Definition**: Increasing the number of frames increases page faults.
- **Impact**: Highlights inefficiencies in certain algorithms like FIFO.

Global and Local Page Replacement

Global page replacement algorithms select a page to evict from any process, while local algorithms restrict eviction to pages belonging to the current process. Global replacement can improve system-wide performance but may lead to issues like process starvation.

Global vs. Local Page Replacement

- **Global Replacement**: Evicts any page in memory; improves overall utilization.
- **Local Replacement**: Limits eviction to pages of the requesting process; ensures fairness.

Working Sets and Thrashing

The working set model identifies the set of pages a process accesses during a timeframe. Proper management of working sets reduces page faults and avoids thrashing, where excessive paging dominates CPU cycles, reducing system efficiency.

Working Sets and Thrashing

- **Working Set**: Pages accessed within a defined timeframe; determines resource needs.
- **Thrashing**: Excessive paging due to insufficient memory allocation.

Page Fault Metrics for Example Sequences

Calculations under different algorithms reveal their efficiency. For a sequence of page references {1, 2, 3, 1, 4, 1, 2, 4, 2, 4, 3} with 3 frames:

- **FIFO**: 7 page faults.
- **OPTIMAL**: 5 page faults.
- **LRU**: 6 page faults.

Performance Analysis

- **FIFO**: High fault rate; prone to Belady's anomaly.
- **OPTIMAL**: Ideal fault rate; requires future knowledge.
- **LRU**: Balances complexity and efficiency; avoids Belady's anomaly.

Summary of Key Concepts

- **Virtual Memory Benefits:** Supports larger processes than physical memory.
- **Page Replacement:** FIFO, LRU, and OPTIMAL manage memory efficiently.
- **Belady's Anomaly:** Demonstrates inefficiencies in specific algorithms.
- **Working Sets:** Aid in memory allocation and reduce thrashing risks.

Virtual memory is a cornerstone of modern operating systems, enabling multitasking and efficient memory usage. Understanding page replacement, working sets, and anomaly detection helps optimize performance in complex environments.

Protection and Security Overview

Protection and security are fundamental aspects of operating systems, ensuring controlled access to resources and safeguarding data from unauthorized actions. Protection mechanisms define how resources are accessed and by whom, while security mechanisms focus on detecting and preventing unauthorized access or data breaches. Together, these mechanisms maintain system integrity, confidentiality, and availability.

By establishing clear access control policies and leveraging encryption techniques, operating systems provide a robust framework for secure and reliable operation.

Key Protection Concepts

- **Protection:** Mechanisms controlling access to resources by users or programs.
- **Security:** Activities aimed at detecting and deterring unauthorized intrusions.

Protection Domains and Access Control

A protection domain defines a set of objects and the operations that can be performed on them. This concept is central to access control, which determines the permissions assigned to users or processes. The access matrix is a model used to represent the relationship between domains, objects, and permissible actions.

Protection Domains

- **Definition:** A set of objects and operations allowed on them.
- **Access Matrix:** Represents the mapping of domains to objects and their permissible operations.

Breaches of Security

Breaches in security occur when unauthorized actions are performed on system resources or data. Common types of breaches include theft of service, unauthorized data modification, and denial of service (DoS) attacks. A breach of integrity refers specifically to unauthorized modifications of data, compromising its correctness or trustworthiness.

Types of Security Breaches

- **Theft of Service:** Unauthorized use of system resources.
- **Breach of Integrity:** Unauthorized modification of data.
- **Denial of Service (DoS):** Preventing legitimate use of a system.

Encryption Techniques

Encryption is critical for ensuring the confidentiality of data, especially during transmission. Symmetric encryption uses the same key for both encryption and decryption, while asymmetric encryption employs a pair of keys—one for encryption and another for decryption. Asymmetric encryption extends symmetric methods by enabling secure key distribution and authentication.

Symmetric vs. Asymmetric Encryption

- **Symmetric Encryption:** Uses a single key for both encryption and decryption.
- **Asymmetric Encryption:** Employs a pair of keys—one for encryption and another for decryption.

Summary of Key Concepts

- **Protection Domains:** Define resource access control and permissible operations.
- **Security Breaches:** Include theft of service, integrity breaches, and denial of service.
- **Encryption Methods:** Symmetric encryption is efficient but requires secure key distribution, while asymmetric encryption provides better security through key pairs.

Protection and security mechanisms are essential for maintaining a system's reliability, confidentiality, and integrity. Understanding these concepts ensures that resources are used appropriately and safeguarded against threats.

