



Department of Computer  
Science

UNIVERSITY OF COLORADO BOULDER



# Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox  
University of Colorado Boulder

These slides adapted from materials provided by the textbook authors.



University of Colorado Boulder

# Deadlock Avoidance

## *Deadlock Prevention*

*prevent at least 1 of  
necessary conditions*

*Mutual exclusion*

*Hold and wait*

*No preemption*

*Circular dependency*

## *Deadlock Avoidance*

*Can we analyze the system  
and detect deadlock?*

*Need to know resource  
usage required by each  
process (and track usage)*

*Let the OS deal with  
deadlock at runtime*

## *Deadlock Avoidance*

*will need to detect deadlock  
within a set of processes*

- *knowledge of the maximum number of resources used by each process*
- *keep track of resource usage*
- *make sure there is always a sequence of processes that could run to completion and NOT exceed the available resources even if process demands their maximum*


## *Safe State*

*Intuition: IF there is a way to satisfy the resource requirements for all tasks, the system cannot deadlock!*

- *find a sequence of process completion that will never exceed the resources available*
- *As each process completes the resources currently allocated would be available to other processes*

## Safe State

- *Most of the system states will be safe*
- *Unsafe states do not necessarily lead to deadlock*
- *Some unsafe states can lead to deadlock*





## Example 1 (initial state)

processes	max needs	allocated (total=12)
P0	10	5
P1	4	2
P2	9	2

sequence <P1, P0, P2> is safe

*Available*

- P1 requests its max  
(has 2, so needs 2 more),  
now holds 4 3
- then P1 releases all 4 5
- P0 requests its max  
(has 5, so needs 5 more),  
now holds 10 0
- then P0 releases all 10 10
- P2 requests its max  
(has 2, so needs 7 more),  
holds 9 3
- then P0 releases all 9 12

## Example 1 (initial state)

processes	max needs	allocated (total=12)
P0	10	5
P1	4	2
P2	9	3

Is the state still safe?

*Available*

- P1 requests its max  
(has 2, so needs 2 more),  
now holds 4
- then P1 releases all 4
- P0 requests its max  
(has 5, so needs 5 more)
- DEADLOCK !!!

## *Dijkstra's Banker's Algorithm*

- *Generalizes deadlock avoidance to multiple resources*
  - *Determines whether the system is in a safe state*
- *Before granting a request, run Banker's Algorithm pretending as if request was granted*
  - *Does the worst-case analysis find such a hypothetical system is in a safe state?*
  - *If so, grant request.*
  - *If not, delay requestor, and wait for more resources to be freed.*

*Available*

*Allocated*

*Max*

*Needed*

	$R_1$	$R_2$	$R_3$	...	$R_M$
			Col $j$		
$P_1$					
...					
$P_N$					
$P_1$					
...					
$P_N$					
$P_1$					
...					
$P_N$					

$Alloc_i$  is shorthand  
for row  $i$  of matrix,  
i.e. the resources  
allocated  
to process  $i$

## *Terminology*

$$V1 = \begin{bmatrix} 1 \\ 7 \\ 3 \\ 2 \end{bmatrix} \quad V2 = \begin{bmatrix} 0 \\ 3 \\ 2 \\ 1 \end{bmatrix}$$

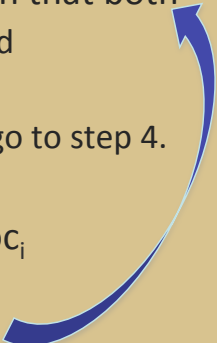
$$V3 = \begin{bmatrix} 0 \\ 10 \\ 2 \\ 1 \end{bmatrix}$$

$$V2 \leq V1$$

But,

$$V3 \not\leq V1$$

# Banker's Algorithm

1. Let Work and Finish be vectors length m and n respectively.
    - Initialize Work = Available
    - Finish[i]=false for  $i=0, \dots, n-1$
  2. Find a process i such that both
    - Finish[i]==false, and
    - $Need_i \leq Work$
    - If no such i exists, go to step 4.
  3.  $Work = Work + Alloc_i$ 
    - Finish[i] = true
    - Go to step 2.
  4. If Finish[i]==true for all i, then the system is in a safe state
- 

Available

A	B	C
3	3	2

	Alloc[i,j]			Max[i,j]		
	A	B	C	A	B	C
P0	0	1	0	7	5	3
P1	2	0	0	3	2	2
P2	3	0	2	9	0	2
P3	2	1	1	2	2	2
P4	0	0	2	4	3	3

Work = Available

while (find unfinished  $P_i$

where  $Need_i \leq Work$ )

Work = Work +  $Alloc_i$

Finish[i] = true

If Finish[i]==true for all i,

then the system is in a safe state

Is  $\langle P1, P3, P4, P2, P0 \rangle$  a safe sequence?

# Complexity

*M resources*

*N processes*

*Algorithm must look at  
each process and consider  
all process sequences*

*Each search must look at  
all resources*

$O(N^2 * M)$

Algorithm:


```
Work = Available
while (find unfinished  $P_i$ 
      where  $Need_i \leq Work$ )
    Work = Work +  $Alloc_i$ 
    Finish[i] = true
```

If Finish[i]==true for all i,  
then the system is in a safe  
state



# Banker's Algorithm

- Adapt to handle requests from processes  
check to see if request would be safe
- Similar to resource allocation graph checking  
See if new edge would generate a cycle



# Solutions to Handling Deadlocks

1. Prevention by OS
  - provide methods to guarantee that at least 1 of the 4 necessary conditions for deadlock does not hold
2. Avoidance by OS
  - the OS is given advanced information about process requests for various resources, use the Banker's algorithm
  - Use resource allocation graphs
3. Detection and Recovery by OS
  - Analyze existing system resource allocation, and see if there is a sequence of releases that satisfies every process' needs.
4. Application-level solutions (OS Ignores and Pretends)
  - it's up to the application programmer to implement mechanisms that prevent, avoid, detect and deal with application-level deadlock

# Deadlock Detection

- When/how often should the detection algorithm run?
  - Depends on how often deadlock is likely to occur
  - Depends on how quickly deadlock grows after it occurs, i.e. how many processes get pulled into deadlock and on what time scale
  - Could check at each resource request – this is costly
  - Could check periodically – but what is a good time interval?
  - Could check if CPU utilization suddenly drops – this might be an indication that there's deadlock, and processes are no longer executing, but what's a good threshold?
  - Could check if resource utilization exceeds some threshold, but what's a good threshold?

# Deadlock Recovery

- After OS has detected which processes are deadlocked:
  - Terminate all processes - draconian
  - Terminate one process at a time until the deadlock cycle is eliminated
    - Check if there is still deadlock after each process is terminated. If not, then stop.
  - Preempt some processes – temporarily take away a resource from current owner and give it to another process but don't terminate process
    - e.g. give access to a laser printer – this is risky if you're in middle of printing a documents
  - Rollback some processes to a checkpoint – assuming that processes have saved their state at some checkpoint



Department of Computer  
Science

UNIVERSITY OF COLORADO BOULDER



# Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox  
University of Colorado Boulder



These slides adapted from materials provided by the textbook authors.