


Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox
University of Colorado Boulder





Department of Computer Science
UNIVERSITY OF COLORADO **BOULDER**



Design and Analysis of Operating Systems CSCI 3753


Distributed Systems

Dr. David Knox
University of
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.

Distributed Systems


Multicore vs Distributed



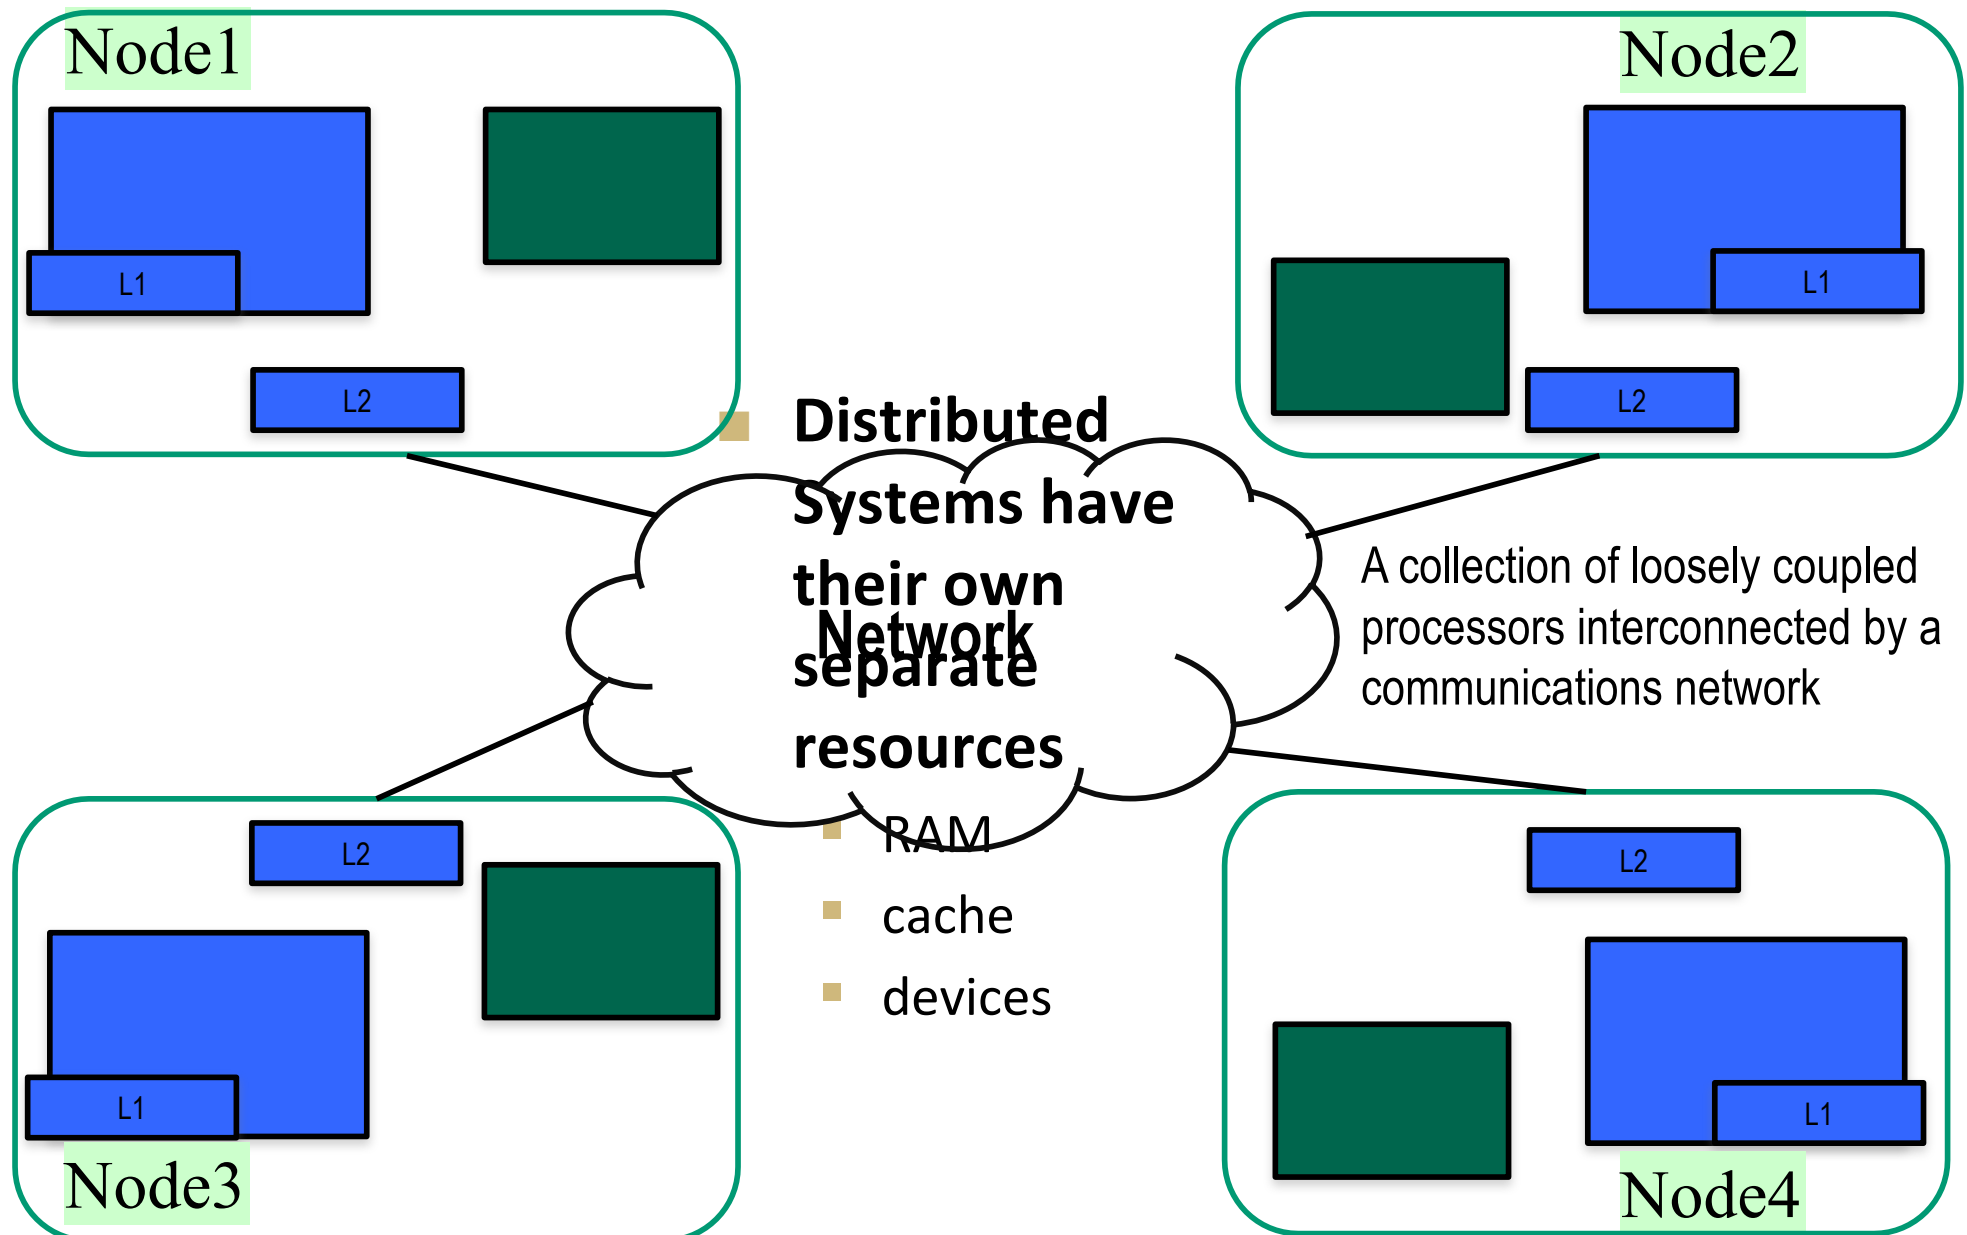
- Multi-core systems have access to shared resources

- RAM
- cache
- devices

Multicore vs Distributed




Multicore vs Distributed



Distributed System


- A collection of loosely coupled processors interconnected by a communications network
- Processors variously called *nodes*, *computers*, *machines*, *hosts*
 - Network is used to communicate between machines
 - Generally a **server** has a resource a **client** node at a different site wants to use



Distributed System

■ Reasons for a Distributed System

- Resource sharing
- Computational speed up (load balancing)
- Reliability (thru redundancy)
- Remote connections thru communications



Network-Operating Systems

- **Users are aware of multiplicity of machines**
- **Access to resources of various machines is done explicitly by:**
 - Remote logging into the appropriate remote machine (telnet, ssh)
 - Remote Desktop (Microsoft Windows)
 - Transferring data from remote machines to local machines, via the File Transfer Protocol (FTP) mechanism
- **Users must change paradigms**
 - establish a **session**
 - give network-based commands
 - More difficult for users

Distributed-Operating Systems

- **Users not aware of multiplicity of machines**
 - Access to remote resources similar to access to local resources
- **Data Migration**
 - transfer data by transferring entire file
 - or transferring only those portions of the file necessary for the immediate task
- **Computation Migration**
 - Transfer the computation, rather than the data, across the system
 - Via remote procedure calls (RPCs)
 - or via messaging system

Distributed-Operating Systems (Cont.)

■ Process Migration

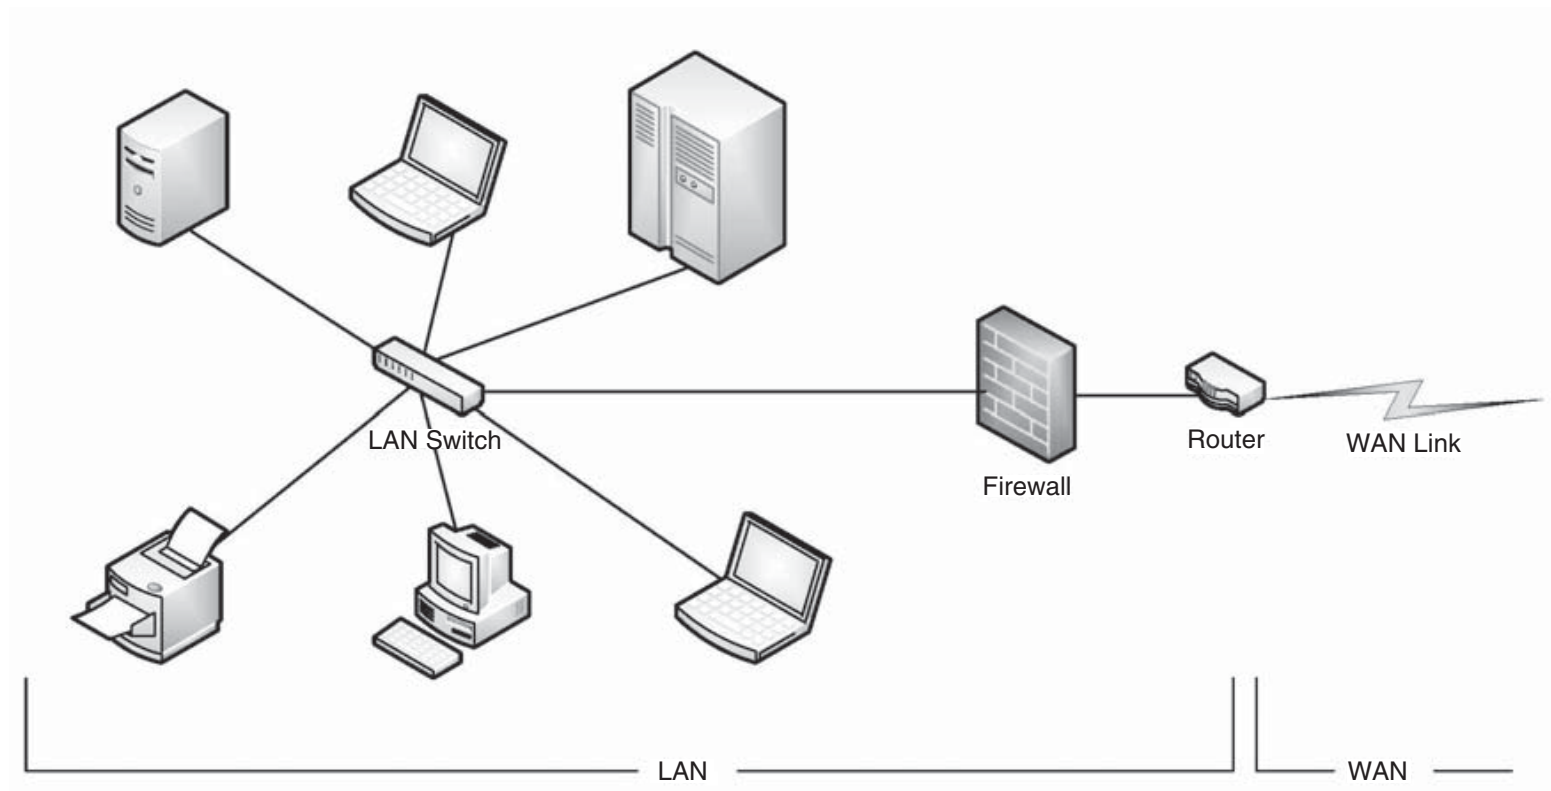
- **Execute** an entire process, or parts of it, at different sites
- **Load balancing** – distribute processes across network to even the workload
- **Computation speedup** – subprocesses can run concurrently on different sites
- **Hardware preference** – process execution may require specialized processor
- **Software preference** – required software may be available at only a particular site
- **Data access** – run process remotely, rather than transfer all data locally

Network Structure

■ Local-Area Network (LAN)

- Designed to cover small geographical area
- Multiple topologies like star or ring
- Speeds from 1Mb per second (Appletalk, bluetooth)
40 Gbps for Ethernet over twisted pair copper or optical fibre
- Consists of multiple computers (mainframes through mobile devices), peripherals (printers, storage arrays), routers (specialized network communication processors) providing access to other networks
- Ethernet most common way to construct LANs
 - Multiaccess bus-based
 - Defined by standard IEEE 802.3
- Wireless spectrum (**WiFi**) increasingly used for networking
 - I.e. IEEE 802.11g standard implemented at 54 Mbps


Local-area Network



Network Types (Cont.)

- **Wide-Area Network (WAN) – links geographically separated sites**
 - Point-to-point connections over long-haul lines (often leased from a phone company)
 - Implemented via **connection processors** known as **routers**
 - Internet WAN enables hosts world wide to communicate
 - Hosts differ in all dimensions but WAN allows communications
 - Speeds
 - T1 link is 1.544 Megabits per second
 - T3 is 28 x T1s = 45 Mbps
 - OC-12 is 622 Mbps
 - WANs and LANs interconnect, similar to cell phone network:
 - Cell phones use radio waves to cell towers
 - Towers connect to other towers and hubs

Communication Processors in a Wide-Area Network



Communication Structure

Design of a communication network must address four basic issues:

- **Naming and name resolution**

- How do two processes locate each other to communicate?

- **Routing strategies**


- How are messages sent through the network?

- **Connection strategies**

- How do two processes send a sequence of messages?

- **Contention**

- The network is a shared resource, so how do we resolve conflicting demands for its use?



Naming and Name Resolution

- Name systems in the network
- Address messages with the process-id
- Identify processes on remote systems by
 <host-name, identifier> pair
- Domain name system (DNS)
 - specifies the naming structure of the hosts, as well as name to address **resolution** (Internet)

Routing Strategies

- **Fixed routing - A path from A to B is specified in advance; path changes only if a hardware failure disables it**
 - Since the shortest path is usually chosen, communication costs are minimized
 - Fixed routing cannot adapt to load changes
 - Ensures that messages will be delivered in the order in which they were sent
- **Virtual routing- A path from A to B is fixed for the duration of one session. Different sessions involving messages from A to B may have different paths**
 - Partial remedy to adapting to load changes
 - Ensures that messages will be delivered in the order in which they were sent

Routing Strategies (Cont.)

■ Dynamic routing

- Path used to send a message from site *A* to site *B* is chosen only when a message is sent
- Usually a site sends a message to another site on the link least used at that particular time
- Adapts to load changes by avoiding routing messages on heavily used path
- Messages may arrive out of order
 - This problem can be remedied by appending a sequence number to each message
- Most complex method to set up

■ Tradeoffs mean all methods are used

- UNIX provides ability to mix fixed and dynamic
- Hosts may have fixed routes and **gateways** connecting networks together may have dynamic routes

Routing Strategies (Cont.)

- Router is communications processor responsible for routing messages
- Must have at least 2 network connections
- Maybe special purpose or just function running on host
- Checks its tables to determine where destination host is, where to send messages
 - Static routing – table only changed manually
 - Dynamic routing – table changed via **routing protocol**

Routing Strategies (Cont.)

- **More recently, routing managed by intelligent software more intelligently than routing protocols**
- **Messages vary in length**
 - simplified design breaks them into **packets** (or **frames**, or **datagrams**)
- **Connectionless message is just one packet**
 - Otherwise need a connection to get a multi-packet message from source to destination



Department of Computer Science
UNIVERSITY OF COLORADO **BOULDER**



Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox
University of
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.