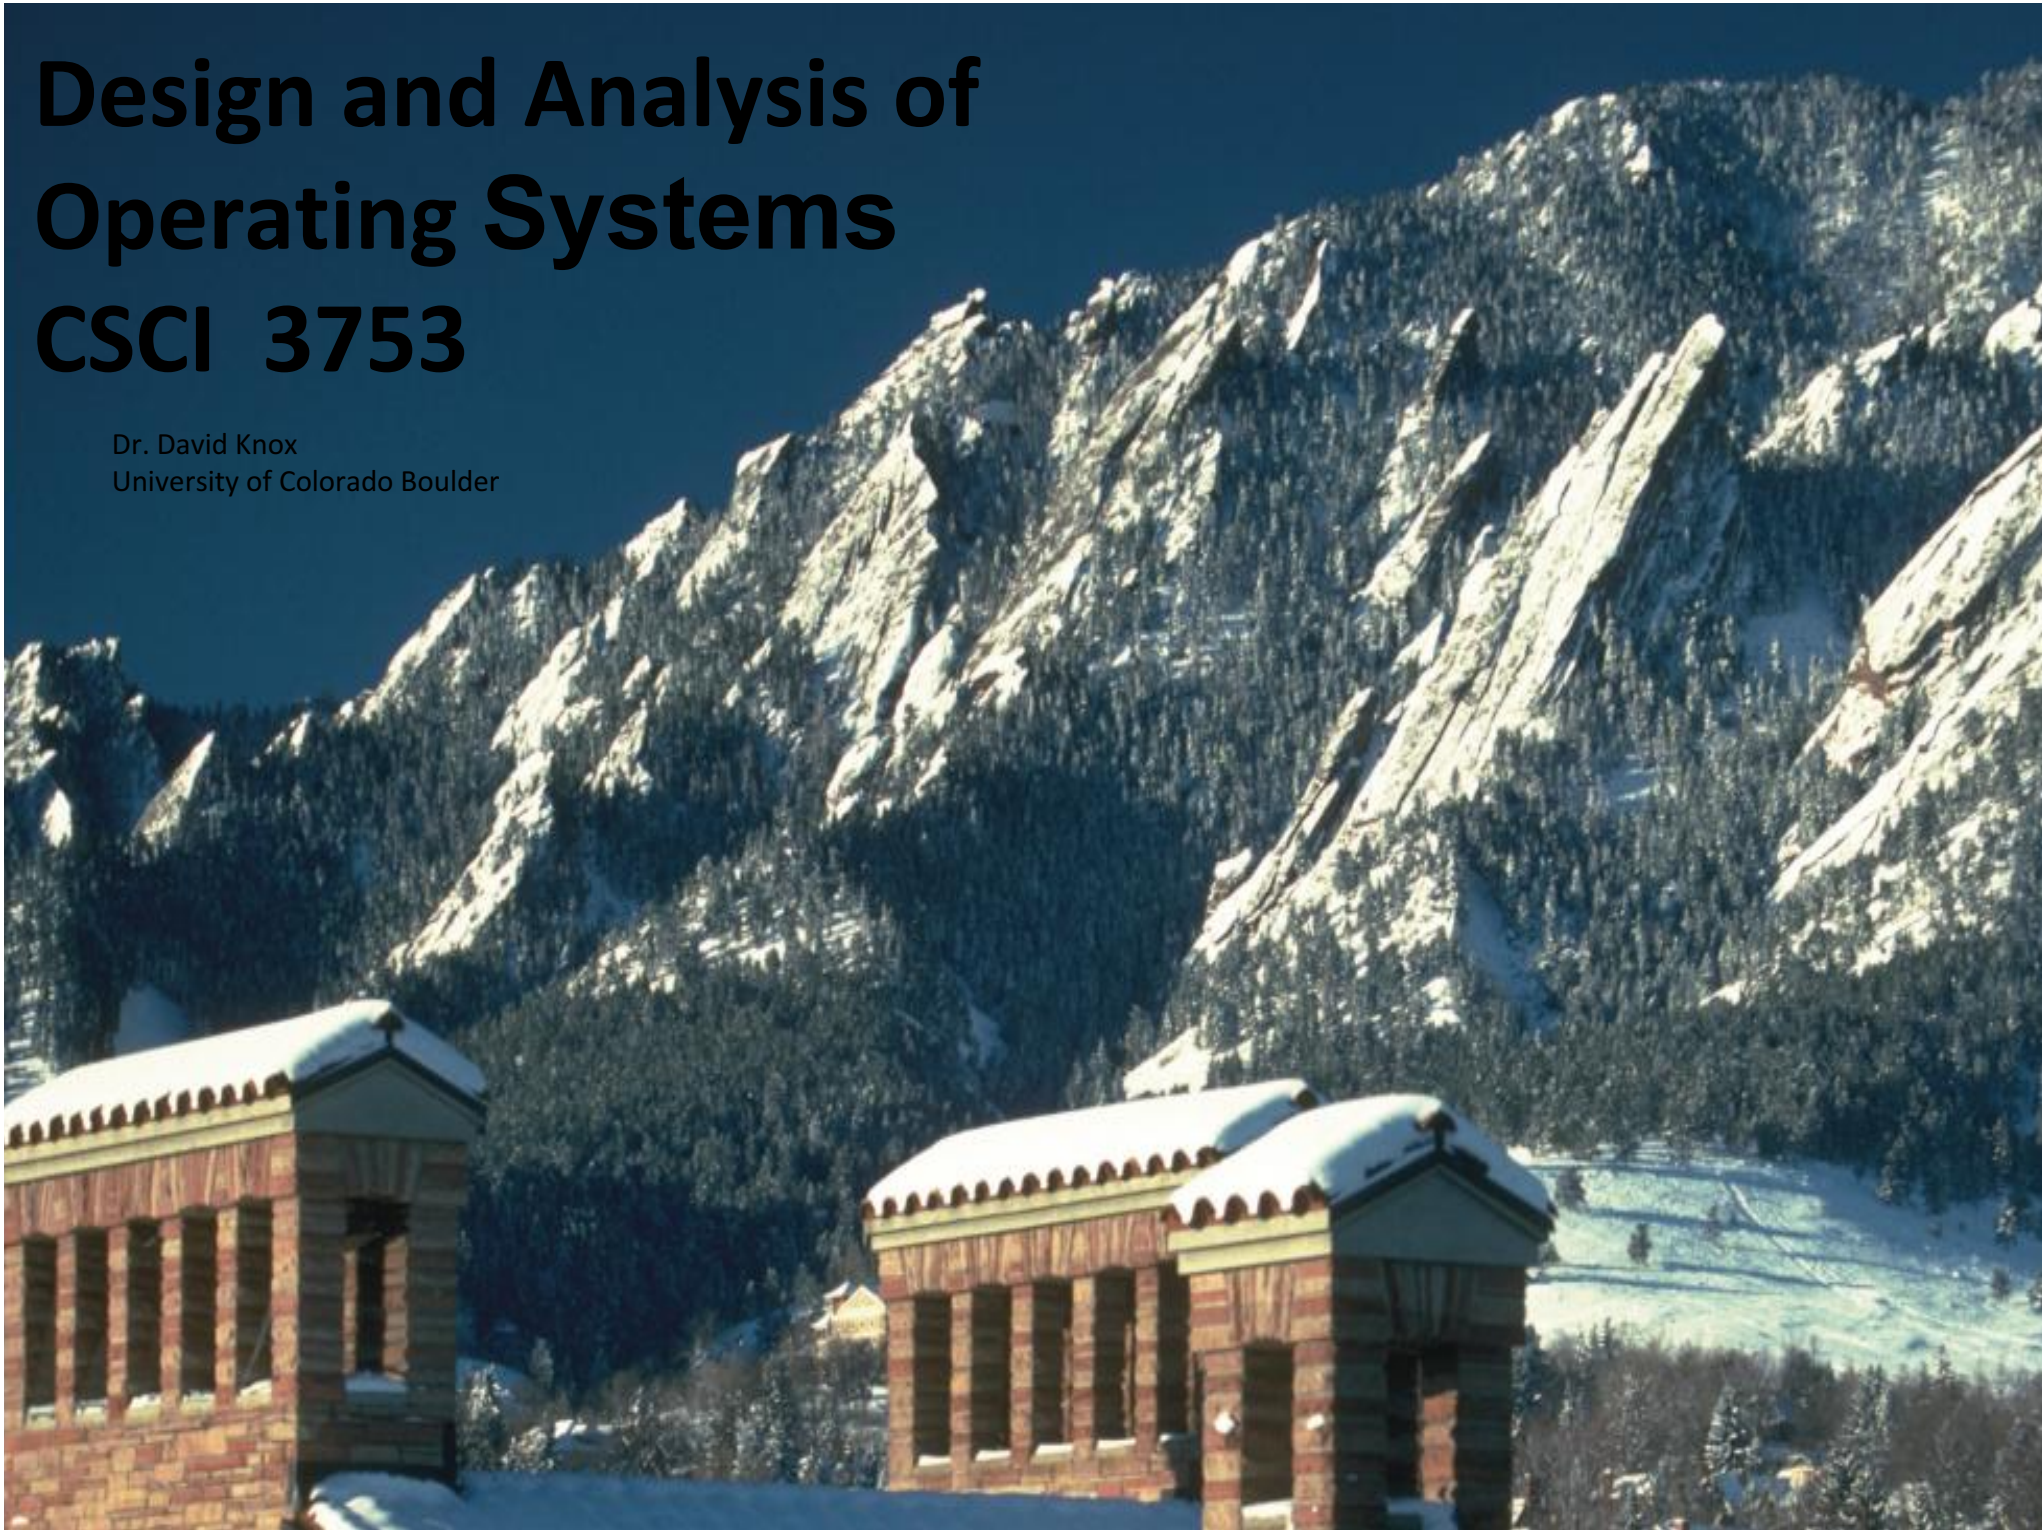


# Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox  
University of Colorado Boulder





Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Design and Analysis of Operating Systems CSCI 3753

## File Allocation

Dr. David Knox  
University of  
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.

# File Allocation

# File Allocation

## ■ Approaches:

### 1. Contiguous file allocation

- a file is laid out contiguously, i.e. if a file is  $n$  blocks long, then a starting address  $b$  is selected and the file is allocated blocks  $b, b+1, b+2, \dots, b+n-1$

### 2. Linked Allocation

- each file is a linked list of disk blocks

### 3. File Allocation Table (FAT) is an important variation of linked lists


- Don't embed the pointers of the linked list with the file data blocks themselves
- Instead, separate the pointers out and put them in a special table – the file allocation table (FAT)
- The FAT is located at a section of disk at the beginning of a volume

### 4. Indexed Allocation

- collect all pointers into a list or table called an *index block*
- the index  $j$  into the list or index block retrieves a pointer to the  $j$ 'th block on disk

# Approach #1: Contiguous File Allocation

File headers		
file	start	length
foo1	0	2
foo2	4	4



- **Advantage: fast performance (low seek times because the blocks are all allocated near each other on disk)**



# Approach #1: Contiguous File Allocation

## ■ Disadvantages:

- **Problem 1:** external fragmentation (observed same problem fitting into RAM)
  - same solutions apply: first fit, best fit, etc.
  - can also compact memory/defragment disk
  - schedule to be performed in the background late at night, etc.
- **Problem 2:** May not know size of file in advance
  - allocate a larger size than estimated
  - if file exceeds allocation, have to copy file to a larger free “hole”
- **Problem 3:** Over-allocation of a “slow growth” file
  - A file may eventually need 1 million bytes of space
  - But initially, the file doesn’t need much  
(may be growing at a very slow rate, e.g. 1 byte/sec)
  - For much of the lifetime of the file, allocating 1 MB wastes allocation

# Contiguous File Allocation

## ■ Disadvantages:

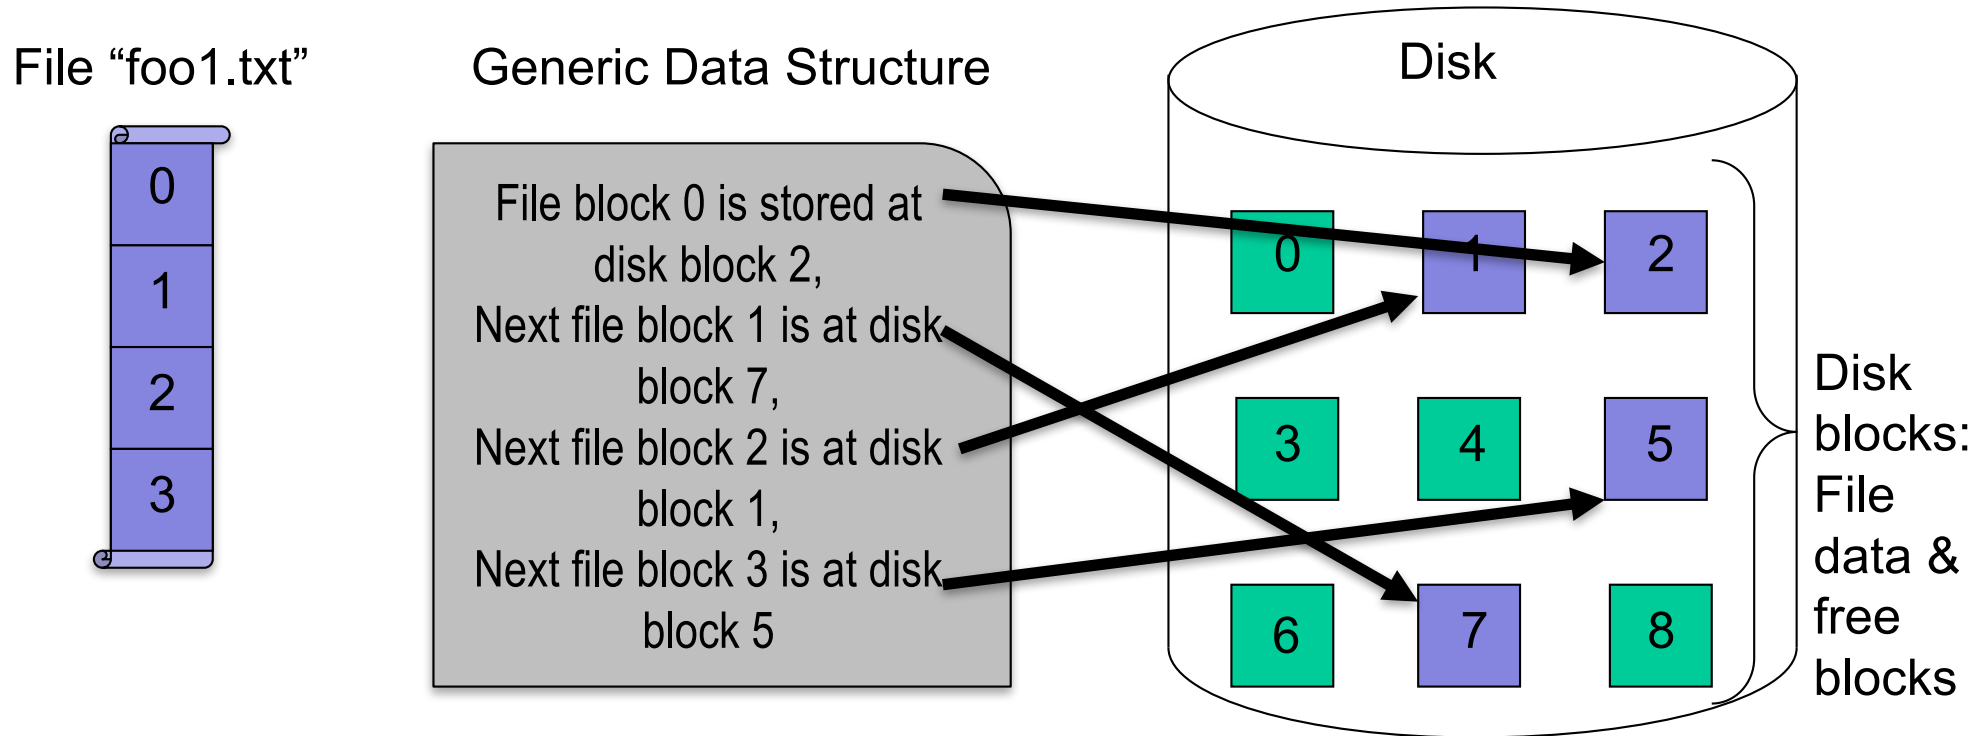
- Problem 3: Over-allocation of a “slow growth” file
  - A file may eventually need 1 million bytes of space
  - But initially, the file doesn’t need much, and it may be growing at a very slow rate, e.g. 1 byte/sec
  - So for much of the lifetime of the file, allocating 1 MB wastes allocation
  - This is a “slow growth” problem.

# General File Allocation

- Page table solved external fragmentation problem for process allocation
- Apply a similar concept to file allocation
  - Divide disk into fixed-sized blocks, just as main memory was divided into fixed-sized physical frames
  - Allow a file's data blocks to be spread across any collection of disk blocks, not necessarily contiguous
  - *Need a data structure to keep track of what block of a file is stored on which block in disk*



# General File Allocation



- **Generic data structure can be:**
  - A Linked list and variants
  - Indexed allocation (somewhat resembles a page table) and variants

# General File Allocation

- **General approach to file allocation solves:**
  - External fragmentation problem
  - Problem of not knowing file size in advance and having to over-estimate
    - Allocate exactly number of disk blocks needed
    - As more disk blocks are needed, easy to allocate exactly the additional number of disk blocks needed from pool of free/unallocated disk blocks
    - and these disk blocks can be anywhere on disk, not necessarily contiguous
  - Slow growth problem: only allocate exactly as many blocks as a slow growth file needs

# General File Allocation

## ■ Examples:


- UNIX FS (UFS, = Berkeley FFS) uses 8 KB blocks.
- Linux' file system ext2fs uses default 1 KB blocks (though 2 and 4 KB supported (and much larger))

# Approach #2: Linked File Allocation

## ■ Linked Allocation

- each file is a linked list of disk blocks
- to add to a file, just modify the linked list either in the middle or at the tail, depending on where you wish to add a block
- to read from a file, traverse linked list until reaching the desired data block


File “foo1.txt”



# Linked File Allocation

## ■ Linked Allocation

- each file is a linked list of disk blocks



# Linked File Allocation

## ■ Advantages:

- solves problems of contiguous allocation
  - no external fragmentation
  - don't need to know size of a file a priori
- Minimal bookkeeping overhead in file header – just a pointer to start of file on disk
  - Compromise is that all the pointer overhead is stored in each disk block
- Good for sequential read/write data access
- Easy to insert data into middle of linked list

# Linked File Allocation

## ■ Problems:

- performance of random (direct) data access is extremely slow for reads/writes
  - because you have to traverse the linked list until indexing into the correct disk block
- Space is required for pointers on disk in every disk block
- reliability is fragile
  - if one pointer is in error or corrupted, then lose the rest of the file after that pointer




# Approach #3: File Allocation Table (FAT)

- **File Allocation Table (FAT) is an important variation of linked lists**
  - Don't embed the pointers of the linked list with the file data blocks themselves
  - Instead, separate the pointers out and put them in a special table – the file allocation table (FAT)
  - The FAT is located at a section of disk at the beginning of a volume

# File Allocation Table

- entries in the FAT point to other entries in the FAT as a linked list, but their values are interpreted as the disk block number
- unused blocks in FAT initialized to 0



# File Allocation Table

- **FAT file systems used in MS-DOS and Win95/98**
  - Bill Gates designed/coded original FAT file system
  - replaced by NTFS (basis of Windows file systems from WinNT through Windows Vista/7)
  - Variants include FAT16, FAT32, etc. FAT16 and FAT32 refer to the size of the address used in the FAT.

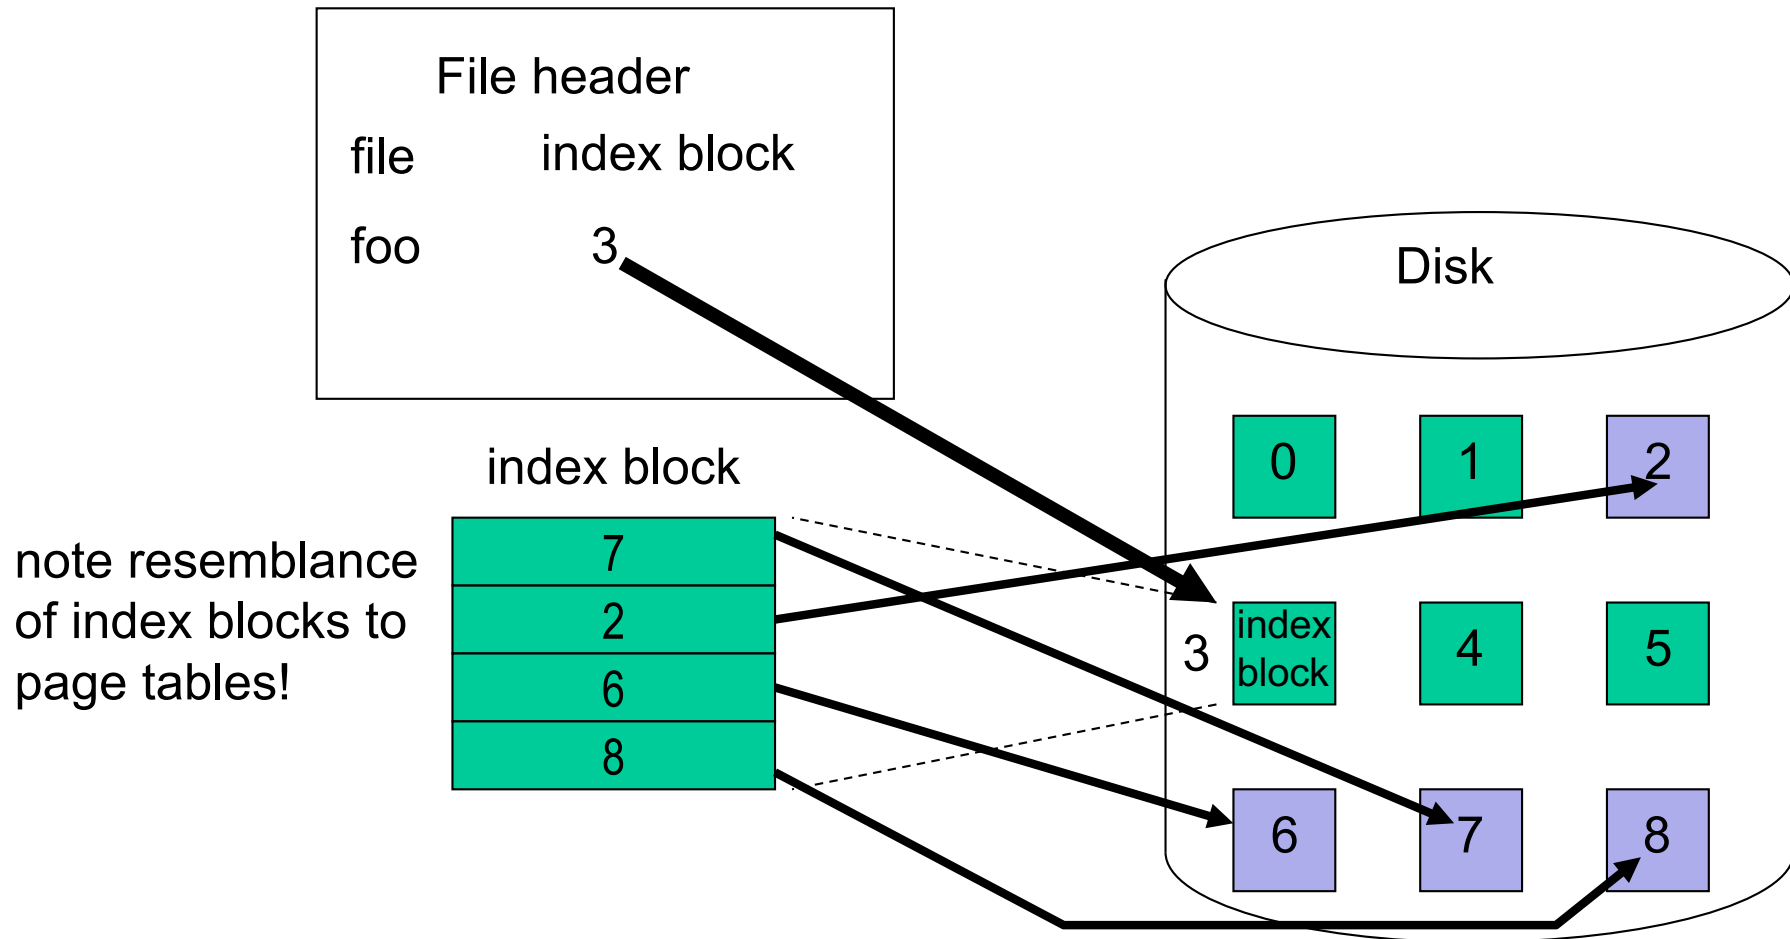
# File Allocation Table

- **Linked list for a file is terminated by a special end-of-file EOF value**
- **Allocating a new block is simple - find the first 0-valued block**
- **Advantage: random Reads/Writes faster than pure linked list**
  - the pointers are all colocated in the FAT near each other at the beginning of disk volume - low disk seek time
- **Still have to traverse the linked list to find location of data – this is a slow operation**

# Approach #4: Indexed Allocation

- **Conceptually, collect all pointers into a list or table called an *index block***
  - the index  $j$  into the list or index block retrieves a pointer to the  $j$ 'th block on disk
  - Looks kind of like a page table, except it's extensible
- **Unlike the FAT**
  - the index block can be stored in any block on disk, not just in a special section at the beginning of disk
  - the index is just a linear list of pointers

# Indexed Allocation



# Indexed Allocation

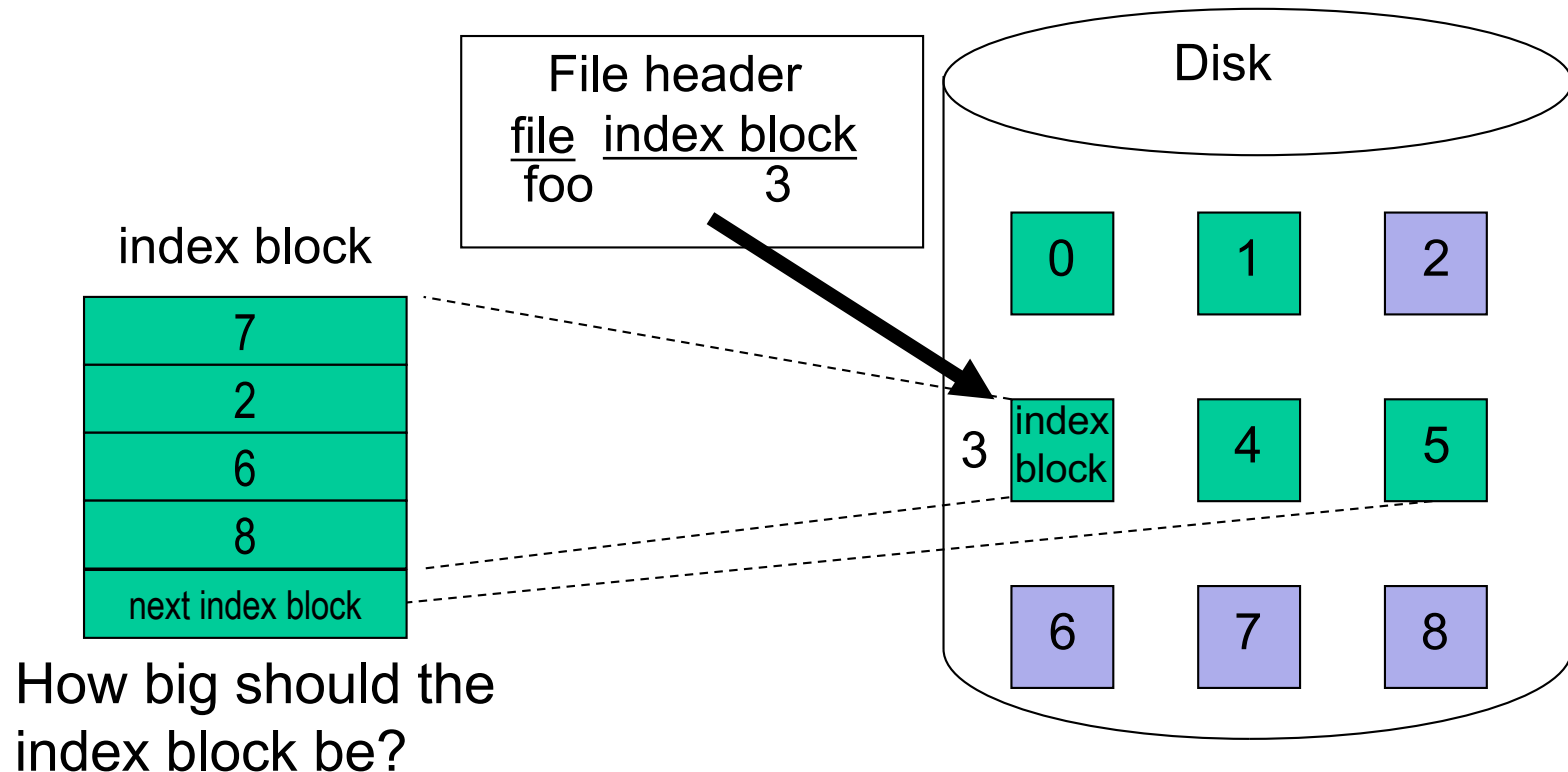
- **Solves many problems of contiguous and linked list allocation:**
  - no external fragmentation
  - size of file not required a priori
  - don't have to traverse linked list for random/direct reads/writes
    - just index quickly into the index block



# Indexed Allocation

## ■ Solutions:

1. Link together index blocks –
  - each index block has link to next index block
2. *Multilevel index* (like hierarchical page tables!)
  - First level is list of all index blocks for file
  - Second level is list of all data blocks in that section of the file



# Indexed Allocation

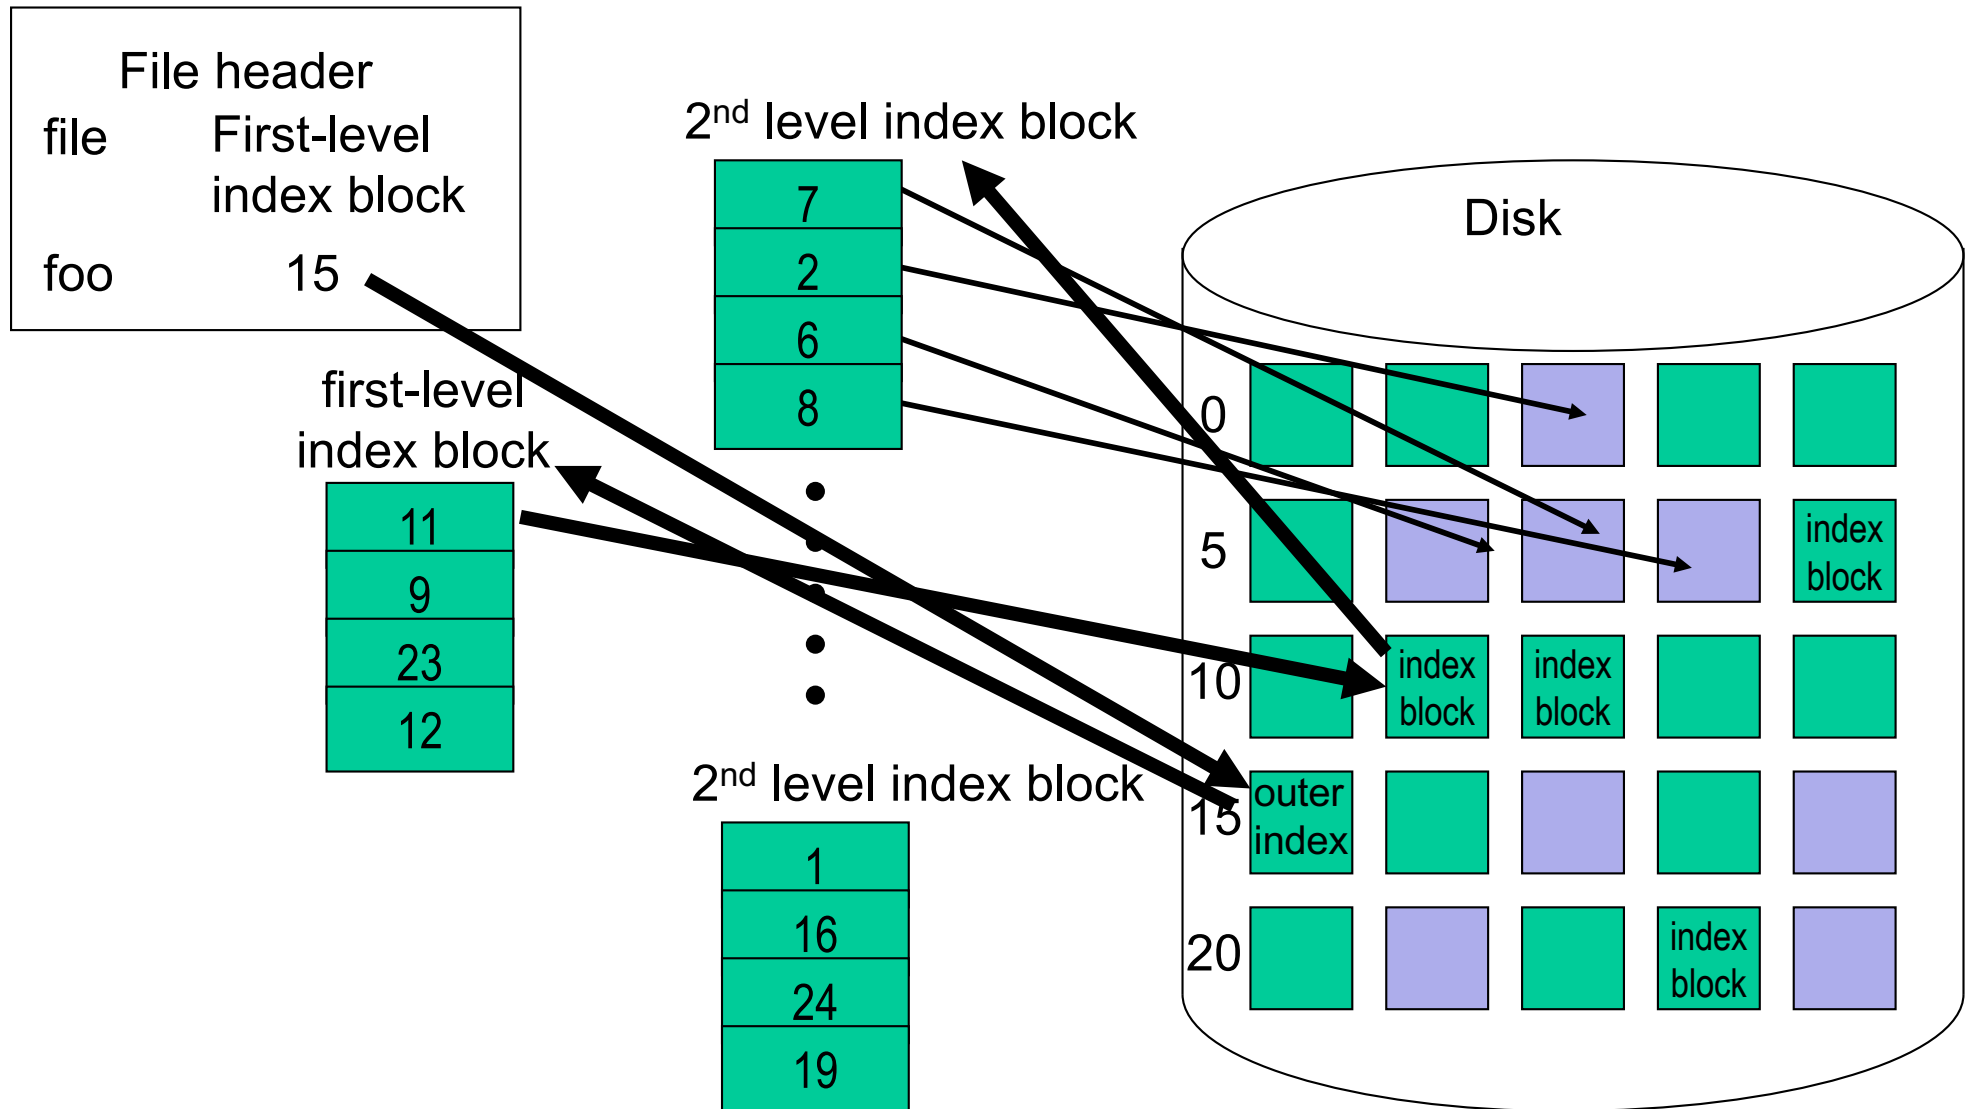
- **Problem: how big should the index block be?**

- if the index block is too large, then there are many wasted/empty entries for small files
- if the index block is too small, then there are not enough entries for large files

- **Solutions:**

- link together index blocks
- *multilevel index* (like hierarchical page tables!)
  - indexing into the first-level index block provides a pointer to second-level index blocks. Indexing into the second-level index block (using a different offset) retrieves a pointer to the actual file block on disk

# Approach #5: Multilevel Indexed Allocation



# Multilevel Indexed Allocation

- **Don't have to allocate unused second-level index blocks!**
- **Maximum file size?**

two levels of index blocks,

1024 pointer entries/block => 1 million addressable data blocks.

If each block is 4 KB, then the largest file size is 4 GB.

- **Problems with multi-level indexing:**
  - accessing small files takes just as long as large files
  - have to go through the same # of levels of indexing, hence same # of disk operations
  - accessing the data of a 100 byte file requires at least 4 block reads

# Approach #6:

## UNIX Multilevel Indexed Allocation

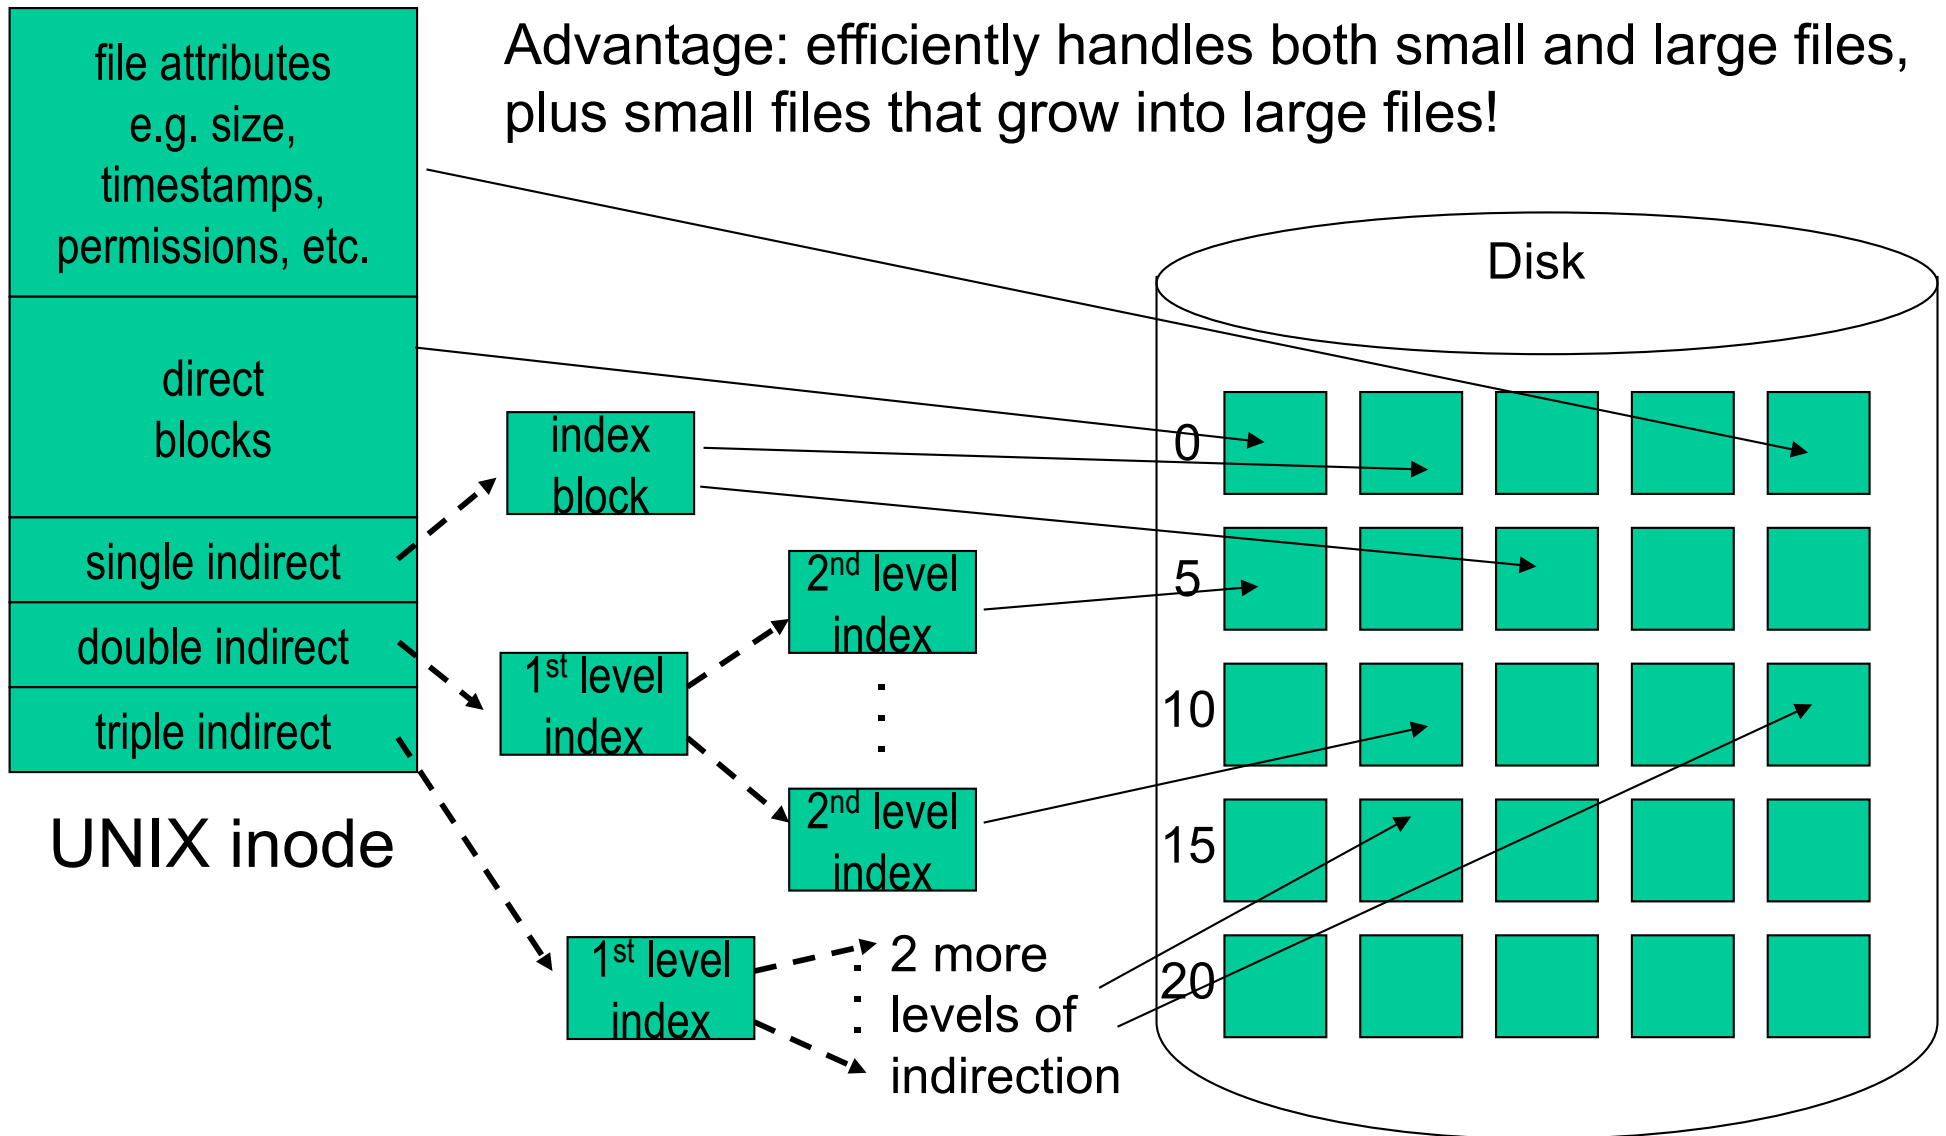
- UNIX (and Linux ext2fs, ext3fs, etc.) uses this variation of multi-level indexing to accommodate large and small files
  - Suppose there are 15 entries in the index block
  - the first 12 entries are pointers to direct blocks of file data on disk
  - the 13th pointer points to a singly indirect block, which is an index block pointing to disk blocks
  - the 14th pointer points a *doubly indirect block* (2 levels of index blocks)
  - the 15th pointer points to a *triply indirect block* (3 levels of index blocks)

# **Approach #6:**

## **UNIX Multilevel Indexed Allocation**

- **UNIX (and Linux ext2fs, ext3fs, etc.) uses this variation of multi-level indexing to accommodate large and small files**

# UNIX Multilevel Indexed Allocation





# UNIX Multilevel Indexed Allocation

- **for small files**

- only uses a small index block of 15 entries, so there is very little wasted memory

- **for large files**


- the indirect pointers allow expansion of the index block to span a large number of disk blocks

# Comparing File Allocation with Process Allocation

- **In both cases, mapping an entity to storage**
  - Process address space allocated frames in RAM via page tables
  - File data is allocated to disk/flash
- **Differences:**
  - Address spaces are fixed in size and known in advance,
  - Files grow/contract over time – files need a mapping/allocation system that is more flexible than page tables, which can't grow
  - Address spaces can be sparse and mostly unused, while file data is all “used”

# Free Space Management

- **Another aspect of managing a file system is managing free space**
  - the file system needs to keep track of what blocks of disk are free/unallocated
  - keeps a free-space “list”
  - In this example, need to keep track that disk blocks 2, 3 and 8 are free/unallocated



# Free Space Management Approaches


## 1. Bit Vector or Bit Map

- Each block is represented by a bit.
- Concatenate all such bits into an array of bits, namely a bit vector.
  - The  $j$ 'th bit indicates whether the  $j$ 'th block has been allocated.
  - if bit = 1, then a block is free, else if bit = 0, then block is allocated

# Free Space Management Approaches


## 2. Linked List

- link together all free blocks
- efficient - keeps track of only the free blocks.
  - bitmap has the overhead of tracking both free and allocated blocks - this is wasteful if memory is mostly allocated
- Faster than bitmap – find 1<sup>st</sup> free block immediately



# Free Space Management Approaches


- **Problem with Linked List free space management:**
  - traversing the free list is slow if you want to allocate a large number of free blocks all at once
    - hopefully this occurs infrequently



# Free Space Management Approaches

## 3. Grouping


- linked list, except store  $n-1$  pointers to free blocks in each list block
- the last block points to the next list block containing more free pointers
- allows faster allocation of larger numbers of free blocks all at once



# Free Space Management Approaches

## 4. Counting -

- grouped linked list, but also add a field to each pointer entry that indicates the number of free blocks immediately after the block pointed to
- even faster allocation of large #'s of free blocks







Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox  
University of  
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.