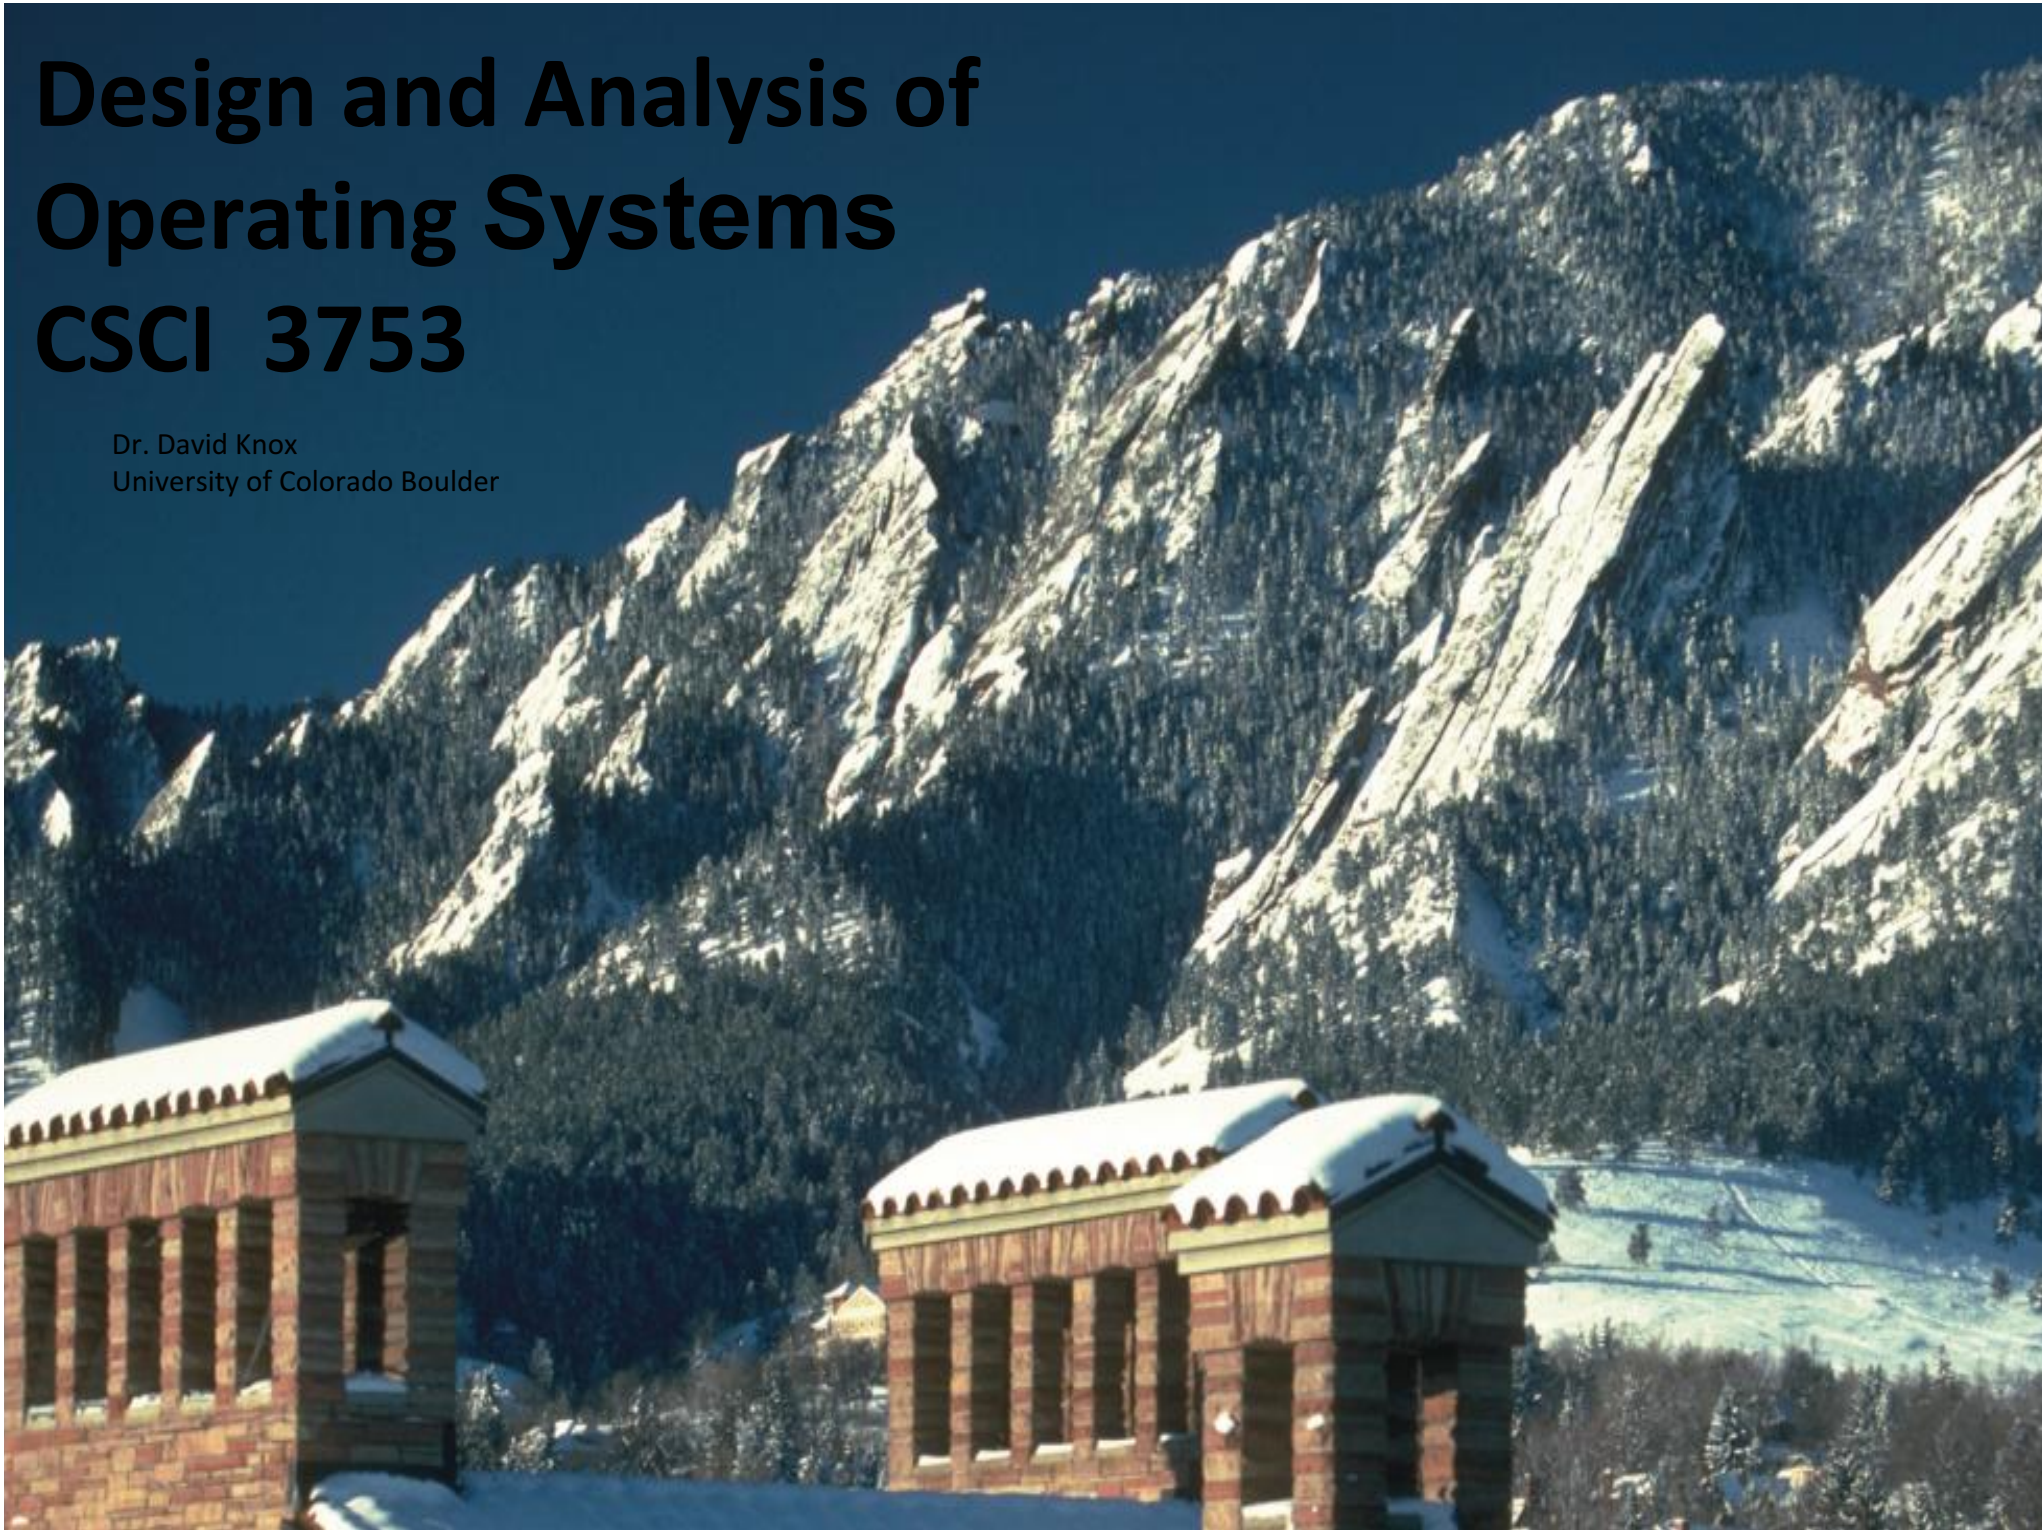


# Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox  
University of Colorado Boulder





Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Design and Analysis of Operating Systems CSCI 3753

## Flash Memory

Dr. David Knox  
University of  
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.

# Flash Memory





# What is Flash Memory?


- **Flash memory is a type of solid state storage**
  - bits are stored in an electronic chip
    - not on a magnetic disk
  - Programmed and erased electrically
    - no mechanical parts
  - Non-volatile storage
    - i.e. “permanent” with caveats
- **Primarily used in solid state drives (SSDs), memory cards, USB flash drives, mobile smartphones, digital cameras, etc.**



# What is Flash Memory?

## ■ Flash memory is a form of EEPROM

- Electrically Erasable Programmable Read Only Memory
- bits can be rewritten again and again using ordinary electrical signals
- non-volatile



## ■ Compare to:

RAM (Random Access Memory) – main memory, volatile

ROM (Read Only Memory) – code is burned in at the factory and can only be read thereafter, no writes

PROM (Programmable ROM) – code can be written once using a special machine

EPROM (Erasable PROM) – code can be written multiple times with a special machine

# What is Flash Memory?

- **Advantages vs. disk:**
  - much faster access (lower latency)
  - more resistant to kinetic shock (& intense pressure, water immersion etc.)
  - more compact
  - lighter
  - lower power (typically 1/3 - 1/2 of disks)



# What is Flash Memory?

- **Disadvantages vs. disk:**
  - more costly per byte,
  - limited lifetime,
  - erases are costly, ...
- **Flash's name is due to its electronic circuit design, which can inexpensively erase a large amount of solid state memory quickly, like a "flash" of light**



# NAND Flash

- Smaller memory cell area (less expensive)
- Word accessible (large granularity, 100-1000 bits/word, good for secondary storage where files don't mind being read out in large chunks/words)
- Slower random byte access (have to read a word)
- Short erasing (~2 ms) and programming times (~5 MB/sec sustained write speed)
- Longer write lifetime

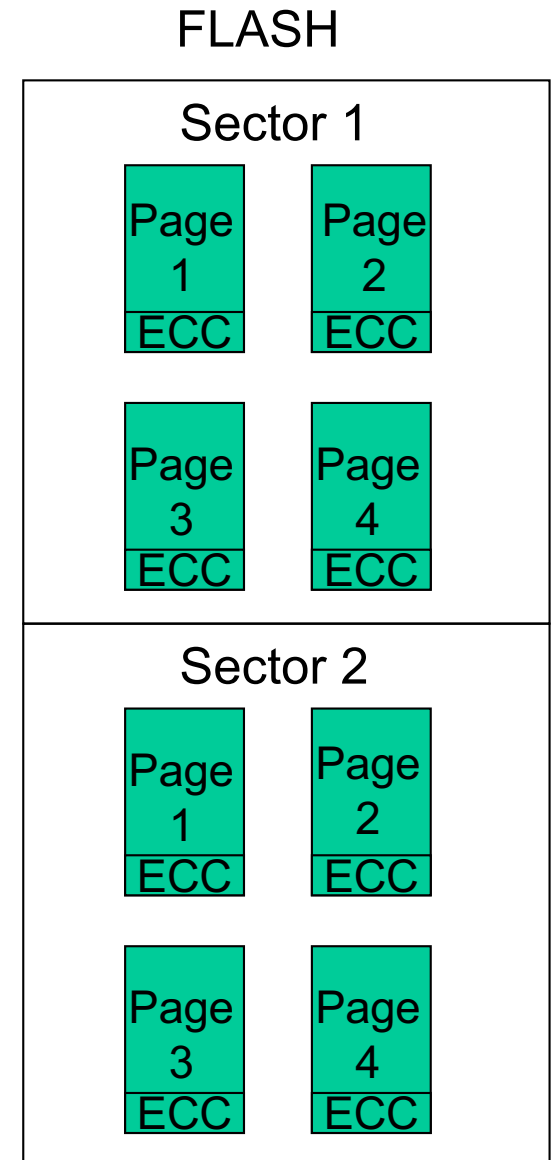


# NOR Flash

- Larger memory cell area
- Byte accessible (good for ROM-like program storage, where instructions need to be read out on a byte-size granularity)
- Faster random byte access
- Longer erasing ( $\sim 750$  ms) and programming times ( $\sim .2$  MB/sec sustained write speed)
- Shorter write lifetime
- iPhone uses both multi-GB NAND flash for multimedia file storage and multi-MB NOR flash for code/app storage


# Organizing Flash Memory

- Flash memory is typically divided into blocks or *sectors*
  - There can be many *pages* per block/sector
    - e.g. 64 pages per block, 16 blocks per flash, and 4 KB per block => 4 MB of flash
  - The last 12-16 bytes of each page is reserved for error detection/correction (ECC) and bookkeeping
    - the flash file system can put information in this space



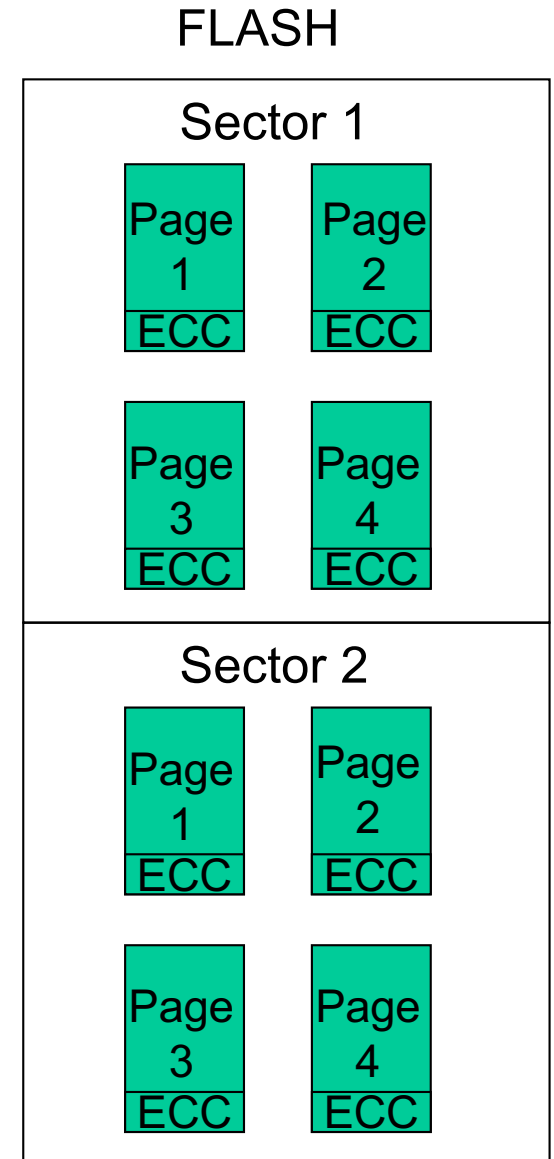
# Operations on Flash Memory

- Reads and writes occur on a sub-page level granularity
  - Can read a byte (NOR) or word (NAND) by giving page + offset
- However, writes are tricky
  - Writes of a byte (NOR) or word (NAND) only proceed immediately if that memory has been cleanly reset/erased and not yet written to
  - “Rewrites” to memory that has already been written to *require an erase first before the write*



# Rewrites Are Costly

- **Unfortunately, erasing in flash can only occur on a sector granularity!**
  - This is the price we pay for cheap “flash” technology
- **Hence, to write just one byte into a memory location that has already been written to**
  - First you have to erase the entire sector containing that byte address
  - Then write the byte!



# Rewrites Are Costly

- **Rewrites happen all the time.**
  - Example: deleting a file
    - De-allocate the flash pages containing the file data
    - If we want to allocate these free flash pages to other files, we have to first erase the entire sectors containing these pages
  
- **In comparison, the cost of rewrites is the same as writes for magnetic disk:**
  - Can simply have magnetic head reset of the orientation of the bits on magnetic media
  - No need to erase before writing



# Rewrites Are Costly

- To rewrite over a page of flash memory *while preserving* the other pages in the same sector:
  - Either:
    1. Copy the entire sector into RAM
    2. In RAM, write the changes to the page that you want to modify
    3. Erase the sector on flash
    4. Write all the sector's pages currently in RAM to the newly erased sector, including the modified page
  - Or:
    - Copy the entire sector into a new clean sector, except for the page to be modified
    - Write the one page to the correct location in the new sector (page is clean)

# Limited Lifetime of Flash Memory

- **Memory wear - Flash memory has a limited write lifetime**

- NOR flash memory can withstand 100,000 write-erase cycles before becoming unreliable
- NAND flash memory can withstand 1,000,000 write-erase cycles
- The strong electric field needed to set and reset bits eventually breaks down the silicon transistor

- **Eventually, your mobile devices can no longer save data!**

- Fortunately, you'll probably upgrade your mobile device long before the lifetime limit is reached

# Wear Leveling Solution

- **To extend the life of flash memory chips, write and erase operations are spread out over the entire flash memory**
  - Keep a count in 12-16 byte trailer of each page of how many writes have been made to that page, and choose the free page with the lowest write count
  - Randomly scatter data over the entire flash memory span, since it doesn't take any longer to extract data from nearby pages as distant pages
  - Many of today's flash chips implement wear leveling in the hardware's device controller

# Problems Applying Standard Disk File Systems to Flash Memory

1. **A typical disk file system stores data/metadata in fixed locations**
  - the directory structure, file headers, allocation structures (e.g. FAT) and free space structures (e.g. bitmap) and even file data are in relatively fixed locations
  - Repeated modifications to these structures in static locations on flash would result in rapidly exhausting the write lifetimes of those locations, hence all of flash

# Problems Applying Standard Disk File Systems to Flash Memory

## 2. Disk file systems are not optimized for rewrites

- Disk file system rewrites are viewed essentially the same as writes, and repeated rapid rewrites can be quickly done because there is little disk seek time
- Flash rewrites incur a much greater penalty, are much longer than writes to clean flash pages, and repeated rapid rewrites in flash would cause disastrously long latency, not to mention reduced lifetime



# Flash File Systems

- **Must be designed to deal with 2 major problems of flash memory:**
  - Rewrites requires erases
  - Limited lifetime
- **Log-structured file systems are well-suited to flash memory!**
  - Writes to a file are added/appended to the end of the log, as are rewrites. The log is the file system.
  - The log grows sequentially and doesn't rewrite over the same location
  - Hence less erasing & also longer lifetime

# Flash File Systems

## ■ Journaling Flash File System (JFFS) and JFFS2

- Is a Log-structured file system similar to log-based recovery seen earlier for file system reliability
- The file system is the log, and is written sequentially over flash pages,
  - The entries in the log contain all the data we need to reconstruct the file (recall the X & Y old and new values)
  - The log may be circular, eventually wrapping around once all free flash pages have been consumed.
    - Free space = distance between head and tail of log
    - As free space decreases, garbage collection is performed to free up some more space.



Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox  
University of  
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.