

# Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox  
University of Colorado Boulder



# Adding a New Device



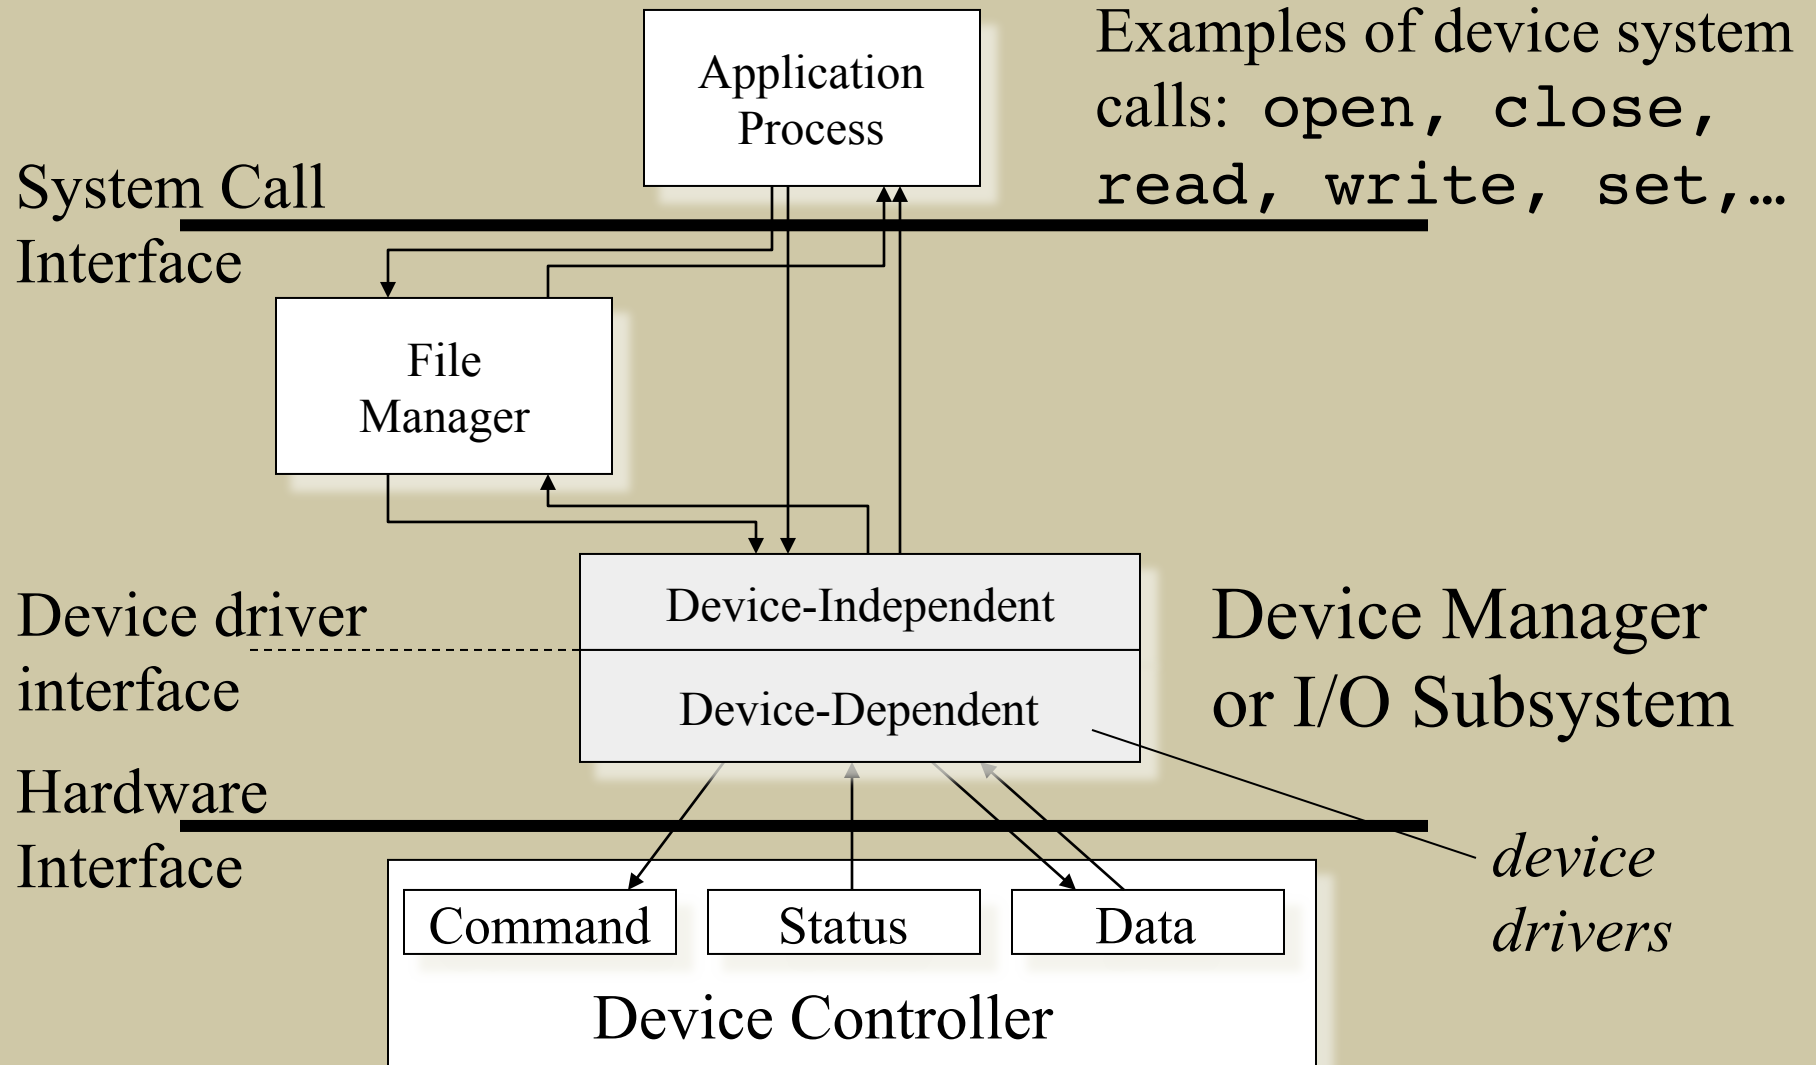
Design and Analysis of  
Operating Systems  
CSCI 3753

Dr. David Knox  
University of Colorado  
Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.



# Device Management Organization



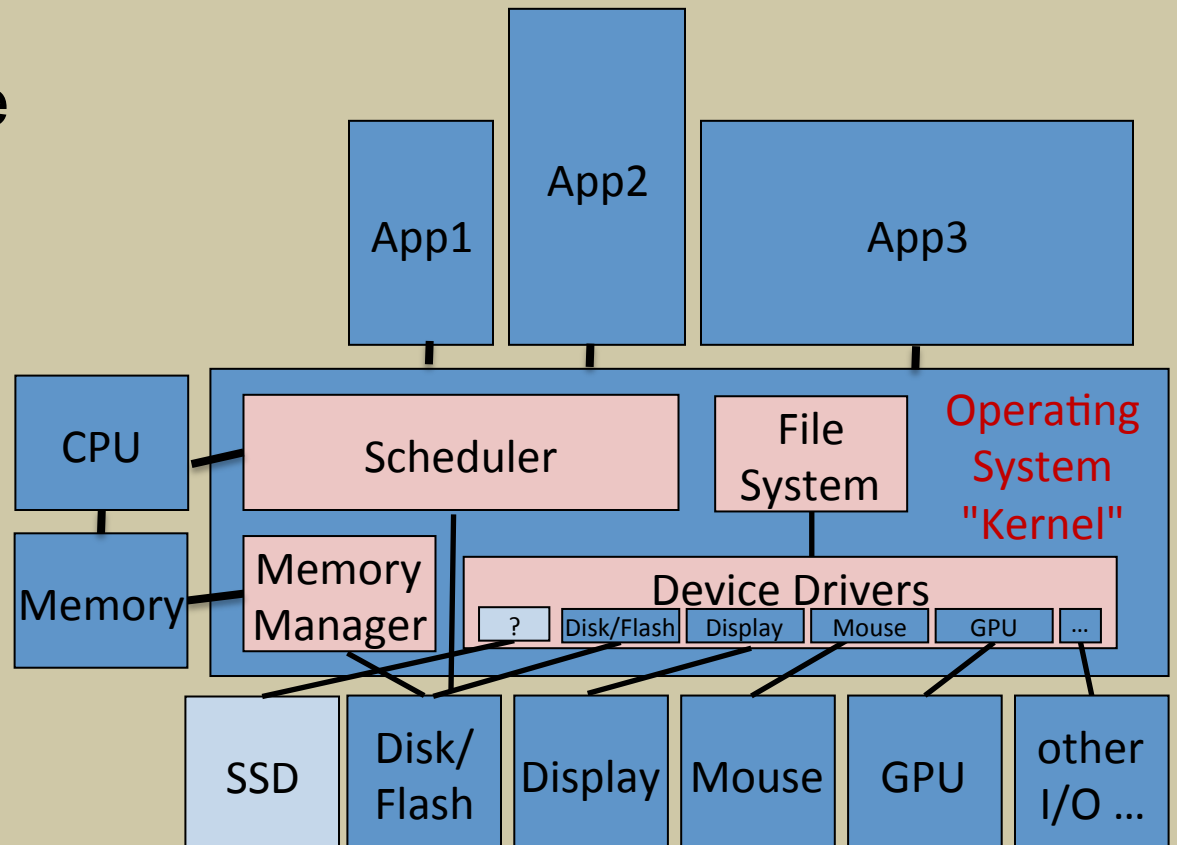
# Devices have both device-independent and device-dependent code

There is special device driver code associated with each different device connected to the system

**1. Device-Independent API**

**2. Device-Dependent** driver code

**3. Device Controller**



# Device Independent Part

- A set of system calls that an application program can use to invoke I/O operations
- A particular device will respond to only a subset of these system calls
  - A keyboard does not respond to *write()* system call
- POSIX set: *open()*, *close()*, *read()*, *write()*, *lseek()* and *ioctl()*

# Device Independent Function Call

Trap Table

func <sub>i</sub> (...)

```
dev_func_i(devID, ...) {  
    // Processing common to all devices  
    ...  
    switch(devID) {  
        case dev0: dev0_func_i(...);  
                    break;  
        case dev1: dev1_func_i(...);  
                    break;  
        ...  
        case devM: devM_func_i(...);  
                    break;  
    };  
    // Processing common to all devices  
    ...  
}
```



# Adding a New Device

- Write device-specific functions for each I/O system call
- For each I/O system call, add a new *case* clause to the *switch* statement in device independent function call



## Trap Table

<code>func<sub>i</sub>(...)</code>

```
dev_func_i(devID, ...) {  
    // Processing common to all devices  
    ...  
    switch(devID) {  
        case dev0: dev0_func_i(...);  
                    break;  
        case dev1: dev1_func_i(...);  
                    break;  
        ...  
        case devM: devM_func_i(...);  
                    break;  
        case devNew: devNew_func_i(...);  
        break;  
    };  
    // Processing common to all devices  
    ...  
}
```





# Adding a New Device

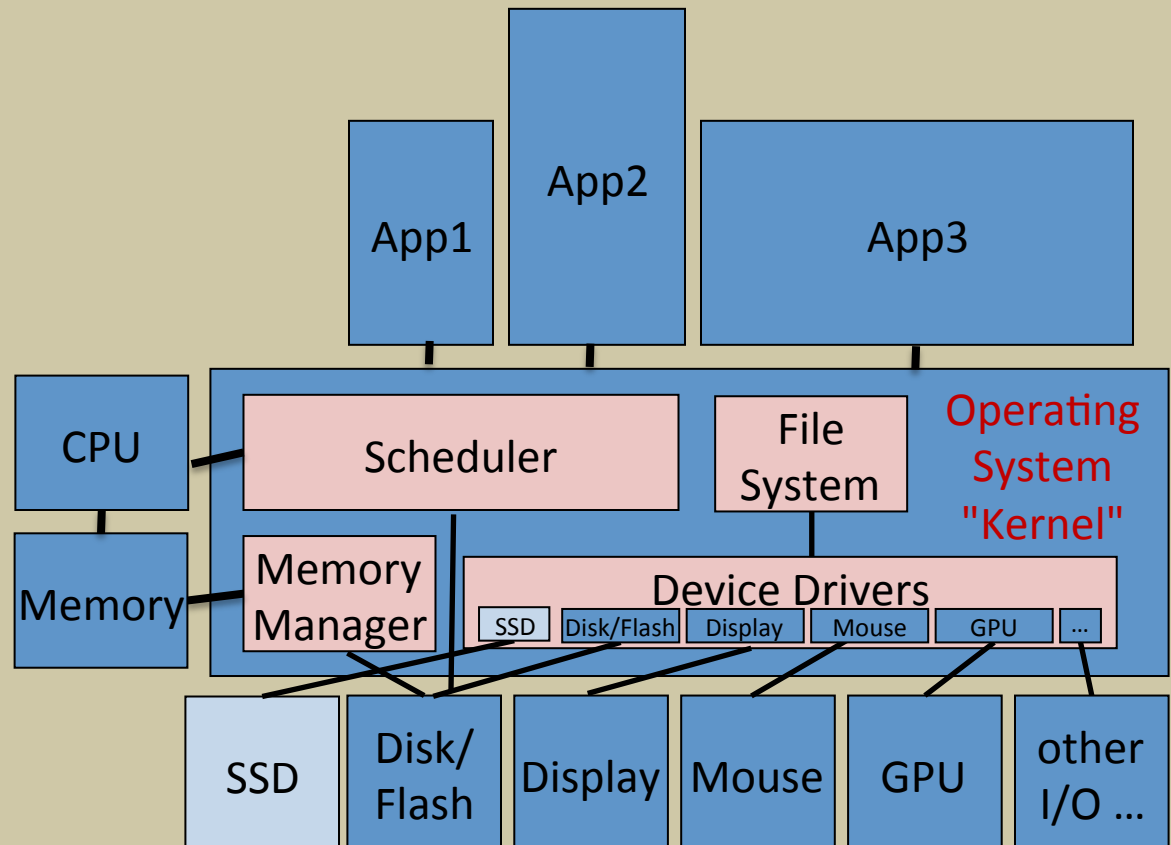
- After updating all `dev_func_*(...)` in the kernel, compile the kernel

Problem: Need to recompile the kernel, every time a new device or a new driver is added

# Devices have both device-independent and device-dependent code

We need a way to dynamically add new code into the OS Kernel when a new device needs to be supported.

1. *Load Device-Dependent driver code into kernel*
2. *Only load the device drivers as needed*
3. *Do not want to recompile kernel for changes in the device driver*



# Loadable Kernel Modules

- A loadable kernel module (LKM) is an object file that contains code to extend a running kernel
- Windows (kernel-mode driver), Linux (LKM), OS X (Kernel extension: kext)
- Without loadable kernel modules, an OS would have to include all possible anticipated functionality already compiled directly into the base kernel – monolithic kernel
- LKMs can be loaded and unloaded from kernel on demand at runtime



# LKMs

- Offer an easy way to extend the functionality of the kernel without having to rebuild or recompile the kernel again
- Simple and efficient way to create programs that reside in the kernel and run in privileged mode
- Most of the drivers today are written as LKMs
- See */lib/modules* for the all the LKMs
- *lsmod*: lists all kernel modules that are already loaded
  - Reads */proc/modules* file

# How to write a kernel module

- Kernel Modules are written in the C programming language
- You must have a Linux kernel source tree to build your module
- You must be running the same kernel you built your module with to run it
- Linux kernel object: *.ko* extension





# Kernel Module: Basics

- A kernel module file has several typical components:
  - `MODULE_AUTHOR("your name")`
  - `MODULE_LICENSE("GPL")`
    - The license must be an open source license (GPL, BSD, etc.) or you will “taint” your kernel.
    - Tainted kernel loses many abilities to run other open source modules and capabilities.



# Kernel Module: Key Operations

- `int init_module(void)`
  - Called when the kernel loads your module.
  - Initialize all your stuff here.
  - Return 0 if all went well, negative if something blew up.
- Typically, `init_module( )` either
  - registers a handler for something with the kernel
  - or replaces one of the kernel functions with its own code(usually code to do something and then call the original function)

# Kernel Module: Key Operations

- `void cleanup_module(void)`
  - Called when the kernel unloads your module.
  - Free all your resources here.



# Hello World Example

```
#include <linux/kernel.h>
#include <linux/module.h>
MODULE_AUTHOR("Shiv Mishra");
MODULE_LICENSE("GPL");

int init_module(void)
{
    printk(KERN_ALERT "Hello world: I am the developer
                                speaking from the Kernel");
    return 0;
}
```

# Hello World Example

```
void cleanup_module(void)
{
    printk(KERN_ALERT "Goodbye from the developer,
                      I am exiting the Kernel");
}
```





# Building Your Kernel Module

- Accompany your module with a 1-line GNU Makefile:
  - `obj-m += hello.o`
  - Assumes file name is “hello.c”
- Run the magic make command:
  - `make -C <kernel-src> M=`pwd` modules`
  - Produces: `hello.ko`
- Assumes current directory is the module source.

# obj-\$(CONFIG\_FOO) += foo.o

- Good definitions are the main part (heart) of the kbuild Makefile
- The most simple kbuild makefile contains one line:  
obj-\$(CONFIG\_FOO) += foo.o
- Tell kbuild that there is one object in that directory named foo.o
  - foo.o will be built from foo.c or foo.S
- \$(CONFIG\_FOO) evaluates to either y (for built-in) or m (for module). If CONFIG\_FOO is neither y nor m, then the file will not be compiled nor linked.



# Loading Your Kernel Module: *insmod*

- Use *insmod* to manually load your kernel module

*sudo insmod helloworld.ko*

- *insmod* makes an *init\_module* system call to load the LKM into kernel memory
- *init\_module* system call invokes the LKM's initialization routine (also called *init\_module*) right after it loads the LKM
- The LKM author sets up the initialization routine to call a kernel function that registers the subroutines that the LKM contains



# Where is our Hello World message

- dmesg
- /var/log/system.log

# Unloading Your Kernel Module

- Use *rmmod* command  
rmmod hello.ko
- Should print the Goodbye message



# Kernel Module Dependencies: *modprobe*

- insmod/rmmod can be cumbersome...
  - You must manually enforce inter-module dependencies.
- *modprobe* automatically manages dependent modules
  - Copy hello.ko into /lib/modules/<version>
  - Run depmod
  - modprobe hello / modprobe -r hello
- Dependent modules are automatically loaded/unloaded.

- *depmod* creates a Makefile-like dependency file, based on the symbols it finds in the set of modules mentioned on the command line or from the directories specified in the configuration file
- This dependency file is later used by *modprobe* to automatically load the correct module or stack of modules

# modinfo command

- .ko files contain an additional .modinfo section where additional information about the module is kept
  - Filename, license, dependencies, ...
- modinfo command retrieves that information  
modinfo hello.ko

# Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox

University of Colorado  
Boulder