

# Design and Analysis of Operating Systems

## CSCI 3753

Dr. David Knox  
University of Colorado Boulder





# Monitors and Conditional Variables



Design and Analysis of  
Operating Systems  
CSCI 3753

Dr. David Knox  
University of Colorado  
Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.



University of Colorado  
Boulder

# Deadlock when using Semaphores

- It can easily occur
  - two tasks, each desires a resource locked by the other process
  - circular dependency
  - programming errors
    - switching order of P() and V()
    - calling wait multiple times
    - forgetting a signal
- Semaphores provide mutual exclusion, but can introduce deadlock



# Monitors

- Semaphores can result in deadlock due to programming errors
  - forgot to add a P() or V(), or misordered them, or duplicated them
- To reduce these errors, introduce high-level synchronization primitives, e.g. *monitors with condition variables*
  - essentially automates insertion of P() and V() for you
  - As high-level synchronization constructs, monitors are found in high-level programming languages like Java and C#
  - underneath, the OS may implement monitors using semaphores and mutex locks



# Monitors

- Declare a monitor as follows (looks somewhat like a C++ class):

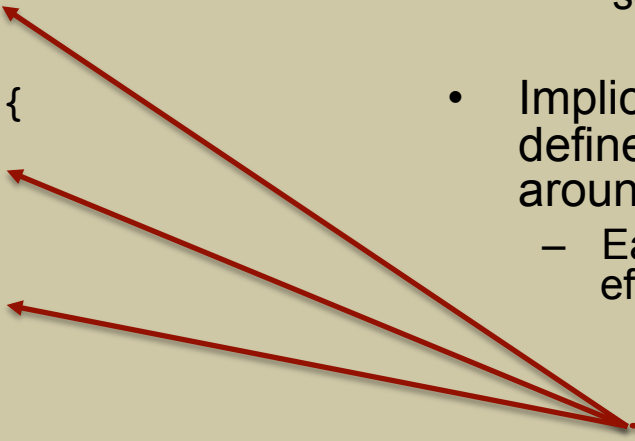
```
monitor monitor_name {  
    // shared local variables  
  
    function f1(...) {  
        ...  
    }  
    ...  
    function fN(...) {  
        ...  
    }  
    init_code(...) {  
        ...  
    }  
}
```

- A monitor ensures that only 1 process/thread at a time can be active within a monitor
  - simplifies programming, no need to explicitly synchronize

- Implicitly, the monitor defines a mutex lock  
semaphore mutex = 1;

- Implicitly, the monitor also defines mutual exclusion around each function
  - Each function's critical code is effectively:

```
function fj(...) {  
    P(mutex)  
    // critical code  
    V(mutex)  
}
```



# Monitors

- Declare a monitor as follows (looks somewhat like a C++ class):

```
monitor monitor_name {  
    // shared local variables  
    int count;  
    data_type data[MAX_COUNT];  
  
    function add_item(...) {  
        ...  
    }  
  
    ...  
    function remove_item(...) {  
        ...  
    }  
  
    init_code(...) {  
        ...  
    }  
}
```

- The monitor's private local variables can only be accessed by local monitor functions
- Each function in the monitor can only access variables declared locally within the monitor and its parameters



# Monitors and Condition Variables

- Previous definition of a monitor achieves
  - mutual exclusion
  - hiding of wait() and signal() from user
  - loses the ability that semaphores had to enforce order
    - wait() and signal() are used to provide mutual exclusion
    - but have lost the unique ability for one process to signal another blocked process using signal()
    - there is no way to have a process sleep waiting on the signal
- In general, there may be times when one process wishes to signal another process based on a condition, much like semaphores.
  - Thus, augment monitors with *condition variables*.



# Condition Variables

- Augment the mutual exclusion of a monitor with an ordering mechanism
  - Recall: Semaphore  $P()$  and  $V()$  provide both mutual exclusion *and* ordering
  - Monitors alone only provide mutual exclusion
- A condition variable provides ordering
  - Used when one task wishes to wait until a condition is true before proceeding
    - Such as a queue being full enough or data being ready
  - A 2<sup>nd</sup> task will signal the waiting task, thereby waking up the waiting task to proceed





# Monitors and Condition Variables

condition *y*;

A condition variable *y* in a monitor allows three operations

- *y.wait()*
  - blocks the calling process
  - can have multiple processes suspended on a condition variable, typically released in FIFO order
  - textbook describes another variation specifying a priority *p*, i.e. call *y.wait(p)*
- *y.signal()*
  - resumes exactly 1 suspended process.
  - If no process is waiting, then function has *no effect*.
- *y.queue()*
  - Returns true if there is at least one process blocked on *y*



# A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }

    void release() {
        busy = FALSE;
        x.signal();
    }

    initialization code() {
        busy = FALSE;
    }
}
```



# Condition Variables Example

Block Task 1 until a condition holds true, e.g. queue is empty

```
condition wait_until_empty;
```

Task 1:

```
wait_until_empty.wait();  
... // proceed after queue  
      empty
```

Task 2:

```
...  
// queue is empty so signal  
wait_until_empty.signal();
```

**Problem:** if Task 2 signals before Task 1 waits, then Task 1 waits *forever* because CV *has no state*!



# Condition Variables vs Semaphores

Both have wait() and signal(), but semaphore's signal()/V()  
*preserves state* in its integer value

```
Semaphore wait_until_empty=0;
```

Task 1:

```
wait(wait_until_empty);  
... // proceed after queue  
      empty
```

Task 2:

```
...  
// queue is empty so signal  
signal(wait_until_empty); //V()
```

if Task 2 signals before Task 1 waits, then Task 1 does *not* wait forever because semaphore has state and remembers earlier signal()





# Complex Conditions

- Suppose you want task T1 to wait until a complex set of conditions set by T2 becomes TRUE
  - use a condition variable

```
condition x;  
int count=0;  
float f=0.0;
```

T1

----

...

```
while(f!=7.0 && count<=0) {  
    x.wait();  
}
```

... // proceed

T2

----

...

```
f=7.0;  
count++;  
x.signal();
```

**Problem:** could be a race condition in testing and setting shared variables



# Complex Conditions (2)

```
lock mutex;  
condition x;  
int count=0;  
float f=0.0;
```

- Surround the test of the conditions with a mutex to atomically test the set of conditions

T1

----

...

```
Acquire(mutex);  
while(f!=7.0 && count<=0) {  
    Release(mutex);  
    x.wait();  
    Acquire(mutex);  
}  
Release(mutex);  
... // proceed
```

T2

----

...

```
Acquire(mutex);  
f=7.0;  
count++;  
Release(mutex);  
x.signal();
```



# Complex Conditions (3)

```
lock mutex;  
condition x;  
int count=0;  
float f=0.0;
```

- pthreads replaces complex sequence of release/wait/acquire() with:

```
pthread_cond_wait(&cond_var, &mutex)
```

T1

----

...

```
Acquire(mutex);  
while(f!=7.0 && count<=0) {  
    pthread_cond_wait(&x, &mutex);  
}  
Release(mutex);  
... // proceed
```

T2

----

...

```
Acquire(mutex);  
f=7.0;  
count++;  
Release(mutex);  
x.signal();
```



# Broadcast Signals

- In some cases, you may want to wake all tasks blocked on a condition variable *x*, not just one
  - *x.signal()* only wakes one task
- *x.broadcast()* is a 3<sup>rd</sup> operation provided for CVs on some systems
  - Wakes all tasks blocked on a CV
  - In pthreads, `pthread_cond_broadcast ( )` wakes all waiting threads





# Monitors

- Declare a monitor as follows (looks somewhat like a C++ class):

```
monitor monitor_name {  
    // shared local variables  
    int count;  
    data_type data[MAX_COUNT];  
    condition y;  
  
    function add_item(...) {  
        ...  
        y.wait();  
        ...  
    }  
  
    ...  
    function remove_item(...) {  
        ...  
        y.signal();  
        ...  
    }  
  
    init_code(...) {  
        ...  
    }  
}
```

- The monitor's local variables can only be accessed by local monitor functions
- Each function in the monitor can only access variables declared locally within the monitor and its parameters



# Monitors and Condition Variables

- Problem: A monitor ensures that only 1 process/thread at a time can be active within a monitor
  - if a process P1 calls `y.signal()` it would wake another process P2 blocked on `y.wait()`
  - But we must avoid having two processes at the same time in the monitor
  - Need “wake-up” semantics on a `y.signal()`
- Two solutions proposed to only have one process active:
  - Hoare semantics, also called signal-and-wait
    - The signaling process P1 waits for P2 to either leave the monitor or wait on another condition, before resuming
  - Mesa semantics, also called signal-and-continue
    - The signaled process P2 continues to wait until the signaling process P1 leaves the monitor or wait on another condition



# Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox

University of Colorado  
Boulder



University of Colorado  
Boulder