

# Design and Analysis of Operating Systems

## CSCI 3753

Dr. David Knox  
University of Colorado Boulder





# Mutexes and Semaphores



Design and Analysis of  
Operating Systems  
CSCI 3753

Dr. David Knox  
University of Colorado  
Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.



University of Colorado  
Boulder

# Mutual exclusion using TS

```
shared boolean lock = FALSE;
shared int count;
Shared data_type buffer[MAX];
```

Code for  $p_1$

```
while(1) {
    produce (nextdata)
    while(count==MAX);
    Acquire(lock);
    buffer[count] = nextdata;
    count++;
    Release(lock);
}
```

Code for  $p_2$

```
while(1) {
    while(count==0);
    Acquire(lock);
    data = buffer[count-1];
    count--;
    Release(lock);
    consume (data);
}
```

CPU is spinning & waiting

How can we eliminate the busy waiting?



# Mutual exclusion using TS

```
shared boolean lock = FALSE;
shared int count;
Shared data_type buffer[MAX];
```

Code for  $p_1$

```
while(1) {
    produce (nextdata)
    while(count==MAX);
    Acquire(lock);
    buffer[count] = nextdata;
    count++;
    Release(lock);
}
```

Code for  $p_2$

```
while(1) {
    while(count==0);
    Acquire(lock);
    data = buffer[count-1];
    count--;
    Release(lock);
    consume (data);
}
```

- Look first at the busy wait when there is no room for data or no data
  - Need a way to pause a process/thread until the space is available or the data is available



# sleep( ) and wakeup( ) primitives

- *sleep()*: causes a running process to block
- *wakeup(pid)*: causes the process whose id is *pid* to move to ready state
  - No effect if process *pid* is not blocked



# Using **sleep( )** and **wakeup( )** primitives

```
// producer – place data into buffer
```

```
while(1) {  
    if (count == MAX) sleep();  
    buffer[count] = nextdata;  
    count++;  
    if (count == 1) wakeup (consumer);  
}
```

```
// consumer – take data out of buffer
```

```
while(1) {  
    if (count == 0) sleep();  
    getdata = buffer[count-1];  
    count--;  
    if (count == MAX-1) wakeup (producer);  
}
```



# Our solution using **sleep( )** and **wakeup( )**

- Still a problem with mutual exclusion of the critical sections containing counter++ and counter-- still exist
  - Could use TS, but that has busy waiting (we will solve this later)
- Possible problem with order of execution:
  - Consumer reads count and count == 0
  - Scheduler schedules the producer to run
  - Producer puts an item in the buffer and signals the consumer to wake up
    - Since consumer has not yet invoked sleep(), the wakeup() invocation by the producer has no effect
  - Scheduler now schedules the Consumer to continue, and the consumer blocks because it missed the wakeup() by the producer
  - Eventually, producer fills up the buffer and blocks
- How can we solve this problem?
  - Could use a mechanism to count the number of sleep() and wakeup() invocations



# Semaphores

- Dijkstra proposed more general solution to mutual exclusion
- Semaphore S is an abstract data type that is accessed only through two standard atomic operations
  - wait( ) (also called P(), from Dutch word *proberen* “to test”)
    - somewhat equivalent to a test-and-set
    - also involves *decrementing* the value of S
  - signal( ) (V(), from Dutch word *verhogen* “to increment”)
    - *increments* the value of S
- OS provides ways to create and manipulate semaphores atomically





# Semaphores

```
typedef struct {  
    int value;  
    PID *list[ ];  
} semaphore;
```

Both wait() and signal()  
operations are atomic

```
wait(semaphore *s) {  
    s→value--;  
    if (s→value < 0) {  
        add this process to s→list;  
        sleep ( );  
    }  
}
```

```
signal(semaphore *s) {  
    s→value++;  
    if (s→value <= 0) {  
        remove a process P from s→list;  
        wakeup (P);  
    }  
}
```



# Mutual Exclusion using Semaphores

## Binary Semaphore

```
semaphore S = 1; // initial value of semaphore is 1
int counter;      // assume counter is set correctly somewhere in code
```

### Process P1:

```
wait(S);
    // execute critical section
    counter++;
signal(S);
```

### Process P2:

```
wait(S);
    // execute critical section
    counter--;
signal(S);
```

- Both processes atomically wait() and signal() the semaphore S, which enables mutual exclusion on critical section code, in this case protecting access to the shared variable *counter*
- This will solve the mutual exclusion like a mutex lock, but will also eliminate the busy waiting



# Problems with Semaphores

shared R1, R2;

semaphore Q = 1; // binary semaphore as a mutex lock for R1

semaphore S = 1; // binary semaphore as a mutex lock for R2

*Process P1:*

wait(S);           (1)

wait(Q);           (4)

modify R1 and R2;

signal(S);

signal(Q);

*Process P2:*

wait(Q);           (2)

wait(S);           (3)

modify R1 and R2;

signal(Q);

signal(S);

- Potential for Deadlock

# Deadlock

- In the previous example,
  - Each process will block on a semaphore
  - The `signal()` statements will never get executed, so there is no way to wake up the two processes
  - There is no rule about the order in which `wait()` and `signal()` operations may be invoked
  - In general, with  $N$  processes sharing  $N$  semaphores, the potential for deadlock grows





# Other problematic scenarios for Semaphores

- A programmer mistakenly follows a wait() with a second wait() instead of a signal()
- A programmer forgets and omits the wait(mutex) or signal(mutex)
- A programmer reverses the order of wait() and signal()



# Another problem with synchronization

shared R1, R2;

semaphore S = 1; // binary semaphore as a mutex lock for R1 & R2

*Process P1:*

wait(S);

modify R1 and R2;

signal(S);

*Process P2:*

wait(S);

modify R1 and R2;

signal(S);

*Process P3:*

wait(S);

modify R1 and R2;

signal(S);

Potential for starvation



# Starvation

- The possibility that a process would never get to run
- For example, in a multi-tasking system the resources could switch between two individual processes
- Depending on how the processes are scheduled, a third process may never get to run
- The third task is being starved of accessing the resource



# Semaphore Solution for Mutual Exclusion

```
semaphore lock = 1; // initial value of semaphore is 1
shared int count;
shared data_type buffer[MAX];
```

Code for  $p_1$

```
while(1) {
    produce (nextdata)
    while(count==MAX);
    wait(lock);
    buffer[count] = nextdata;
    count++;
    signal(lock);
}
```

Code for  $p_2$

```
while(1) {
    while(count==0);
    wait(lock);
    data = buffer[count-1];
    count--;
    signal(lock);
    consume (data);
}
```

- Busy waiting removed from the mutual exclusion when waiting on lock
- Will show a complete solution in another video





# Does the Kernel need Synchronization?

- At any time, many kernel mode processes may be active
  - Share kernel data structures
  - Notice that even though user processes have their own address spaces, race conditions can still arise when they execute in kernel mode, e.g. executing a system call
- Preemptive and non-preemptive kernels
  - Preemptive kernel: allows a process to be preempted while running in kernel mode
    - Race conditions can occur
  - Non-preemptive kernel: does not allow a process to be preempted while running in kernel mode
    - Race conditions cannot occur



# Windows Synchronization

- Kernel level
  - Single processor system: temporarily mask interrupts for all interrupt handlers that may also access a shared resource
  - Multiprocessor system: use spin lock (busy waiting)
- User level
  - Dispatcher objects: mutex locks, semaphores, ...



# Linux Synchronization

- Kernel level
  - Prior to version 2.6, non-preemptive kernel, but later versions are fully preemptive
  - Atomic integers: all math operations on atomic integers are performed without interruptions

```
atomic_t counter;  
atomic_set(&counter, 5);  
atomic_add(10, &counter);  
...
```
  - Mutex locks, spin locks and semaphores, enabling/disabling interrupts on single processor systems
- User level
  - futex, semop( ): system call





University of Colorado  
Boulder

CSCI 3753 University of Colorado Boulder



# pthread Synchronization

- Mutex locks
  - Used to protect critical sections
- Some implementations provide semaphores through POSIX SEM extension
  - Not part of Pthread standard

```
#include <pthread.h>
```

```
pthread_mutex_t m; //declare a mutex object
```

```
pthread_mutex_init (&m, NULL); // initialize mutex object
```

```
//thread 1  
pthread_mutex_lock (&m);  
    //critical section code for th1  
pthread_mutex_unlock (&m);
```

```
//thread 2  
pthread_mutex_lock (&m);  
    //critical section code for th2  
pthread_mutex_unlock (&m);
```



# pthread mutex

- pthread mutexes can have only one of two states: lock or unlock
- Important restriction
  - Mutex ownership: Only the thread that locks a mutex can unlock that mutex
  - So, mutexes are strictly used for mutual exclusion while binary semaphores can also be used for synchronization between two threads



# POSIX semaphores

```
#include <semaphore.h>
```

```
int sem_init(sem_t *sem, int pshared, unsigned int value);  
// pshared: 0 (among threads); 1 (among processes)
```

```
int sem_wait(sem_t *sem); //same as wait( )
```

```
int sem_post(sem_t *sem); //same as signal( )
```

```
sem_getvalue( ), sem_close( )
```



# Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox

University of Colorado  
Boulder



University of Colorado  
Boulder