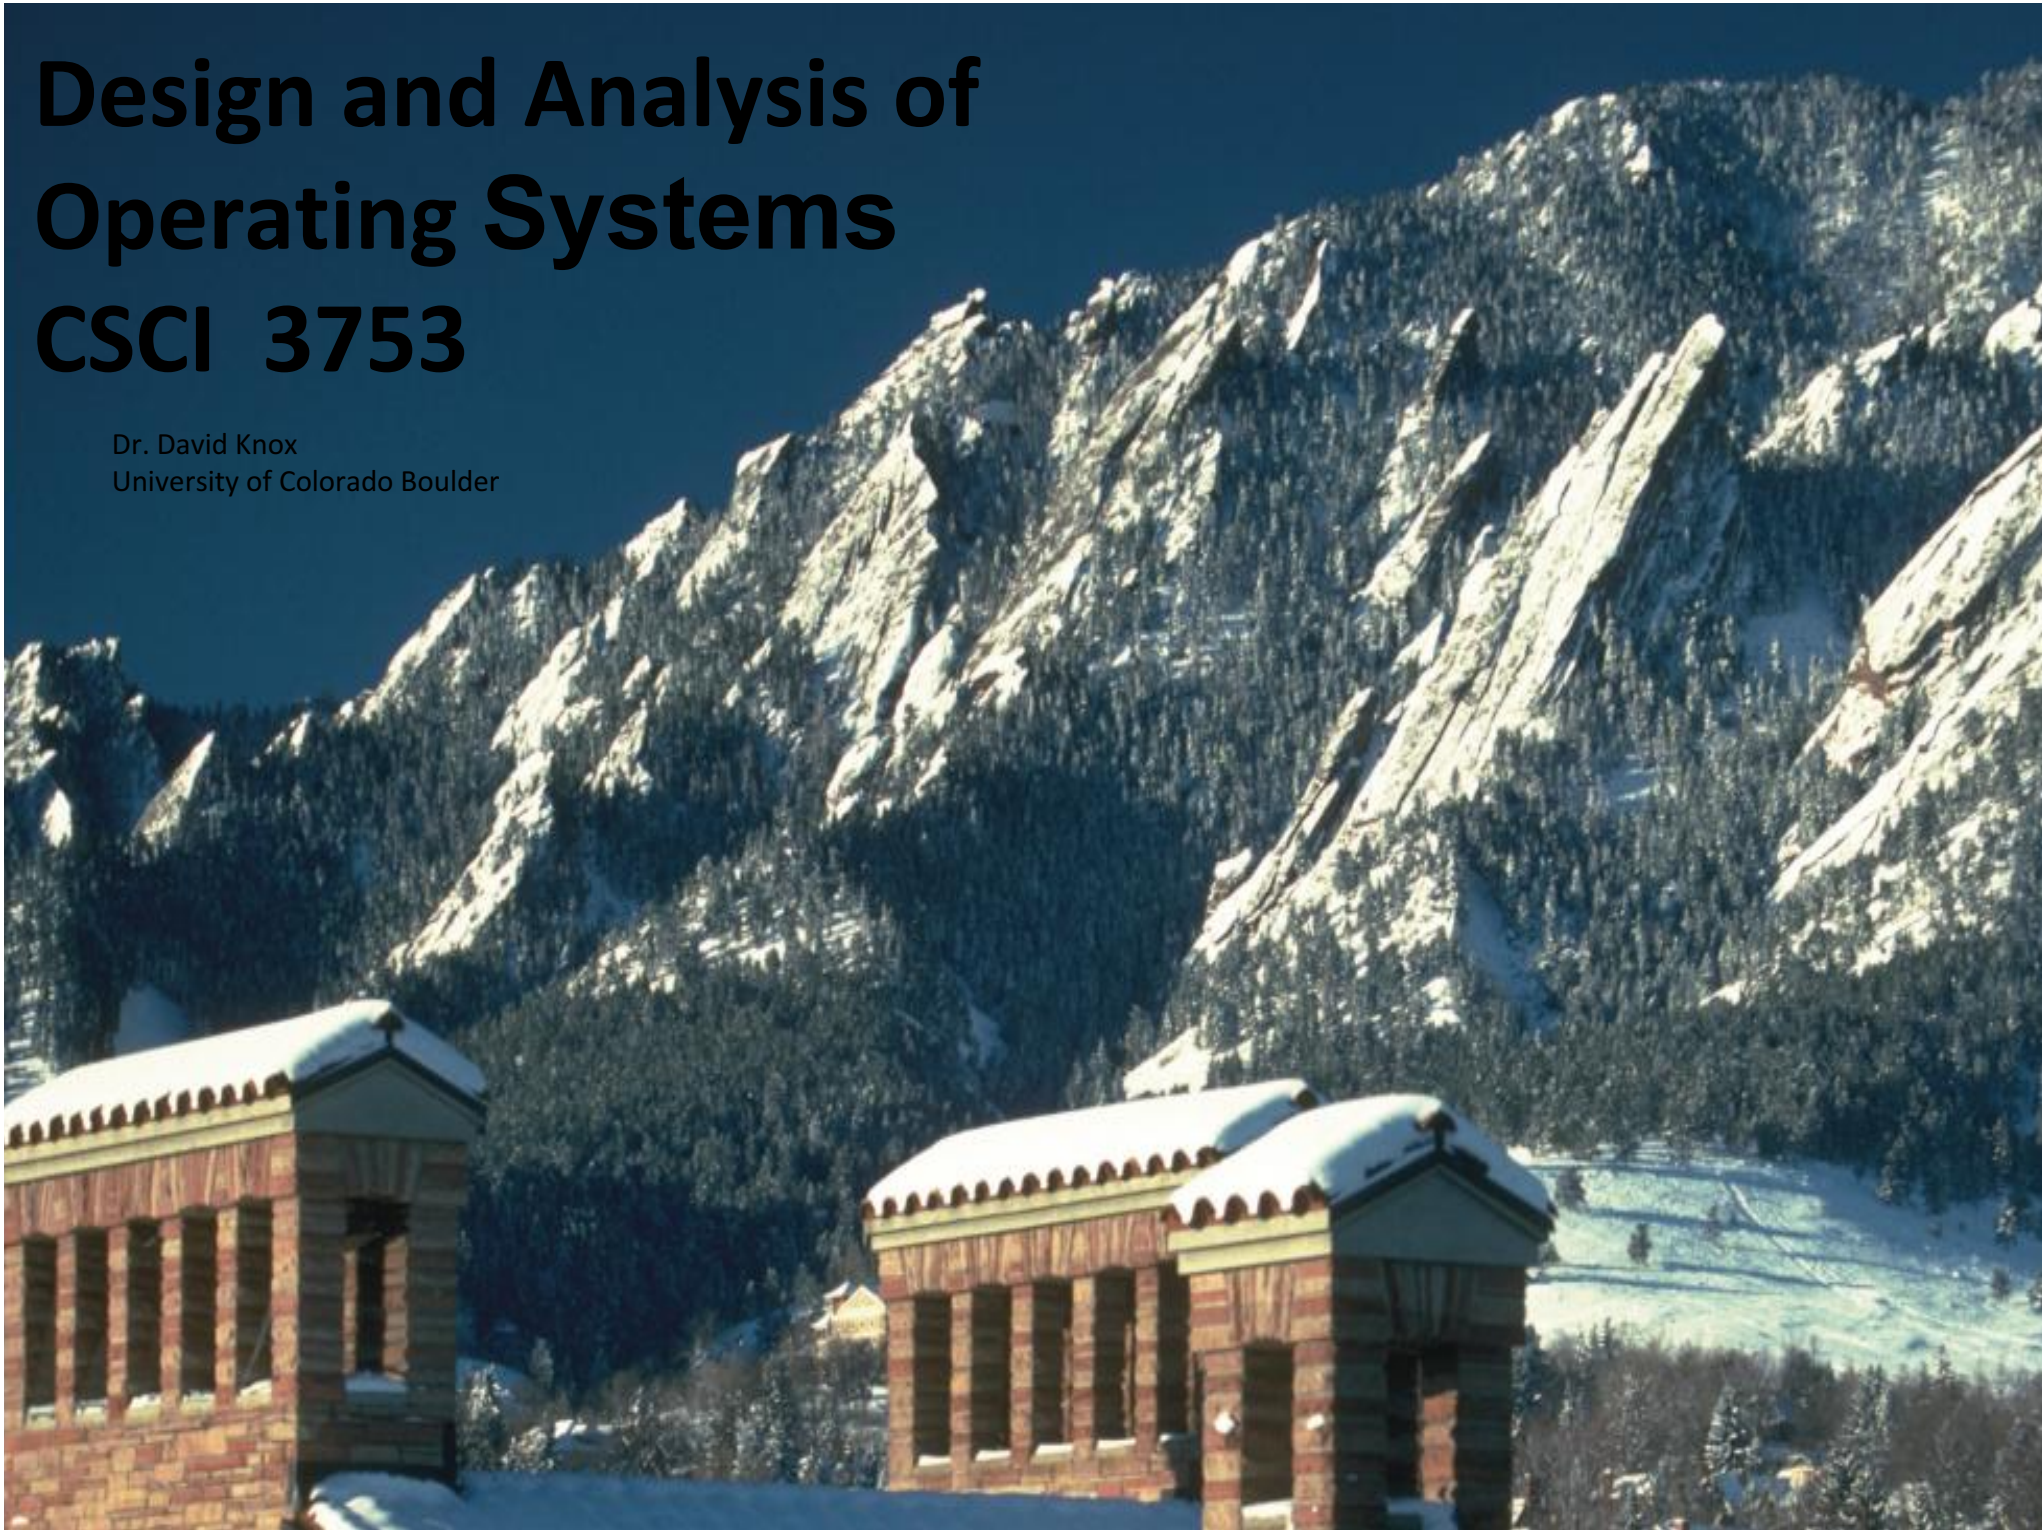


# Design and Analysis of Operating Systems CSCI 3753

Dr. David Knox  
University of Colorado Boulder





Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# **Design and Analysis of Operating Systems CSCI 3753**

**File System Performance,  
Reliability, and Fault Recovery**

Dr. David Knox  
University of  
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.

# **File System Performance, Reliability, and Fault Recovery**

# File System Performance

- **Approaches to improve performance in a file system:**
  - In memory:
    - file header: caching FCB information about open files in memory improves performance (faster access)
    - directory:
      - caching directory entries in memory improves access speed.
      - And hash the directory tree to quickly find an entry and see if it's in memory.

# File System Performance

- On disk:
  - file data: indexed allocation is generally faster than traversing linked list allocation
  - free block list: counting, grouped, linked list allows fast allocation of large # of files

# File System Performance

- **Other potential optimizations:**
  - the disk controller can also have its own cache that stores file data/FCBs/etc. for fast access
  - Cache file data in memory
  - Smarter layout on disk: keep an inode/FCB near file data to reduce disk seeks, and/or file data blocks near each other
  - *read ahead:*
    - if the OS knows this is sequential access, then read the requested page and several subsequent pages into main memory cache in anticipation of future reads



# File System Performance

## ■ Some other potential optimizations:

- *asynchronous writes*: delay writing of file data until sometime later.
- Advantages:
  - removes disk I/O wait time from the critical path of execution,
    - e.g. a write(X) to a file can return quickly rather than waiting for completion of disk I/O
    - allowing the program to move forward in its execution
  - This allows a disk to schedule writes efficiently
    - grouping nearby writes together
  - May avoid a disk write if the data has been changed again soon
  - Note: in certain cases, you may prefer to enforce synchronous writes
    - e.g. when modifying file metadata in the FCB on an open() call

# Memory-Mapped Files

- **Map some parts of a file on disk to pages of virtual memory**
  - Use `mmap()` call
  - First read/write to file data not in memory will on-demand page in the desired blocks of file into main memory pages
  - Subsequent reads/writes to that file data are served quickly by memory
  - Later, flush changes in file to disk
  - Uses virtual memory system
    - fast compared to `read()/write()` system calls to change file data contents




# **File System Reliability and Fault Recovery**

# File System Reliability & Fault Recovery

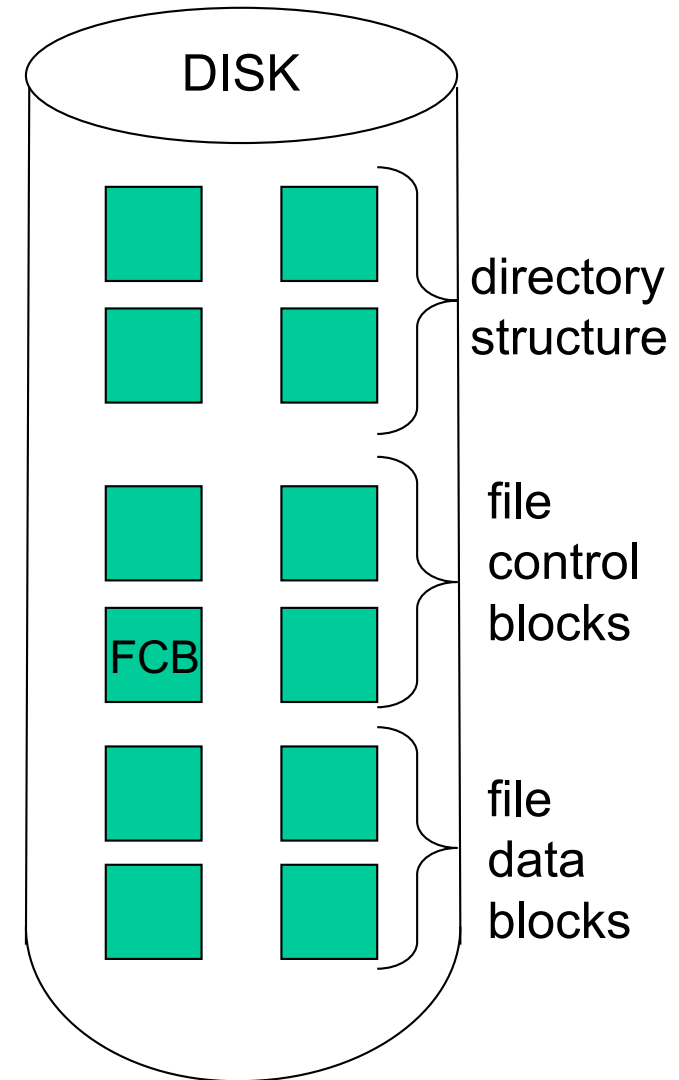
- **In general, OS should gracefully recover from hardware or software failure**
  - The file system needs to be engineered to ensure reliability/fault recovery
- **Problem: File system is quite fragile to system crashes**
  - There is a portion of the file system that is cached in memory
  - This portion may be inconsistent with the complete file system stored on disk

# File System Fragility

In-Memory  
OS File Manager



- All in-memory file system data are lost on a power loss:
  - Directories, file metadata, and file data may all be cached in memory
  - They may all be modified
  - These modifications are lost if they weren't saved to disk



# File System Reliability & Fault Recovery

- Problem #1: asynchronous writes produce inconsistency between in-memory and on-disk file system
- **Promised writes of filed data that were delayed by asynchrony may be lost**
- **Asynchronous writes of directory metadata can create inconsistency between the file system on disk and the writes cached in RAM**
  - Directory information in RAM can be more up to date than disk
  - cached writes may be lost if there is a system failure
    - such as a power loss
    - in this case, the promised writes will not be executed

# File System Reliability & Fault Recovery

- **To address asynchronous write inconsistency,**
  - UNIX caches directory entries for reads
  - But UNIX does not cache any data write that changes metadata or free space allocation

*These changes to critical metadata are written synchronously (immediately) to disk, before the data blocks are written*

- **Problem #2: Even if all writes are synchronous, there is still a consistency problem:**
  - any of the individual synchronous/asynchronous writes to disk can fail halfway through the operation, leaving a half-written directory entry, FCB, or file data block.

# File System Reliability & Fault Recovery

- **Problem #3: Complex operations can create inconsistency while waiting for them to complete**
  - e.g. a file create() involves many operations, and may be interrupted at any time in mid-execution
  - file create() updates the directory, FCB, file data blocks, and free space management
  - if there is a failure after creating the FCB, then the file system is in an inconsistent state because the file data has not yet been saved on disk,
    - i.e. the directory says there is a file and points to the FCB, but the FCB is incomplete because its index block hasn't been fully allocated

# Reliability/Fault Recovery Solutions

- **Approach: file systems can run a consistency checker like fsck in UNIX or chkdsk in MSDOS**
  - in linked allocation, would check each linked list and all FCB's to see if they are consistent with the directory structure.
    - similar checks for indexed allocation
  - Check each allocated file data block to see that its checksum is valid
- **Disadvantages:**
  - This is heavyweight, and takes a long time to check the entire file system.
  - This can detect an error, but doesn't ensure recovery or correction



# Reliability/Fault Recovery Solutions

- **Approach: *log-based recovery* is a solution that helps you *recover from file system failures*:**
  - OS maintains a log or journal on disk of each operation on the file system
  - called log-based or journaling file systems
- **The log on disk is consulted after a failure to reconstruct the file system**
  - In a journaling file system, the log is seen as a separate entity from the file data.
  - In a log-structured system, the log *\*is\** the file system, and there are no separate structures for storing file data and metadata – it's all in the log.

# Log-Based Recovery

# Log-Based Recovery

- Each operation on the file system is written as a record to the log on disk *before* the operation is actually performed on data on disk
  - this is called *write-ahead logging*
- The file system has a sequence of records of operations in the log about what was intended in case of a crash
  - The log contains a sequence of statements like “I’m about to write this directory entry/file header/file data block”,
  - and “I just finished writing this directory/FH/data”.


# Log-Based Recovery

- **Operations are grouped in sets called *transactions***
  - e.g. a file create() has many steps.
    - You either want the entire file created, or not at all if it fails at any step along the way
  - Need the set of steps as a single logical unit
    - It is performed in its entirety or not at all
- **Transactions are viewed as atomic**
  - either succeed in their entirety or not at all

# Log-Based Recovery


- **A transaction  $T_i$  looks like the following:**
  - begins with  $\langle T_i \text{ starts} \rangle$
  - followed by a sequence of records like  $\text{write}(X)$ ,  $\text{read}(Y)$ , ... needed to complete the transaction, e.g. a file  $\text{create}()$
  - ends with  $\langle T_i \text{ commits} \rangle$
- **Write each of these operations to the log**

# Log-Based Recovery



- Each log record of an operation within a transaction consists of:
  - transaction name  $T_i$
  - data item name, e.g. X
  - old value
  - new value
- Both the old and new values must be saved in order for the system to recover from crashes in mid-transaction

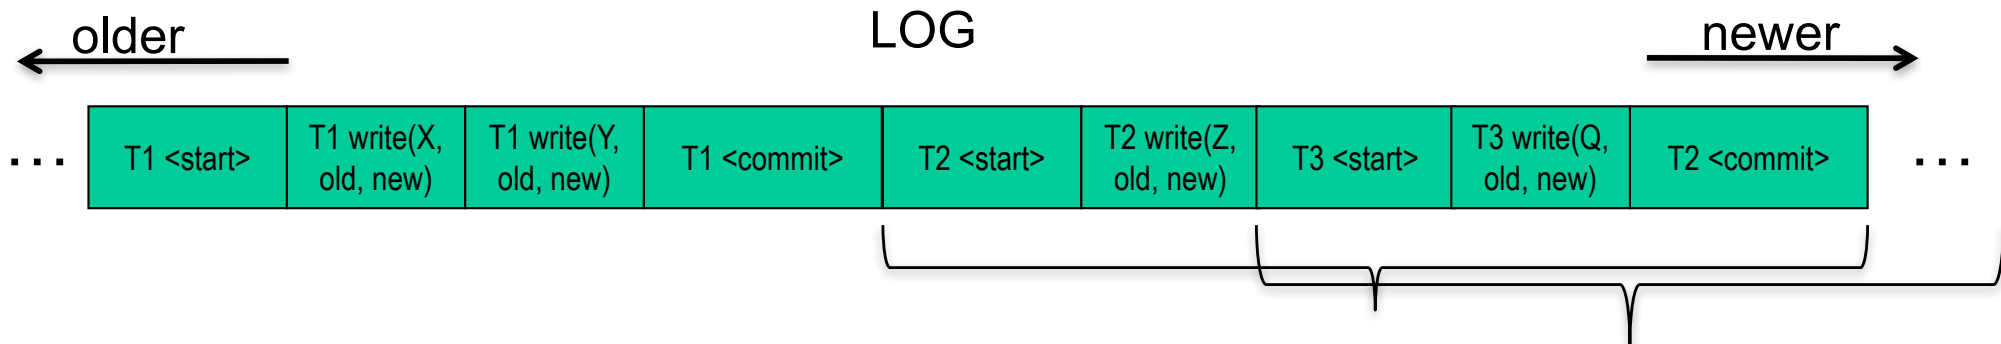
# Log-Based Recovery



- **A transaction is not considered complete until it is committed in the log**
  - once the <commit> appears in the log, then even if the system crashes after this point, there is enough information in the log to fully execute the transaction upon recovery
  - therefore, once the <commit> appears in the log, it is OK to return from the system call that called file create() or file write()

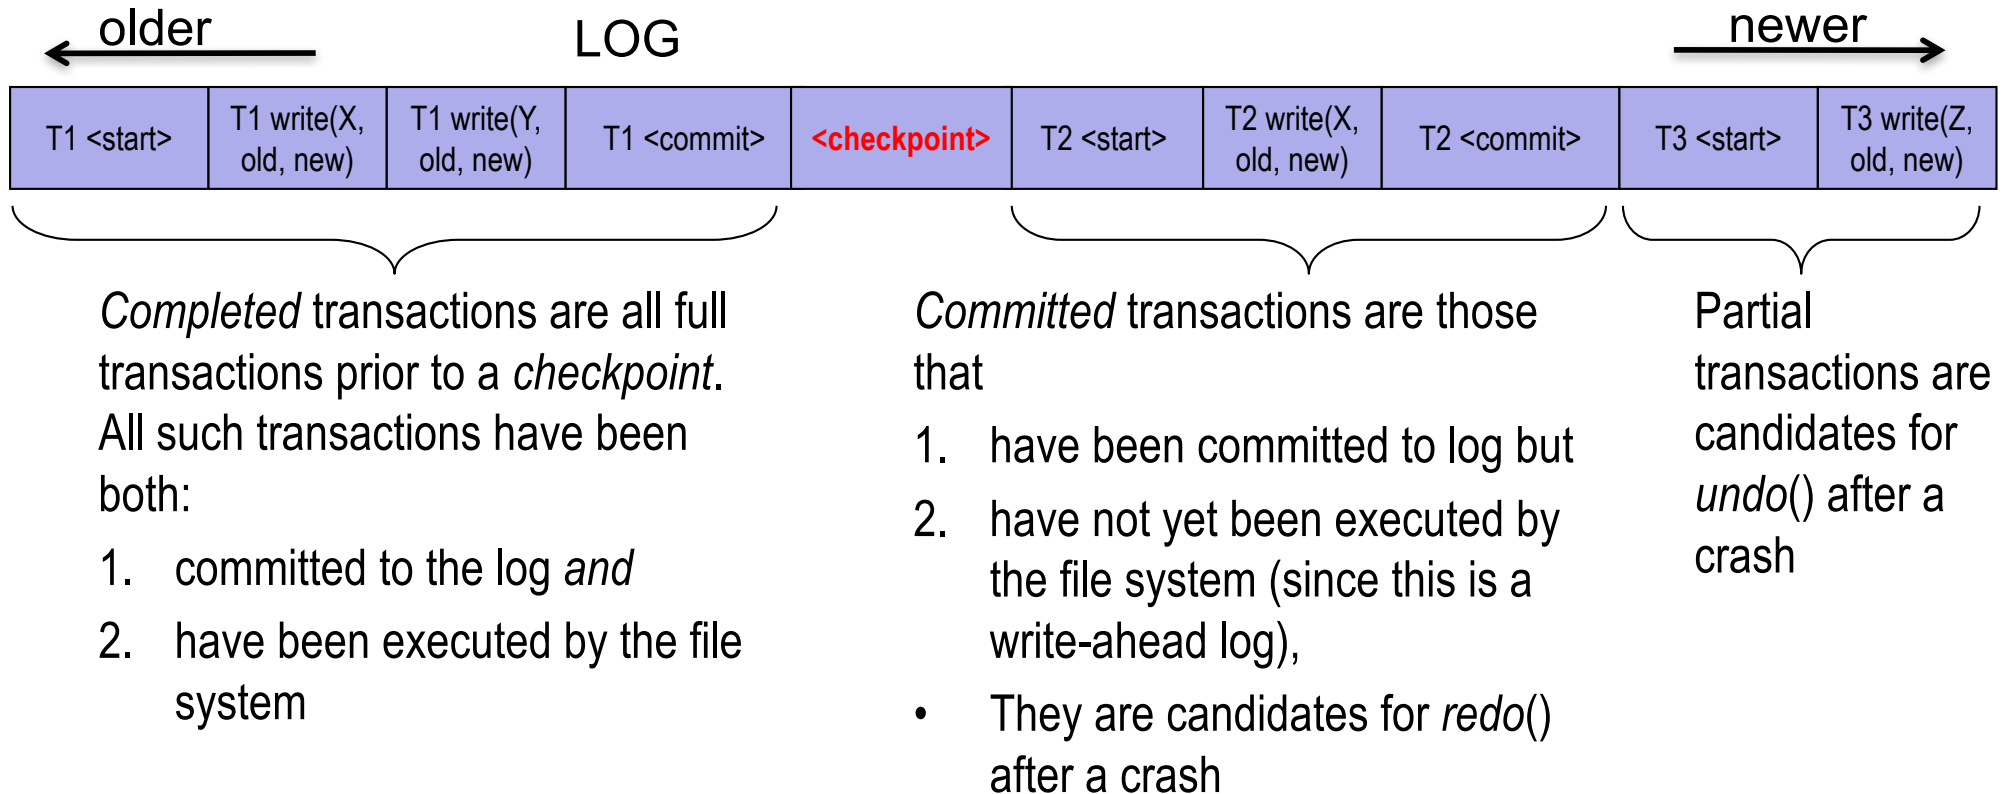


# Log-Based Recovery

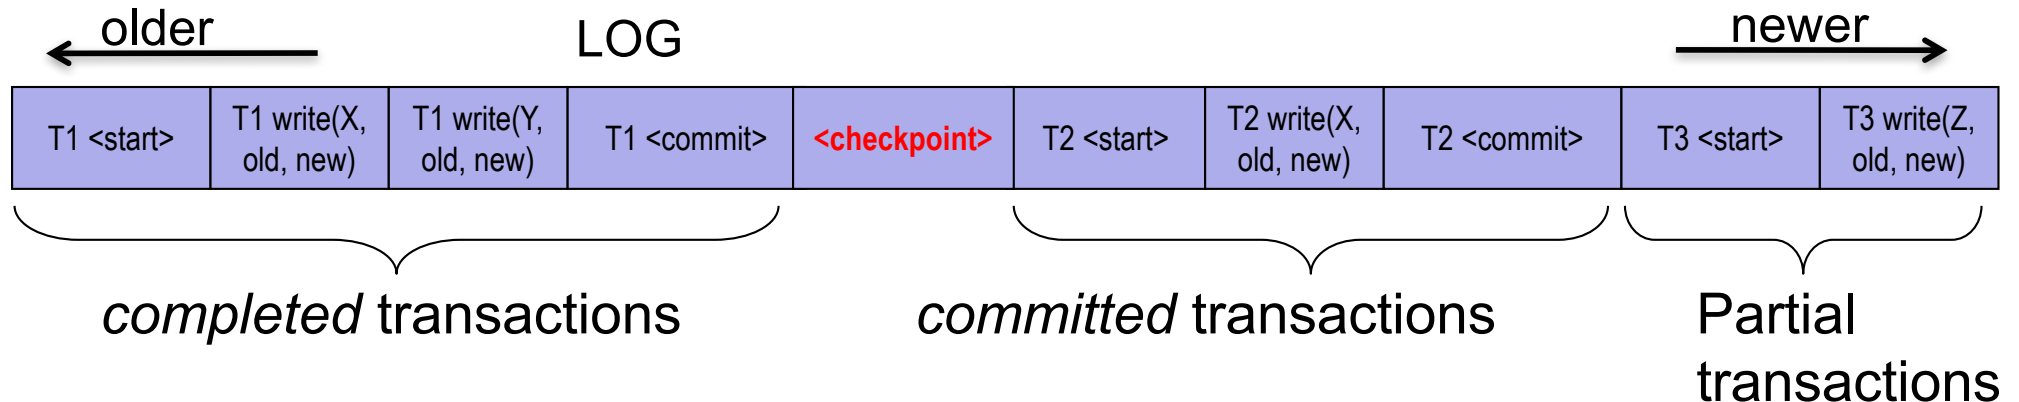


- Operations in different transactions can overlap in the log
- For asynchronous writes, the actual write(X) to disk may occur much later than the entry written to the log

# Log-Based Recovery

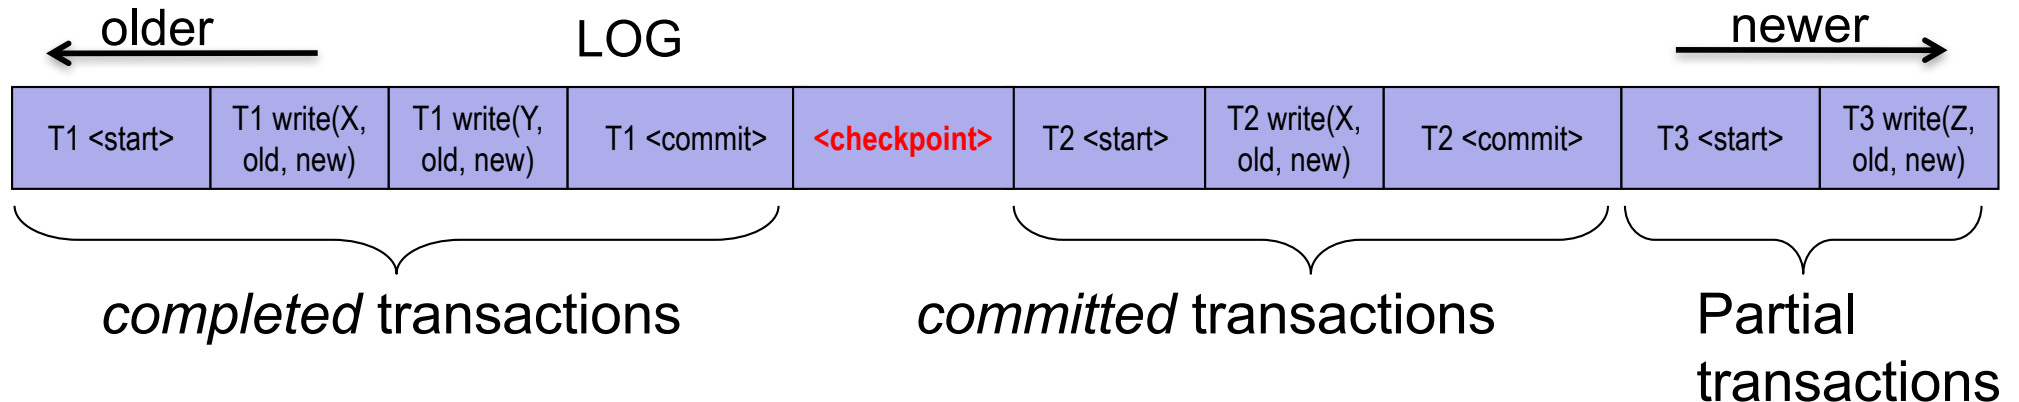


# Log-Based Recovery



- The checkpoint indicates that all full transactions (those with a <start> and <commit>) prior to the checkpoint (to the left) have been written to *both* the disk *and* log
- Committed transactions are full transactions in the log that are to the right of the most recent checkpoint, and thus have not been written to disk yet

# Log-Based Recovery



- In normal operation, the file system will
  - periodically *replay* committed transactions in the log onto disk,
  - Then add a new checkpoint to the log,
  - thus all committed transactions to the left of the newly added checkpoint are converted into completed transactions
  - completed transactions can be removed from the log, or just written over if it's a circular log

# Log-Based Recovery

- On a failure, the OS looks for the latest checkpoint in the log, and redo()'s committed transactions and undo()'s partial transactions from that point on
  - redo() transaction  $T_i$  if the log contains both  $\langle T_i \text{ starts} \rangle$  and  $\langle T_i \text{ commits} \rangle$  and these transactions appear after a checkpoint
  - undo() transaction  $T_k$  if the log contains  $\langle T_k \text{ starts} \rangle$  but not  $\langle T_k \text{ commits} \rangle$ 
    - this is called an aborted transaction
    - during recovery, such a transaction is rolled back to its former state

# Log-Based Recovery

- **Note on an undo(),**
  - the log records contain both the old and new data,
  - so there is enough information in the log to restore or roll back a file or object to its old state if the transaction was not completed
- **OS can cleanly recover the file system**
  - from a checkpoint, reapply changes
- **There is a Write-Twice Penalty with Log-Based**
  - Use log-structured file systems to avoid this penalty
    - The journal is the file system
    - There are no separate data structures
    - The journal contains enough information to be entire file system
  - Flash file systems like JFFS, JFFS2, YAFFS use log-structured file systems

# Journaling File Systems

- **Some file systems like NTFS only write changes to the metadata of a file system to the log**
  - file headers and directory entries only
  - NOT any changes to file data
  - The journal is separate from the main file system
- **Copy-on-write file systems (such as ZFS and Btrfs)**
  - avoid in-place changes to file data by writing out the data in newly allocated blocks
  - followed by updated metadata that would point to the new data and disown the old
  - followed by metadata pointing to updated metadata repeatedly to the root of the file system hierarchy
  - has the same correctness-preserving properties as a journal, without the write-twice overhead.
  - add metadata (checksums) to insure data integrity



# Journaling File Systems

## ■ Linux's ext4fs can be parameterized to operate in 3 modes:

1. *Journal mode*: both metadata and file data are logged. This is the safest mode, but there is the latency cost of two disk writes for every write.
2. *Ordered mode*: only metadata is logged, not file data, and it's guaranteed that file contents are written to disk before associated metadata is marked as committed in the journal.
  - This way, you don't have say a file header pointing to a new block of data, yet that data has not yet been written to disk. This is the default on many Linux distributions.
3. *Writeback mode*: only metadata is logged, not file data, and no guarantee file data written before metadata, so files can become corrupted.
  - This is riskiest mode/least reliable but fastest.



Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Design and Analysis of Operating Systems CSCI 3753



Dr. David Knox  
University of  
Colorado Boulder

Material adapted from: Operating Systems: A Modern Perspective : Copyright © 2004 Pearson Education, Inc.