

CSPB 3202 Artificial Intelligence

Optimization and Tips for Neural Network Training

Geena Kim



Outline and Keywords

Optimization methods

Stochastic Gradient Descent

- Learning rate
- Momentum
- Decay

Tips for neural network training

- General tips for overfitting
- Regularization methods
 - Dropout
 - Batch normalization

Advanced Gradient Descent methods

- Learning rate scheduling
- Nesterov momentum

- AdaGrad
- AdaDelta
- RMSprop
- Adam

Gradient Descent

Optimization Goal

Find a set of (optimized) weights which minimize the error (or loss function) at the output

Weight update rule


$$\dot{W}_{nm}^L \leftarrow W_{nm}^L - \alpha * \delta W_{nm}^L \frac{\partial \mathcal{L}}{\partial W_{nm}^L}$$

Handwritten notes: δW_{nm}^L is circled in red. $\frac{\partial \mathcal{L}}{\partial W_{nm}^L}$ is written in red above the equation. $n+1$ is written in red below the equation.

$$W_{ij} \leftarrow W_{ij} - \alpha \frac{\partial \mathcal{L}}{\partial W_{ij}}$$


Handwritten notes: A red arrow points from the δW_{nm}^L term in the first equation to the δW_{nm}^L term in the second equation. A red arrow points from the $\frac{\partial \mathcal{L}}{\partial W_{nm}^L}$ term in the first equation to the $\frac{\partial \mathcal{L}}{\partial W_{ij}}$ term in the second equation. A red arrow points from the α term in the first equation to the α term in the second equation. A red arrow points from the W_{nm}^L term in the first equation to the W_{ij} term in the second equation. A red arrow points from the \dot{W}_{nm}^L term in the first equation to the W_{ij} term in the second equation.

Global minimum vs. local minimum



Gradient Descent

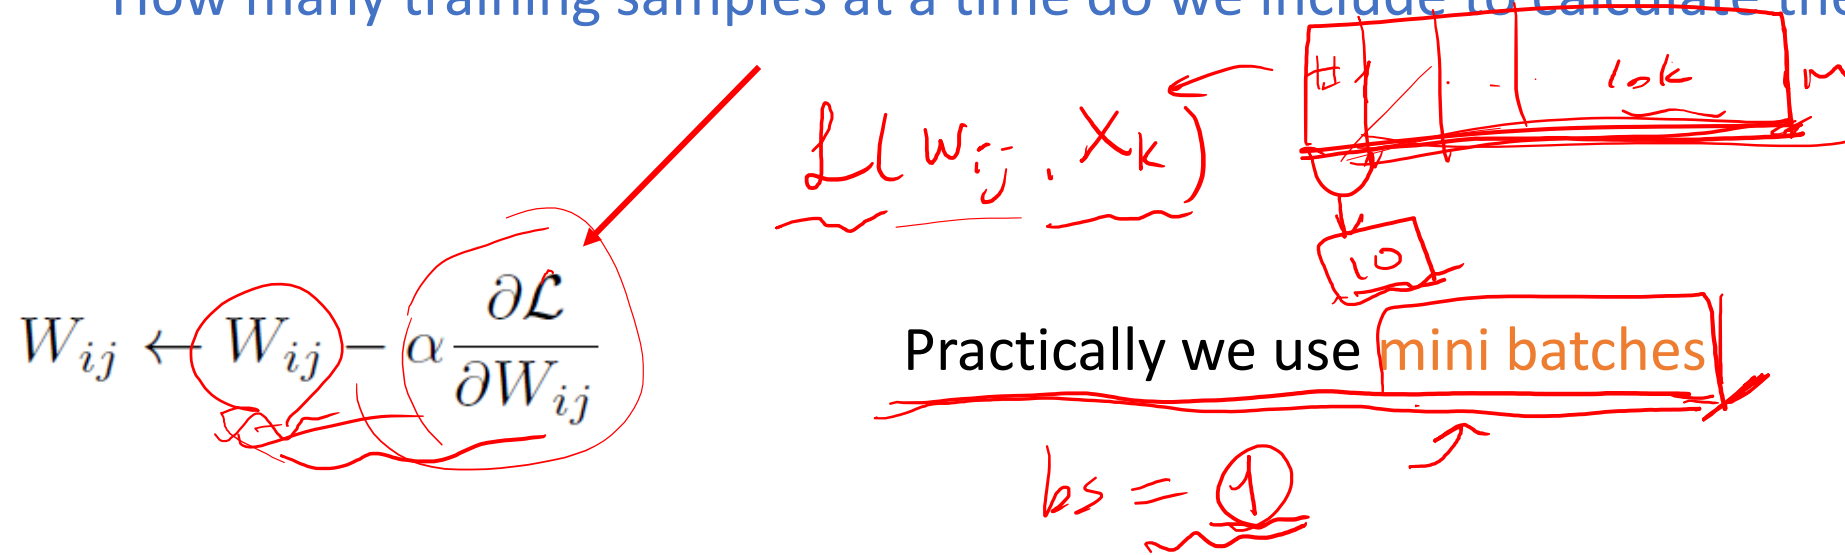
Error surface in the multi dimension can be complicated!



also... cliffs and plateaus.

Stochastic Gradient Descent

How many training samples at a time do we include to calculate the error?



Training speed and accuracy vs. minibatch size

Stochastic Gradient Descent

With decreasing learning rate (Learning rate scheduling)

Algorithm 8.1 Stochastic gradient descent (SGD) update

Require: Learning rate schedule $\epsilon_1, \epsilon_2, \dots$

Require: Initial parameter $\theta \leftrightarrow w$

$k \leftarrow 1$

while stopping criterion not met **do**

Sample a minibatch of m examples from the training set $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ with corresponding targets $\mathbf{y}^{(i)}$.

Compute gradient estimate: $\hat{\mathbf{g}} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

Apply update: $\theta \leftarrow \theta - \epsilon_k \hat{\mathbf{g}}$

$k \leftarrow k + 1$

end while

Stochastic Gradient Descent with momentum

SGD with learning rate alone is slow to converge

Adding a momentum (moving average) can make it faster

$$v \leftarrow \alpha v - \epsilon \nabla_{\theta} \left(\frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta), y^{(i)}) \right), \quad - \sum_{(x)} \frac{\partial L}{\partial \theta}$$
$$\theta \leftarrow \theta + v.$$

warning- different notations used (from deeplearningbook.org)

```
tf.keras.optimizers.SGD(  
    learning_rate=0.01, momentum=0.0, nesterov=False, name='SGD', **kwargs  
)
```

lo

SGD tuning parameters

```
tf.keras.optimizers.SGD(  
    learning_rate=0.01, momentum=0.0, nesterov=False, name='SGD', **kwargs  
)
```

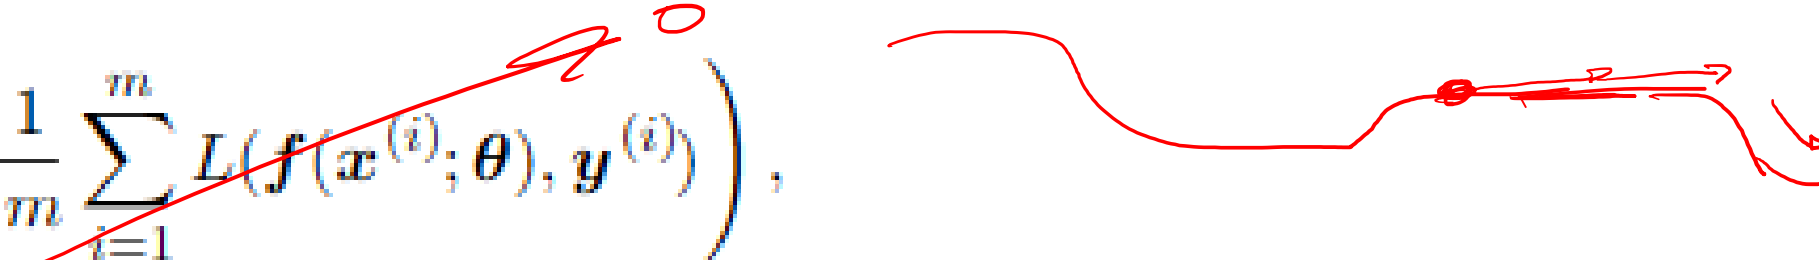
Popular options to tweak

- learning_rate: the base learning rate
- momentum
- decay
- nesterov
- (advanced) callback

Stochastic Gradient Descent with momentum

SGD with learning rate alone is slow to converge

Adding a momentum (moving average of a weight) can make it faster

$$\begin{aligned} \underline{v} &\leftarrow \underline{\alpha v} - \epsilon \nabla_{\theta} \left(\frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)}) \right), \\ \theta &\leftarrow \theta + v. \end{aligned}$$


** see what happens when the gradient is 0 (on plateau)

Stochastic Gradient Descent with decay

Learning rate scheduling using decay

For iteration k (epoch)





$$\underline{\epsilon_k = (1 - \alpha)\epsilon_0 + \alpha\epsilon_\tau} \quad \underline{\alpha = \frac{k}{\tau}} \quad \frac{k}{\tau}$$

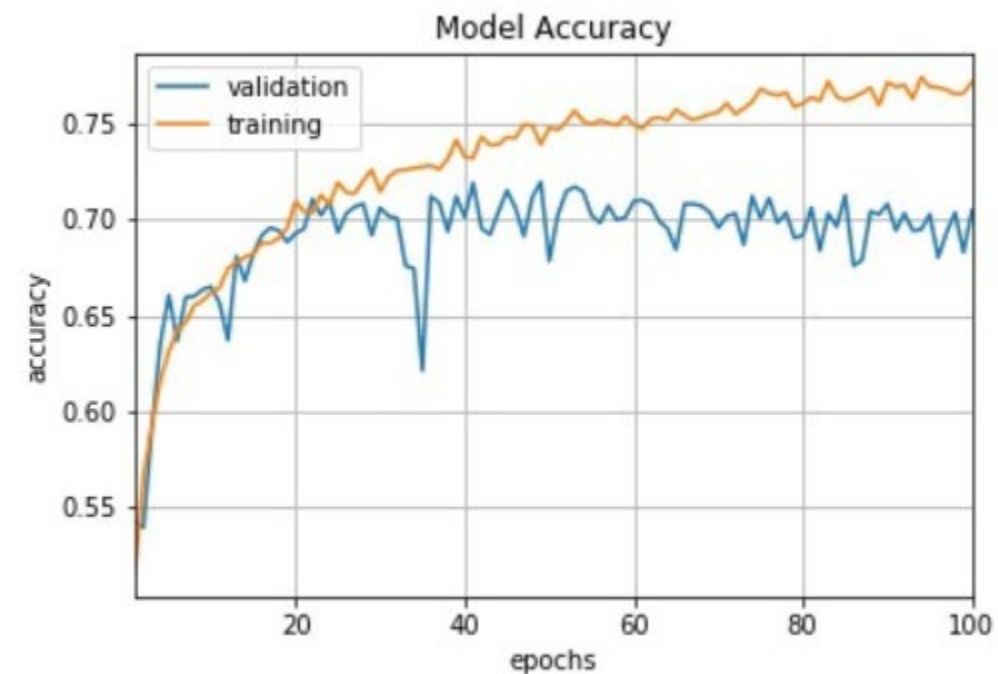
** In the algorithm pseudocode k is for step (each mini batch),
and decay learning rate by step,
but normally we decrease learning rate each epoch



warning- different notations used (from deeplearningbook.org)


Learning rate scheduling

```
tf.keras.optimizers.SGD(  
    learning_rate=0.01, momentum=0.0, nesterov=False, name='SGD', **kwargs  
)
```

learning_rate=0.1, 
momentum=0, decay=0, nesterov=False 



learning_reate=0.1, 
momentum=0.8, decay=learning_rate/epochs 



Learning rate scheduling (custom)

```
tf.keras.callbacks.LearningRateScheduler(  
    schedule, verbose=0  
)
```

```
def step_decay(epoch):  
    initial_lrate = 0.1  
    drop = 0.5  
    epochs_drop = 10.0  
    lrate = initial_lrate * math.pow(drop,  
        math.floor((1+epoch)/epochs_drop))  
    return lrate
```

Ex2

drop lr by half every 10 epochs

```
lrate = LearningRateScheduler(step_decay)
```

```
# This function keeps the learning rate at 0.001 for the first ten epochs  
# and decreases it exponentially after that.  
def scheduler(epoch):  
    if epoch < 10:  
        return 0.001  
    else:  
        return 0.001 * tf.math.exp(0.1 * (10 - epoch))  
  
callback = tf.keras.callbacks.LearningRateScheduler(scheduler)  
model.fit(data, labels, epochs=100, callbacks=[callback],  
    validation_data=(val_data, val_labels))
```

Ex1

Learning rate scheduling (custom)



```
tf.keras.callbacks.LearningRateScheduler(  
    schedule, verbose=0  
)
```

```
def step_decay(epoch):  
    initial_lrate = 0.1  
    drop = 0.5  
    epochs_drop = 10.0  
    lrate = initial_lrate * math.pow(drop,  
        math.floor((1+epoch)/epochs_drop))  
    return lrate
```

Ex2

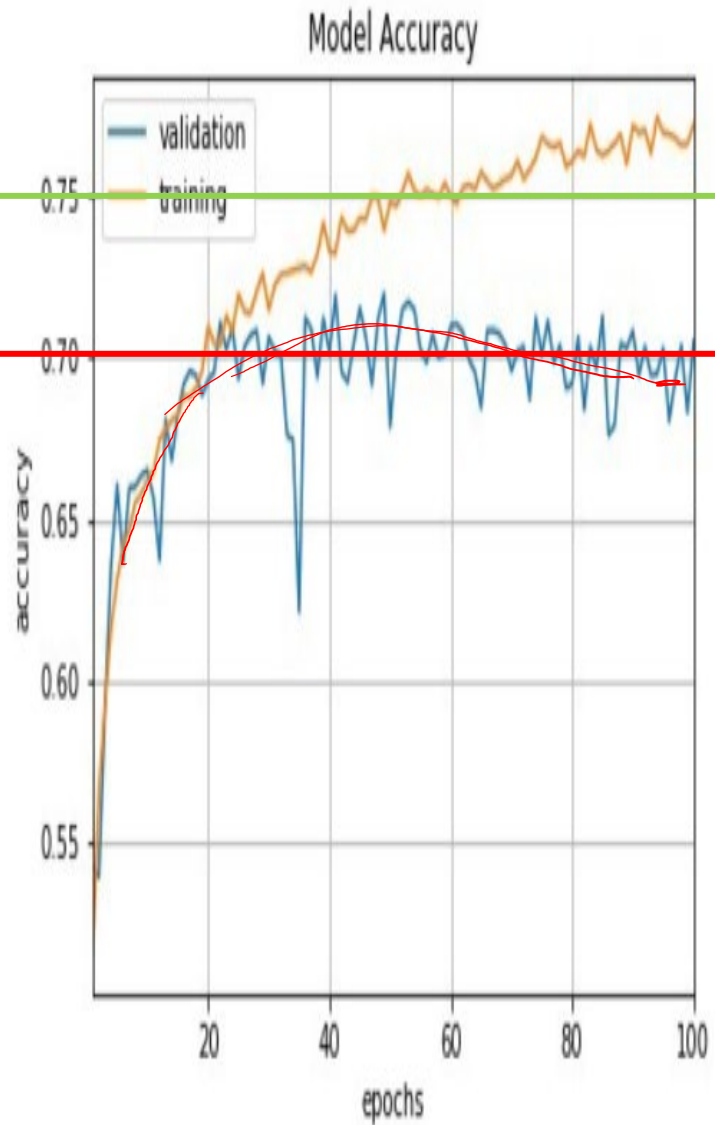
drop lr by half every 10 epochs

```
lrate = LearningRateScheduler(step_decay)
```

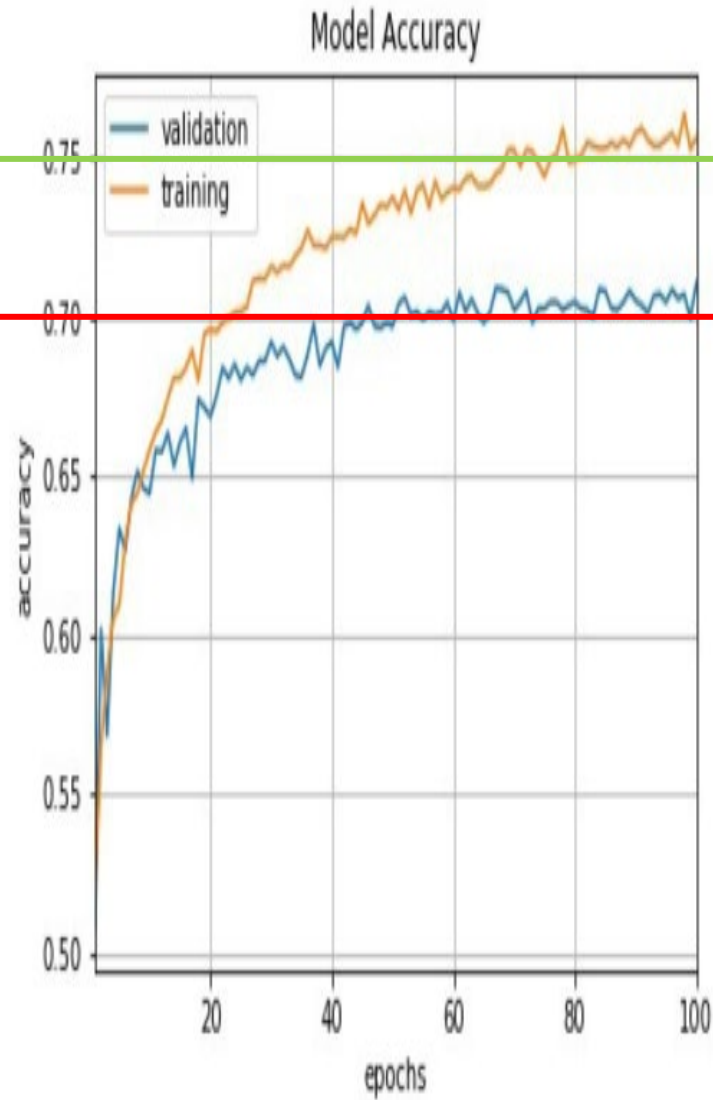


comparison


Base (fixed lr)



with momentum and decay



custom learning rate schedule (step drop)



Nestrov momentum

Nestrov momentum does early correction on gradient

It's supposed to make converge faster, but on SGD it doesn't do much

Regular momentum

$$\begin{aligned} \mathbf{v} &\leftarrow \alpha \mathbf{v} - \epsilon \nabla_{\boldsymbol{\theta}} \left(\frac{1}{m} \sum_{i=1}^m L(\mathbf{f}(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}) \right) \\ \boldsymbol{\theta} &\leftarrow \boldsymbol{\theta} + \mathbf{v}. \end{aligned}$$

Nestrov momentum

$$\begin{aligned} \mathbf{v} &\leftarrow \alpha \mathbf{v} - \epsilon \nabla_{\boldsymbol{\theta}} \left[\frac{1}{m} \sum_{i=1}^m L\left(\mathbf{f}(\mathbf{x}^{(i)}; \boldsymbol{\theta} + \alpha \mathbf{v}), \mathbf{y}^{(i)}\right) \right], \\ \boldsymbol{\theta} &\leftarrow \boldsymbol{\theta} + \mathbf{v}, \end{aligned}$$

Adagrad

learning rate is normalized by the sqrt of the total sum of the gradient

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii}} + \epsilon} \cdot g_{t,i}$$

Advanced optimization

Adadelta

learning rate is normalized by the RMS of the gradient
Weight change is proportional to the RMS ratio

$$\Delta\theta_t = -\frac{\eta}{\underbrace{RMS[g]_t}_{\sim}} \underbrace{g_t}_{\sim}$$

$$\Delta\theta_t = -\frac{RMS[\Delta\theta]_{t-1}}{RMS[g]_t} g_t$$
$$\theta_{t+1} = \theta_t + \Delta\theta_t$$

Advanced optimization

RMSprop

Variant of Adadelta

RMSprop takes a moving average when it calculate the RMS of the gradient

$$E[g^2]_t = \underbrace{0.9}_{\eta} E[g^2]_{t-1} + \underbrace{0.1}_{\eta} g_t^2$$
$$\theta_{t+1} = \theta_t - \frac{\eta g_t}{\sqrt{E[g^2]_t + \epsilon}}$$

An overview of gradient descent optimization algorithms

<https://arxiv.org/pdf/1609.04747.pdf>

warning- different notations used

Advanced optimization

Adaptive Moment Estimation (Adam)

Mimics momentum for gradient and gradient-squared

m_t and v_t are estimates of the first moment (the mean) and the second moment (the uncentered variance) of the gradients

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \end{aligned}$$

$$w \leftarrow w - \eta \cdot g$$

$$w = \frac{\eta}{\sqrt{E[g^2] + \epsilon}} g$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t + \epsilon}} \hat{m}_t$$

momentum

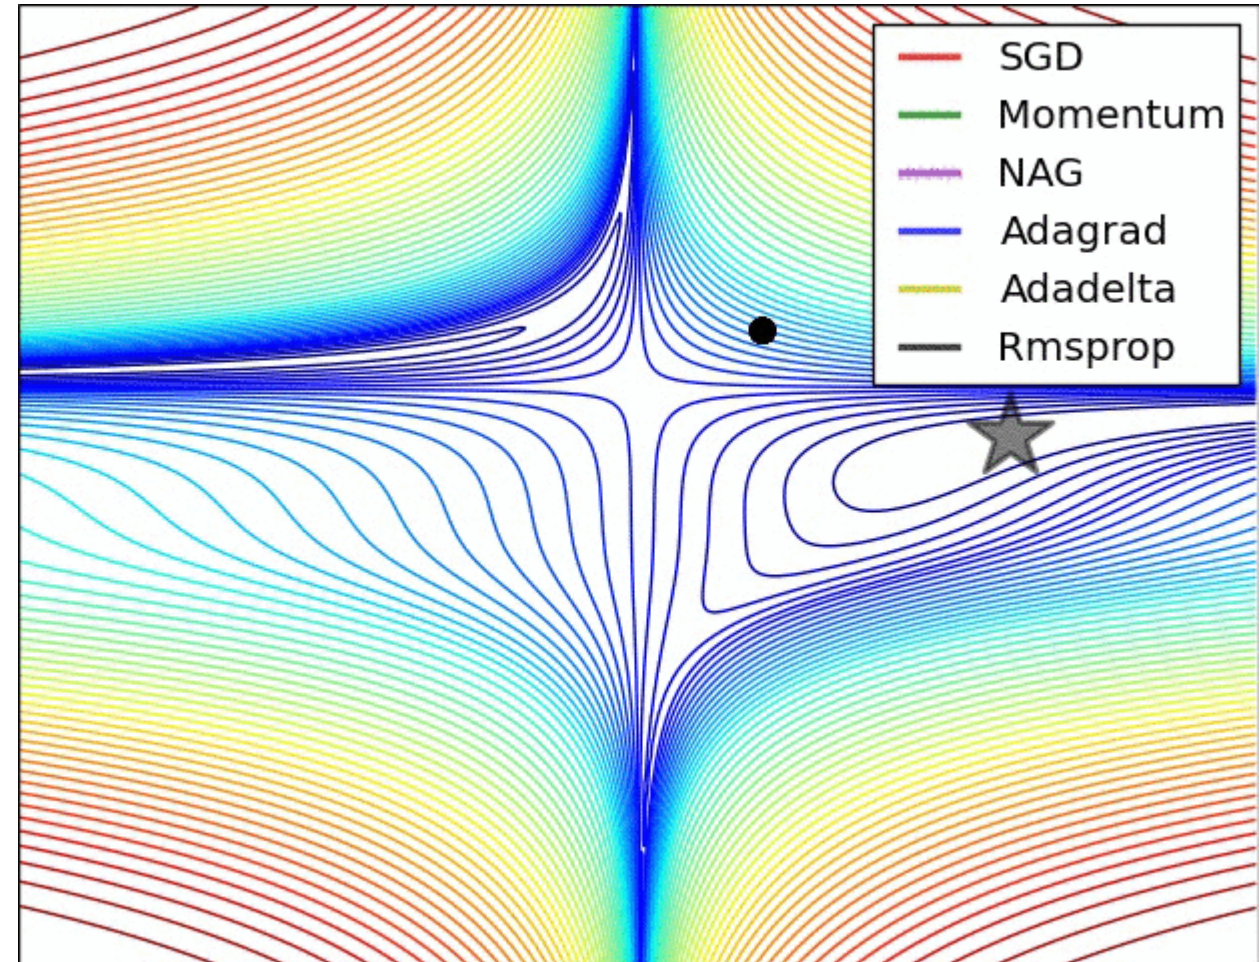
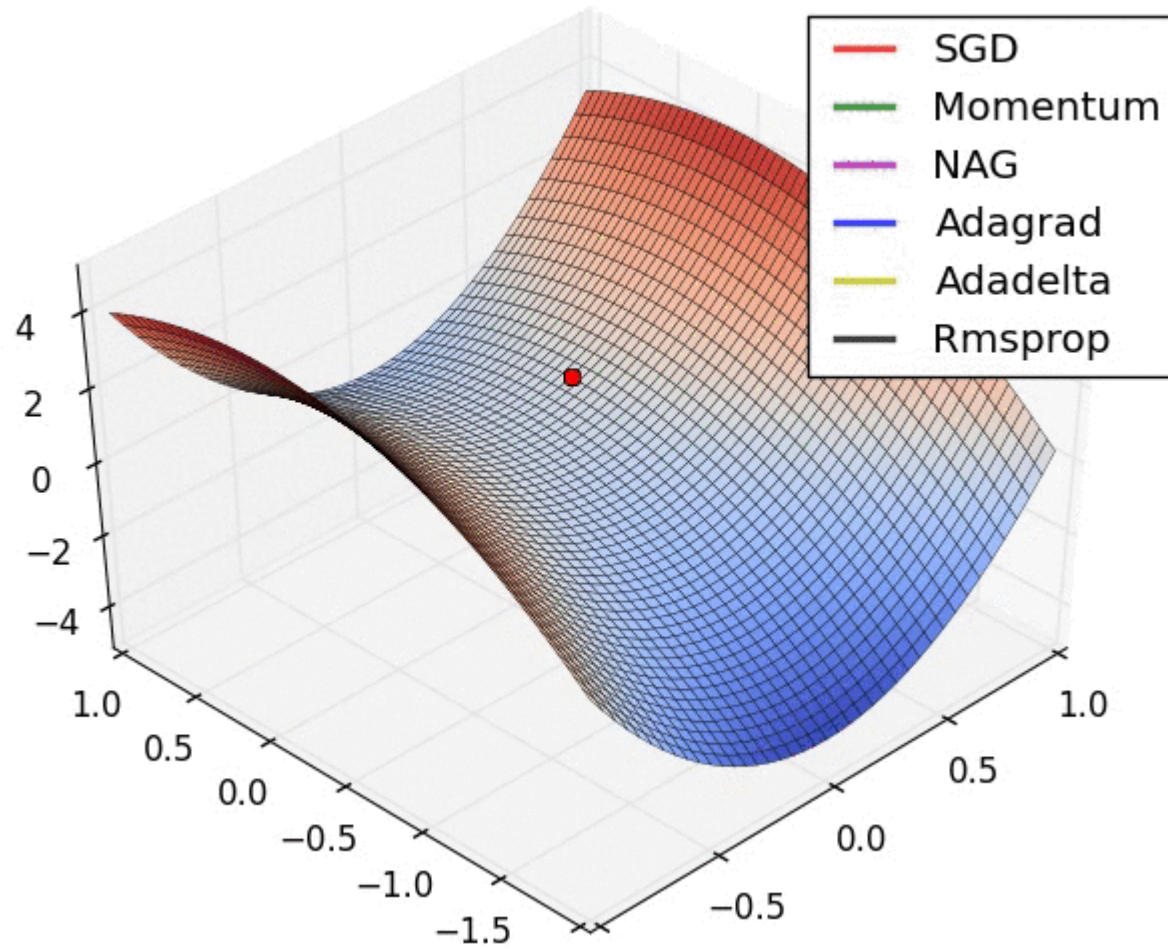
RMS g $E(g^2)$

An overview of gradient descent optimization algorithms

<https://arxiv.org/pdf/1609.04747.pdf>

warning- different notations used

Advanced optimization



animated image source: <https://imgur.com/a/Hqolp>

Tips for training NN

Monitor overfitting as epoch goes

Train hyperparameter tuning : learning rate and other hyperparams

Architecture hyperparameter tuning : NN architecture, # layers, # neurons, activation ft, etc

Try different optimization methods

Regularization: Dropout and Batch Normalization,
or add L1/L2 reg on the loss

$$\sum |a| \rightarrow \sum a^2$$

Monitoring Overfitting in Training







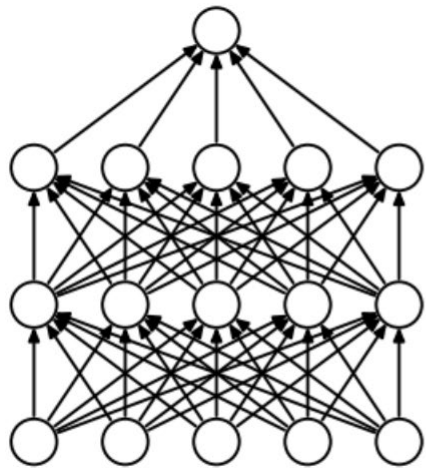
Diagram credit: Fei-Fei Li

Monitoring Overfitting in Training




Ways to reduce overfitting

Dropout [1]




(a) Standard Neural Net



(b) After applying dropout.

Batch normalization [2]



[1] <http://jmlr.csail.mit.edu/papers/volume15/srivastava14a/srivastava14a.pdf>

[2] <https://arxiv.org/pdf/1502.03167.pdf>