

## 11.1 Direct-address tables

Direct addressing is a simple technique that works well when the universe  $U$  of keys is reasonably small. Suppose that an application needs a dynamic set in which each element has a key drawn from the universe  $U = \{0, 1, \dots, m-1\}$ , where  $m$  is not too large. We shall assume that no two elements have the same key.

To represent the dynamic set, we use an array, or *direct-address table*, denoted by  $T[0..m-1]$ , in which each position, or *slot*, corresponds to a key in the universe  $U$ . Figure 11.1 illustrates the approach; slot  $k$  points to an element in the set with key  $k$ . If the set contains no element with key  $k$ , then  $T[k] = \text{NIL}$ .

The dictionary operations are trivial to implement:

DIRECT-ADDRESS-SEARCH( $T, k$ )

1 **return**  $T[k]$


DIRECT-ADDRESS-INSERT( $T, x$ )

1  $T[x.\text{key}] = x$

DIRECT-ADDRESS-DELETE( $T, x$ )

1  $T[x.\text{key}] = \text{NIL}$

Each of these operations takes only  $O(1)$  time.



**Figure 11.1** How to implement a dynamic set by a direct-address table  $T$ . Each key in the universe  $U = \{0, 1, \dots, 9\}$  corresponds to an index in the table. The set  $K = \{2, 3, 5, 8\}$  of actual keys determines the slots in the table that contain pointers to elements. The other slots, heavily shaded, contain NIL.

For some applications, the direct-address table itself can hold the elements in the dynamic set. That is, rather than storing an element's key and satellite data in an object external to the direct-address table, with a pointer from a slot in the table to the object, we can store the object in the slot itself, thus saving space. We would use a special key within an object to indicate an empty slot. Moreover, it is often unnecessary to store the key of the object, since if we have the index of an object in the table, we have its key. If keys are not stored, however, we must have some way to tell whether the slot is empty.

### Exercises

#### 11.1-1

Suppose that a dynamic set  $S$  is represented by a direct-address table  $T$  of length  $m$ . Describe a procedure that finds the maximum element of  $S$ . What is the worst-case performance of your procedure?

#### 11.1-2

A **bit vector** is simply an array of bits (0s and 1s). A bit vector of length  $m$  takes much less space than an array of  $m$  pointers. Describe how to use a bit vector to represent a dynamic set of distinct elements with no satellite data. Dictionary operations should run in  $O(1)$  time.

#### 11.1-3

Suggest how to implement a direct-address table in which the keys of stored elements do not need to be distinct and the elements can have satellite data. All three dictionary operations (INSERT, DELETE, and SEARCH) should run in  $O(1)$  time. (Don't forget that DELETE takes as an argument a pointer to an object to be deleted, not a key.)

#### 11.1-4 ★

We wish to implement a dictionary by using direct addressing on a *huge* array. At the start, the array entries may contain garbage, and initializing the entire array is impractical because of its size. Describe a scheme for implementing a direct-address dictionary on a huge array. Each stored object should use  $O(1)$  space; the operations SEARCH, INSERT, and DELETE should take  $O(1)$  time each; and initializing the data structure should take  $O(1)$  time. (*Hint:* Use an additional array, treated somewhat like a stack whose size is the number of keys actually stored in the dictionary, to help determine whether a given entry in the huge array is valid or not.)